



University of Zurich^{UZH}

High Performance Computing | Lecture 2

Douglas Potter
Jozef Bucko
Noah Kubli



Last
Week

What is HPC?

What is a Supercomputer?

Connecting to Piz Daint & Eiger

SSH & X11

Editors

Unix / Linux

Today's Topics



Revision Control Systems

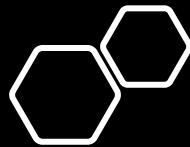
How compilers work

Piz Daint “Environment”

Batch Queues

Processes & Threads

Compiling and running code



Online Resources

<http://goalkicker.com/>

- [Git Notes for Professionals](#)
- [Linux Notes for Professionals](#)
- [Bash Notes for Professionals](#)
- [C Notes for Professionals](#)

[GNU Make](#)

- [GNU Autoconf](#)
- [CMake](#)

[CSCS User Portal](#)

- [SLURM Jobscript Generator](#)

GitHub

Source Code Management

What/Why of Revision Control

You change something and now it's broken

You forgot what you changed

You start something and want to go back

You want to know exactly what was changed

You want to know who changed something

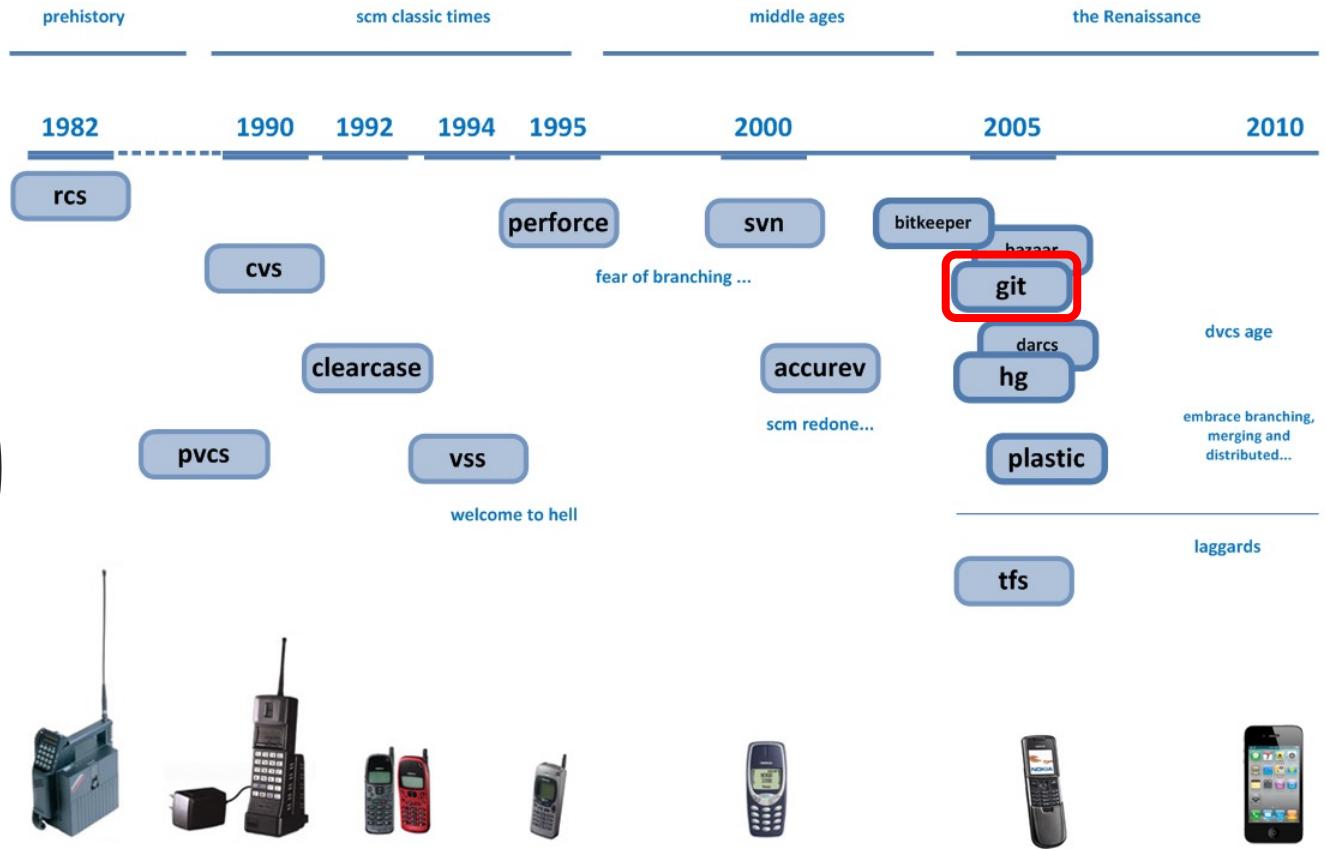
You want to "undo" an old change

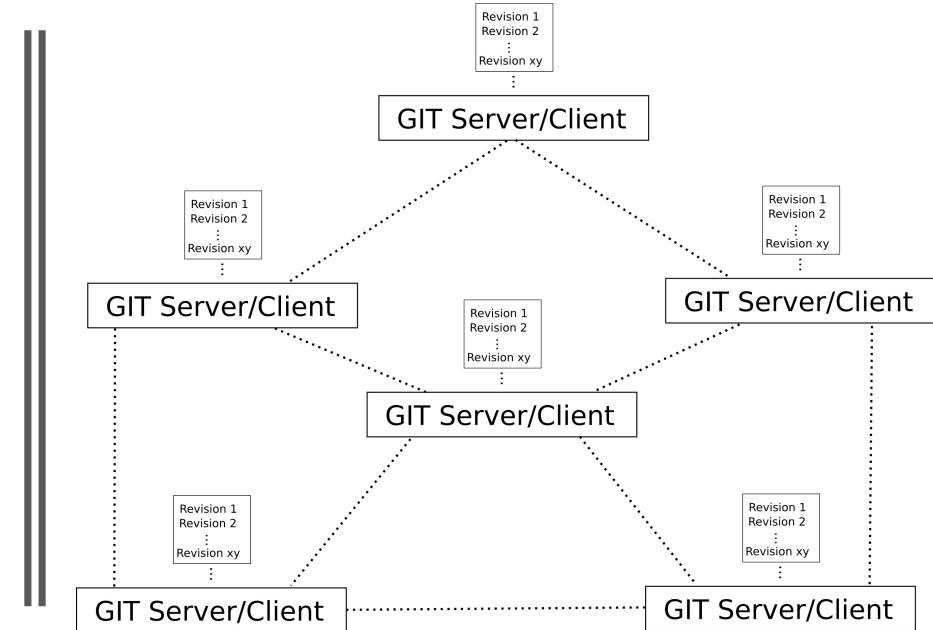
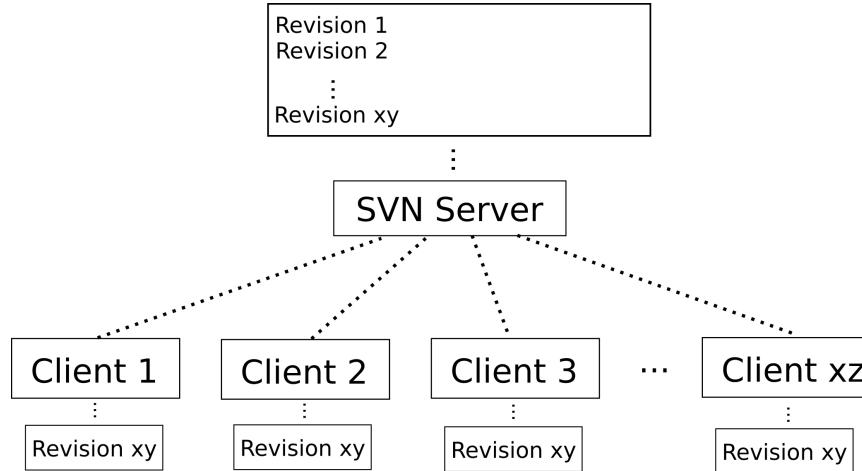
You "accidentally" something 😊

You want to work in parallel (different versions)

You want to work as part of a team

SCM History

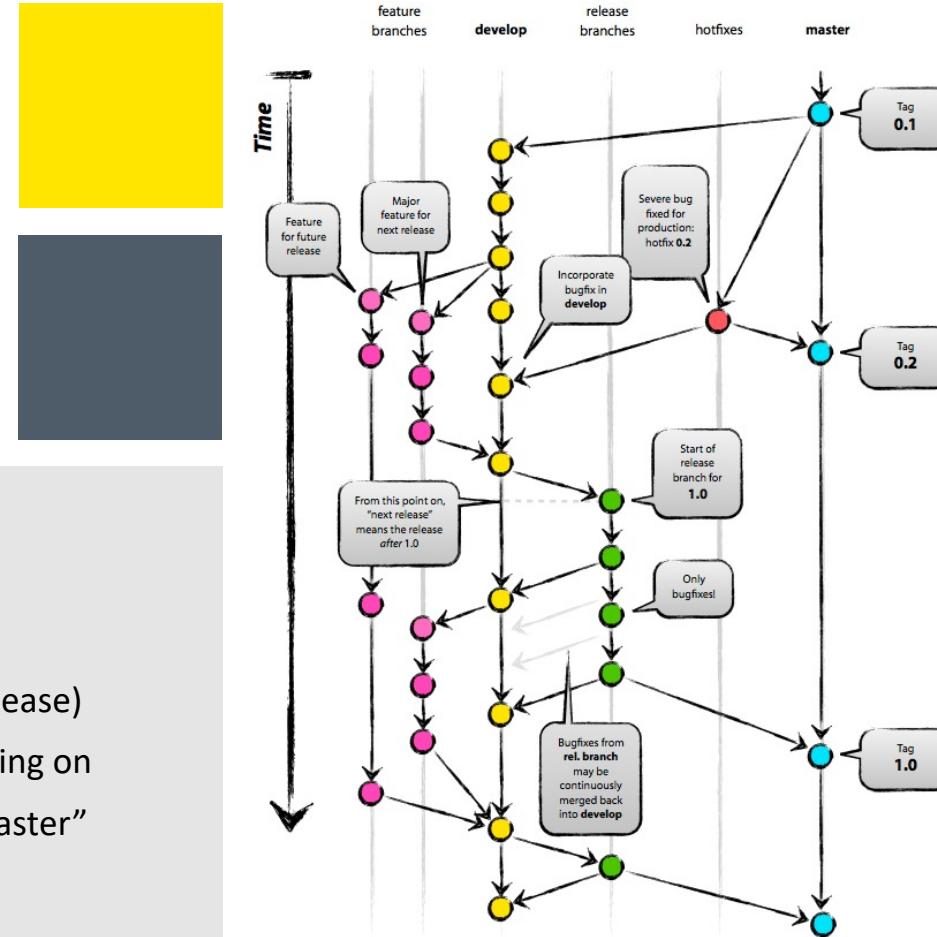




Git versus Subversion

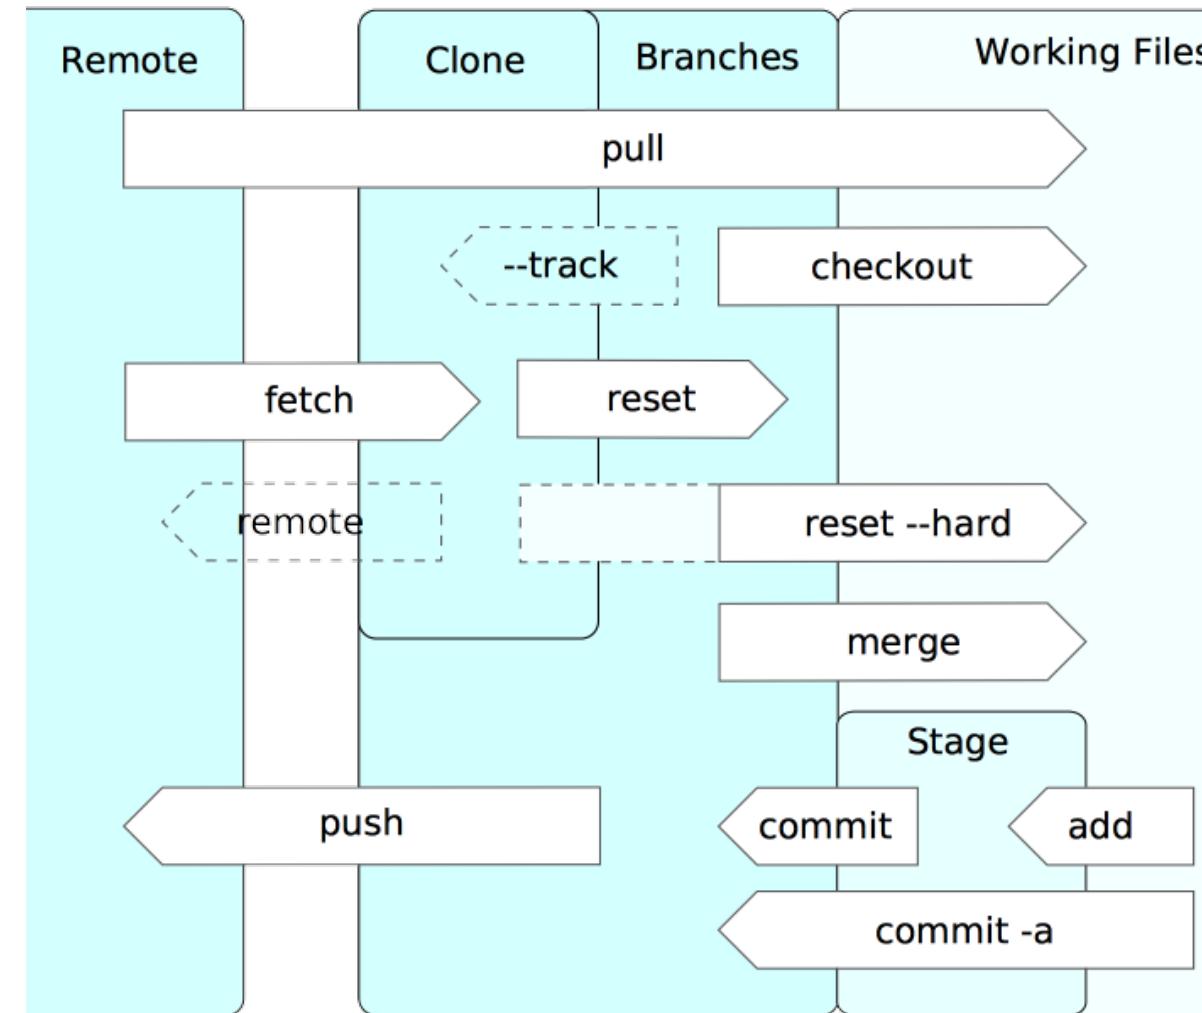
Branches

- git-flow (both an extension and a philosophy)
- “master” branch – always stable working code
- “develop” branch – work in progress (waiting release)
- “feature” branches – new features you are working on
- “release” branches – versions that end up in “master”



Basic Git Operations

- “Working Files” are what you are probably used to.
- Clone/Branches are a local copy of the complete history.
- Remote
 - github.com
 - bitbucket.org
 - gitlab.uzh.ch
 - Or any other computer (or directory even!)



Almost Linear Development

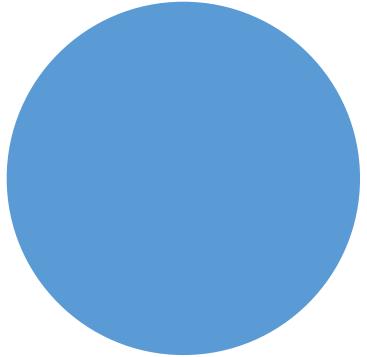
Author	Commit	Message	Date	Builds
Doug Potter	4697bf1	Add openpa COPYRIGHT file	2018-02-14	✓
Joachim Stadel	1dfc65e	Went one step too far in the loop since D1[i+1] was set!	2018-02-13	✓
Doug Potter	1a4269e	MSC macros only under MSC	2018-02-13	✓
Doug Potter	58292c7	Now should compile under native Windows	2018-02-13	✓
Doug Potter	c222b31	More native windows compilation work (still in progress)	2018-02-13	✓
Doug Potter	191fb00	You cannot use a variable for size of local variables	2018-02-13	✓
Doug Potter	ab78c44	Why was mpi included here?	2018-02-13	✓
Doug Potter	71014b5 M	merge	2018-02-12	✓
Doug Potter	edff8a7	Some updates to support native windows	2018-02-12	✓
Doug Potter	dfe9661	Start of integrating base CUDA into MDL	2018-02-08	✓
Doug Potter	5d62d1c	Set sensible defaults when present	2018-02-07	✓
Doug Potter	0062ef1	Simple Windows instructions	2018-02-07	✓
Doug Potter	f1c8375	Add pkdgrav3 to the installation	2018-02-07	✓
Doug Potter	fc06662	Compilation on Windows	2018-02-07	✓
Doug Potter	c262c4f	A little to overzealous with the separation (for now)	2018-02-06	✓
Doug Potter	a6b5eeb	Separate generic CUDA routines better	2018-02-06	✓
Doug Potter	69fb03b	compilation fixes for Cray under cmake	2018-02-06	✓
Doug Potter	76acd63	Fix string literal problems	2018-02-06	✓
Doug Potter	4b8ce39	cmake should now work (perhaps everywhere)	2018-02-06	✓
Doug Potter	607ba6d	Preliminary cmake support (not yet completely ready)	2018-02-06	✓
Doug Potter	0fe6d50	Remove null/pthread build targets (there weren't used)	2018-02-05	✓
Doug Potter	63917ff	Features/macros cleanup	2018-02-05	✓
Doug Potter	1159819	Removed COOLING and Fortran support	2018-02-05	✓
Doug Potter	ab9106f	CHANGESOFT is not used	2018-02-05	✓
Doug Potter	82e60f0	Removed global group assignment code. This will be done diff...	2018-02-02	✓
Doug Potter	c60199c	Removed nLowTot/nHighTot from the pst; they were never used.	2018-02-02	✓
Doug Potter	f1d372f	Make argv processing optional by passing NULL	2018-02-01	✓
Doug Potter	a967fb1	mdl2 can now be compiled stand-alone with cmake	2018-01-31	✓

Parallel Development

Doug Potter	dcbc312	Remove old Fof parameters	2018-01-30	✓
Doug Potter	c4bee8d	Removed old group finding	2018-01-30	✓
Doug Potter	5034d85	Removed obsolete null and pthread versions	2018-01-30	✓
Doug Potter	9b25957	Moved mac directory into mdl2	2018-01-30	✓
Doug Potter	6f4fd5a	this Makefile is not used	2018-01-30	✓
Doug Potter	35f476a	Remove DeepestPot	2018-01-26	✓
Doug Potter	2d59b4f	Removed HSDKD	2018-01-26	✓
Doug Potter	1e6313a	inline functions need inline	2018-01-26	
Doug Potter	466d9b0 M	Merge branch 'master' of bitbucket.org:dpotter/pkdgrav3	2018-01-26	✓
Doug Potter	95831d3	Some fixes for Python3.	2018-01-26	
Doug Potter	eb21e5c	OpenCL include fix	2018-01-26	
Mischa Knabe...	8ac9d3a M	Merged in mknab/pkdgrav3 (pull request #5) Master	2018-01-25	✓
Mischa Knabe...	47a8861	Deleted MK_csmExp2Hub as it was not used in the end	2018-01-25	
Mischa Knabe...	5a9d0d7	Removed comparison code at the end of cosmo.c (was lab...	2018-01-25	
Mischa Knabe...	cf888e0	Renamed MyD_RK4 to csmComoveGrowth (NOTICE: this f...	2018-01-25	
Mischa Knabe...	e8ad609	Changed types for D1, D2, f1 and f2 in MyD_RK4 from float...	2018-01-25	
Mischa Knabe...	c35b92b	Redefined RK_f1, RK_f2, RK_g1 and RK_g2 as static double ...	2018-01-25	
Mischa Knabe...	437b0c8	Removed preprocessor directives that are not used anymore	2018-01-25	
Mischa Knabe...	de36a1f M	Merge branch 'master' of bitbucket.org:mknab/pkdgrav3	2018-01-25	
Mischa Knabe...	e085d30	EUCLID_cosmo.c removed	2018-01-25	
Mischa Knabe...	d9bfdfe M	Merged dpotter/pkdgrav3 into master	2018-01-25	
Joachim Stadel	e7a83da	Minor changes to the healpix function calls to support 64b...	2018-01-25	✓
Joachim Stadel	9cf2fac	Removed the calculation of FoF within the msrOutput call ...	2018-01-25	
Joachim Stadel	ce4cb8e	Chnage to allow FoF group finding in "analysis mode" nSt...	2018-01-25	
Joachim Stadel	6d220b4	Some quick, but not very thorough, changes to allow the c...	2018-01-25	
Mischa Knabe...	c66f2f1	2LPT code added to cosmo.c, cosmo.h and ic.hxx	2018-01-25	
Mischa Knabe...	c778cf4 M	Merged dpotter/pkdgrav3 into master	2018-01-25	
Joachim Stadel	7c5bf4e	Fixed bug in the calculation of the total/half mass radius af...	2018-01-24	
Mischa Knabe...	9b41e70	Commented out some print statements.	2018-01-09	

GIT Example

```
dpotter@daint103:~> cat .gitconfig          (Set this first!)
[user]
    email = douglas.potter@uzh.ch
    name = Doug Potter
dpotter@daint103:~> cd cpi
dpotter@daint103:~/cpi> ls
cpi_mpi.c  cpi_openmp.c  Makefile
dpotter@daint103:~/cpi> git init           (Done only once)
Initialized empty Git repository in /users/dpotter/cpi/.git/
dpotter@daint103:~/cpi> git status        (Output abbreviated)
Untracked files:
    Makefile
    cpi_mpi.c
    cpi_openmp.c
dpotter@daint103:~/cpi> git add Makefile cpi_mpi.c cpi_openmp.c
dpotter@daint103:~/cpi> git commit -m "Initial import"
[master (root-commit) b0f821d] Initial import
  3 files changed, 145 insertions(+)
  create mode 100644 Makefile
  create mode 100644 cpi_mpi.c
  create mode 100644 cpi_openmp.c
```



Compilers | c & Fortran

High Performance Computing



“Hello World” in C

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("Hello world\n");
    return 0;
}
```



“Hello World” in “Machine Code”

```
.data
S0:     .text "Hello world\n"

.code
main:    mova S0,r0
          call printf
          movi #0,r0
          ret
```

```
int main(int argc,char *argv[]) {
    printf("Hello world\n");
    return 0;
}
```

Real Assembler Output

```
.section      __TEXT,__text,regular,pure_instructions
.macosx_version_min 10, 13
.globl        _main
.p2align     4, 0x90
_main:          ## @main
.cfi_startproc
## BB#0:
pushq        %rbp
Lcfi0:         .cfi_def_cfa_offset 16
Lcfi1:         .cfi_offset %rbp, -16
                movq        %rsp, %rbp
Lcfi2:         .cfi_def_cfa_register %rbp
                leaq        L_str(%rip), %rdi
                callq        _puts
                xorl        %eax, %eax
                popq        %rbp
                retq
                .cfi_endproc

L_str:          .section      __TEXT,__cstring,cstring_literals
                ## @str
                .asciz      "Hello world"
.subsections_via_symbols
```

```
int main(int argc, char *argv[])
{
    printf("Hello world\n");
    return 0;
}
```



“man puts”

PUTS(3) BSD Library Functions Manual

PUTS(3)

NAME

puts -- output a line to a stream

LIBRARY

Standard C Library (libc, -lc)

SYNOPSIS

```
#include <stdio.h>
int      puts(const char *s);
```

DESCRIPTION

The function puts() writes the string s, and a terminating newline character to the stream stdout.

Hexadecimal

5F

Binary

0101 | 1111

128
64
32
16
8
4
2
1
01011111

Decimal

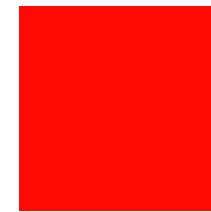
95

Binary & Hex



Boolean Logic

```
leaq L_str(%rip), %rdi
callq _puts
xorl %eax, %eax
popq %rbp
retq
```



A	0	1	0	1
B	0	0	1	1
	↓	↓	↓	↓
false	0	0	0	0
$A \wedge B \Leftrightarrow \bar{A} \leftrightarrow B \Leftrightarrow A \rightarrow \bar{B} \Leftrightarrow \bar{A} \downarrow \bar{B}$	0	0	0	1
$A \leftrightarrow B \Leftrightarrow \bar{A} \wedge B \Leftrightarrow A \downarrow \bar{B} \Leftrightarrow \bar{A} \rightarrow \bar{B}$	0	0	1	0
B	0	0	1	1
$A \rightarrow B \Leftrightarrow \bar{A} \downarrow B \Leftrightarrow A \wedge \bar{B} \Leftrightarrow \bar{A} \leftrightarrow \bar{B}$	0	1	0	0
A	0	1	0	1
$A \oplus B \Leftrightarrow \bar{A} \leftrightarrow B \Leftrightarrow A \leftrightarrow \bar{B} \Leftrightarrow \bar{A} \oplus \bar{B}$	0	1	1	0
$A \vee B \Leftrightarrow \bar{A} \rightarrow B \Leftrightarrow A \leftarrow \bar{B} \Leftrightarrow \bar{A} \uparrow \bar{B}$	0	1	1	1
$A \downarrow B \Leftrightarrow \bar{A} \leftrightarrow B \Leftrightarrow A \leftrightarrow \bar{B} \Leftrightarrow \bar{A} \wedge \bar{B}$	1	0	0	0
$A \leftrightarrow B \Leftrightarrow \bar{A} \oplus B \Leftrightarrow A \oplus \bar{B} \Leftrightarrow \bar{A} \leftrightarrow \bar{B}$	1	0	0	1
\bar{A}	1	0	1	0
$A \rightarrow B \Leftrightarrow \bar{A} \vee B \Leftrightarrow A \uparrow \bar{B} \Leftrightarrow \bar{A} \leftarrow \bar{B}$	1	0	1	1
\bar{B}	1	1	0	0
$A \leftarrow B \Leftrightarrow \bar{A} \uparrow B \Leftrightarrow A \vee \bar{B} \Leftrightarrow \bar{A} \rightarrow \bar{B}$	1	1	0	1
$A \uparrow B \Leftrightarrow \bar{A} \leftarrow B \Leftrightarrow A \rightarrow \bar{B} \Leftrightarrow \bar{A} \vee \bar{B}$	1	1	1	0
true	1	1	1	1



Simple Arithmetic Function

```
float f(float a, float b,
        float c, float d,
        float e) {
    return (a-b)*c + d/e;
}
```

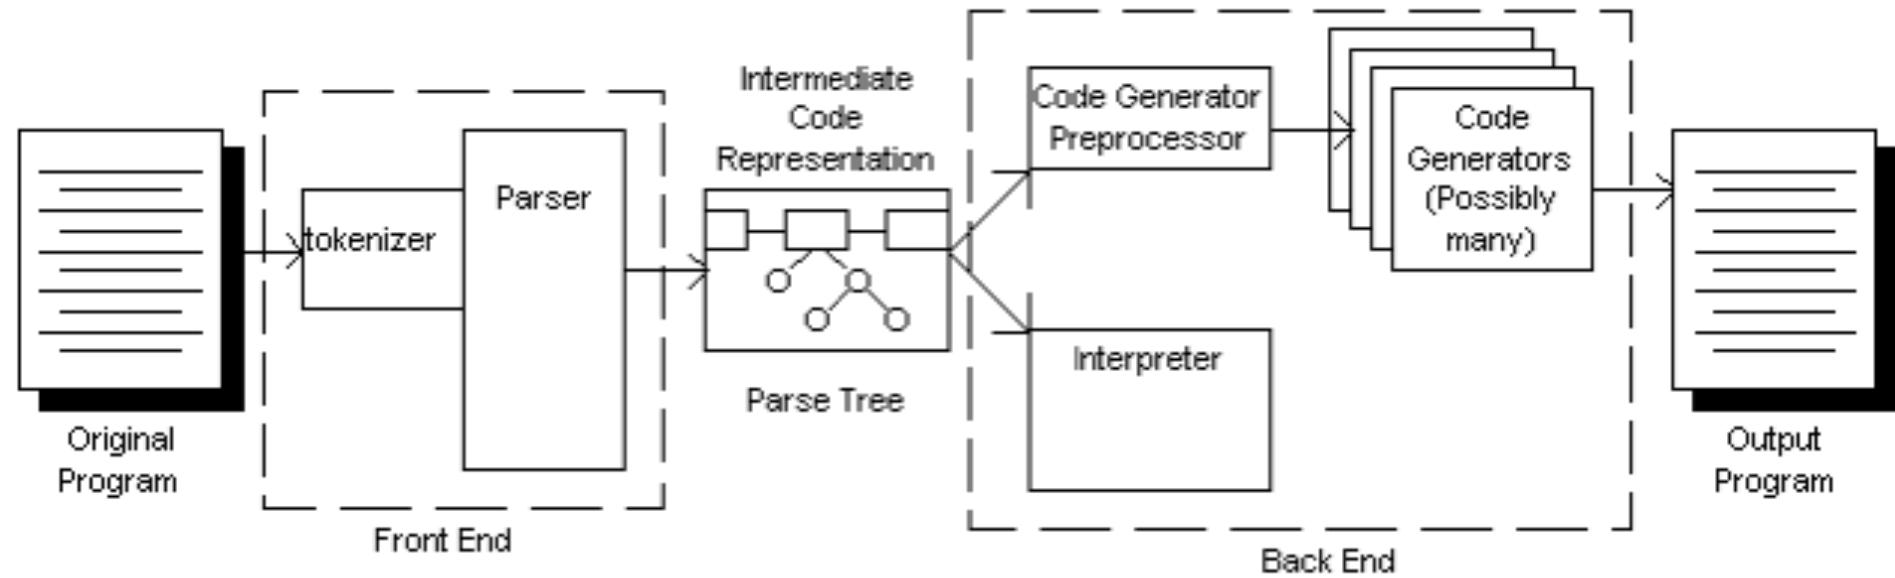
Disassembly of section .text:

```
0: 55          push   %rbp
1: 48 89 e5    mov    %rsp,%rbp
4: f3 0f 5c c1 subss  %xmm1,%xmm0
8: f3 0f 59 c2 mulss  %xmm2,%xmm0
c: f3 0f 5e dc divss  %xmm4,%xmm3
10: f3 0f 58 c3 addss  %xmm3,%xmm0
14: 5d          pop    %rbp
15: c3          retq
```

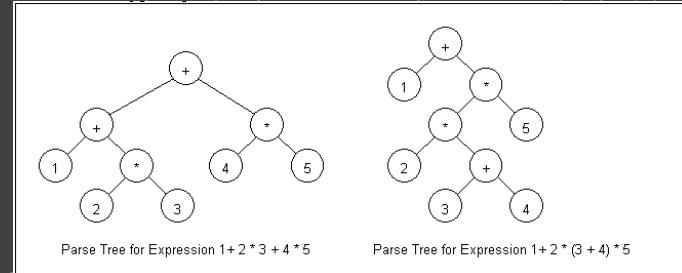
_f:

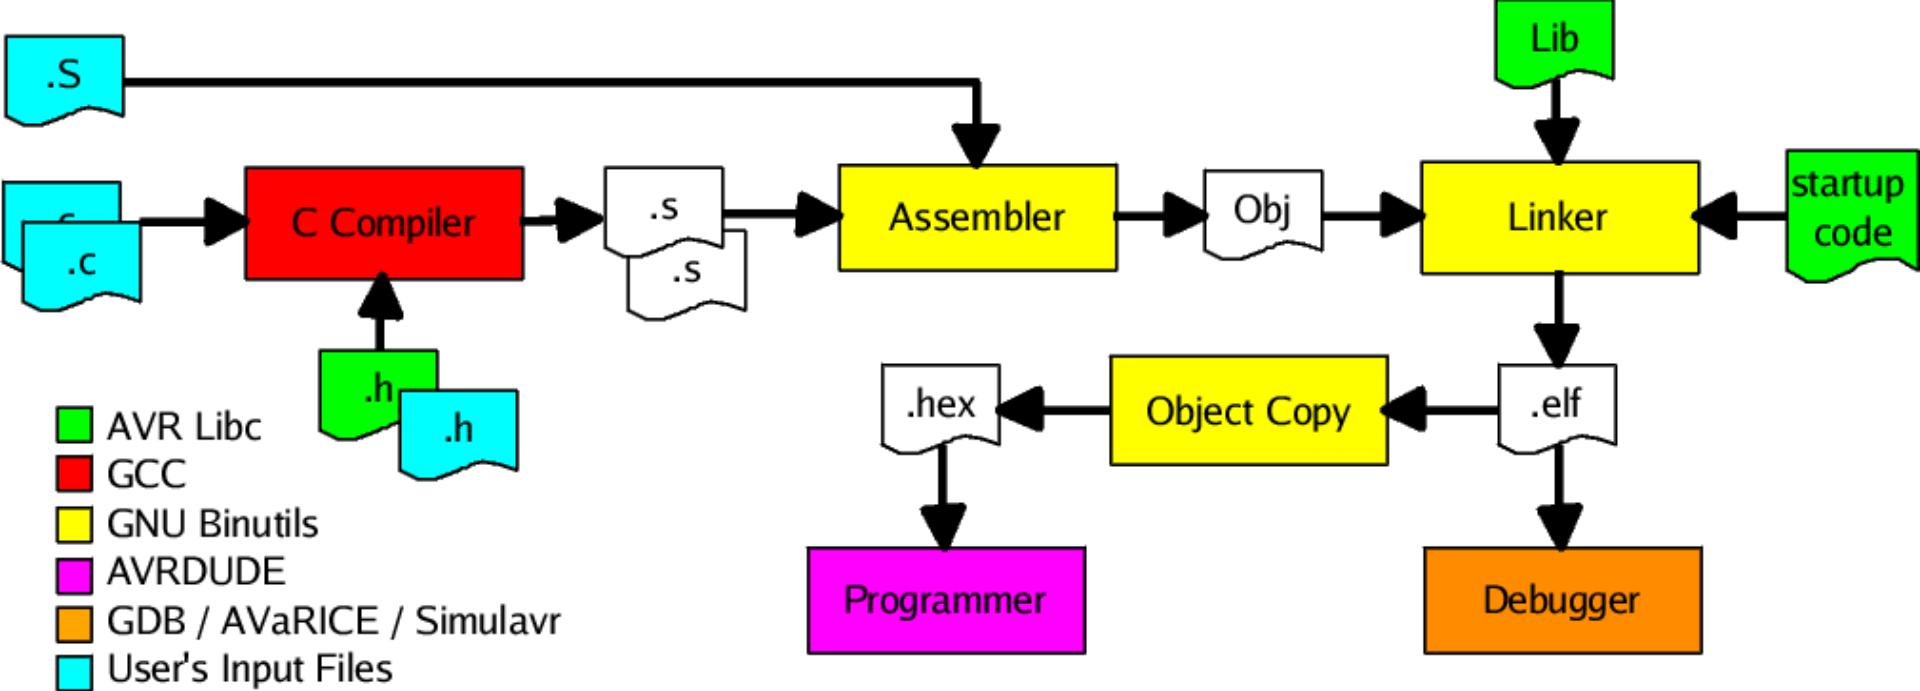
```
pushq  %rbp
movq   %rsp, %rbp
subss  %xmm1, %xmm0
mulss  %xmm2, %xmm0
divss  %xmm4, %xmm3
addss  %xmm3, %xmm0
popq   %rbp
retq
```

gcc -O3 -S -c -o test.s test.c



Anatomy of a Compiler





Compilation Process

High Performance Computing

“Modules” on Piz Daint

```
dpotter@ela3:~> ssh daint
Last login: Thu Feb 22 08:21:39 2018 from 148.187.1.7
=====
IMPORTANT REMINDER FOR USERS of CSCS facilities

help@cscs.ch - +41 91 610 82 10 - http://user.cscs.ch
=====

Please load 'daint-gpu' module for using the GPU/Haswell nodes
or
load 'daint-mc' module for the Multicore/Broadwell nodes

For more info, please refer to the User Portal:
http://user.cscs.ch/getting started/running jobs/piz daint

dpotter@daint105:~> module load daint-mc
dpotter@daint105:~> module swap PrgEnv-cray PrgEnv-gnu
```

Compiler Suites

[Cray Compiler](#)
([PrgEnv-cray](#))

Available on Cray Computers

[GNU Compiler](#)
([PrgEnv-gnu](#))

Free for everyone

[Intel Compiler](#)
([PrgEnv-intel](#))

[Now free for everyone](#)

[Portland](#)
([PrgEnv-pgi](#))

Important for OpenACC (GPU)

[Clang](#)
(Apple & Cray)

Open Source & GNU Compatible

[Visual Studio](#)
(Windows)

Not usually High Performance



List Loaded Modules

```
dpotter@daint105:~> module list
Currently Loaded Modulefiles:
 1) modules/3.2.10.6
 2) eproxy/2.0.16-6.0.4.1_3.1__g001b199.ari
 3) gcc/5.3.0
 4) craype-broadwell
 5) craype-network-aries
 6) craype/2.5.12
 7) cray-mpich/7.6.0
 8) slurm/17.02.9+git20180119.b04278-1
 9) xalt/daint-2016.11
10) daint+mc
11) cray-libscl/17.06.1
12) udreg/2.3.2-6.0.4.0_12.2__g2f9c3ee.ari
13) ugni/6.0.14-6.0.4.0_14.1__ge7db4a2.ari
14) pmi/5.0.12
15) dmapp/7.1.1-6.0.4.0_46.2__gb8abda2.ari
16) gni-headers/5.0.11-6.0.4.0_7.2__g7136988.ari
17) xpmem/2.2.2-6.0.4.0_3.1__g43b0535.ari
18) job/2.2.2-6.0.4.0_8.2__g3c644b5.ari
19) dvs/2.7.2.2.32-6.0.4.1_7.1__ged1923a
20) alps/6.4.1-6.0.4.0_7.2__g86d0f3d.ari
21) rca/2.2.11-6.0.4.0_13.2__g84de67a.ari
22) atm/2.1.1
23) perftools-base/6.5.1
24) PrgEnv-gnu/6.0.4
dpotter@daint105:~> module avail
----- /opt/cray/pe/perftools/6.5.1/modulefiles -----
perftools          perftools-lite-events perftools-lite-hbm    perftools-nwpc
perftools-lite     perftools-lite-gpu      perftools-lite-loops
----- /apps/daint/UES/jenkins/6.0.UP04/mc/easybuild/modules/all -----
Amber/16-CrayGNU-17.08-parallel(default)          PLUMED/2.3.0-CrayGNU-17.08
Amber/16-CrayGNU-17.08-serial                   Paraver/4.6.4.rc1(default)
Boost/1.65.0-CrayGNU-17.08                      Perl/5.26.1-bare(default)
Boost/1.65.0-CrayGNU-17.08-python2            PyExtensions/2.7-CrayGNU-17.08
Boost/1.65.0-CrayGNU-17.08-python3(default)       PyExtensions/3.5-CrayGNU-17.08(default)
... and many more
dpotter@daint105:~> module avail gcc
----- /opt/modulefiles -----
gcc/4.9.3           gcc/6.1.0             gcc/7.1.0
gcc/5.3.0(default)        gcc/6.2.0
```



List Loaded Modules (Eiger)

```
[eiger] [dpotter@nid001765 ~]$ module list
Currently Loaded Modules:
 1) craype-x86-rome      4) perftools-base/21.12.0      7) craype/2.7.13      10) cray-libsci/21.08.1.2
 2) libfabric/1.15.0.0    5) xpmem/2.4.4-2.3_9.1_gff0e1d9.shasta  8) cray-dsmml/0.2.2     11) PrgEnv-cray/8.3.0
 3) craype-network-ofi   6) cce/13.0.0                      9) cray-mpich/8.1.12
```

```
[eiger] [dpotter@nid001080 ~]$ cc --version
Cray clang version 13.0.0 (24b043d62639ddb4320c86db0b131600fdbc6ec6)
Target: x86_64-unknown-linux-gnu
Thread model: posix
InstalledDir: /opt/cray/pe/cce/13.0.0/cce-clang/x86_64/share/../bin
```

```
[eiger] [dpotter@nid001765 ~]$ module swap PrgEnv-cray PrgEnv-gnu
Due to MODULEPATH changes, the following have been reloaded:
 1) cray-mpich/8.1.12
```

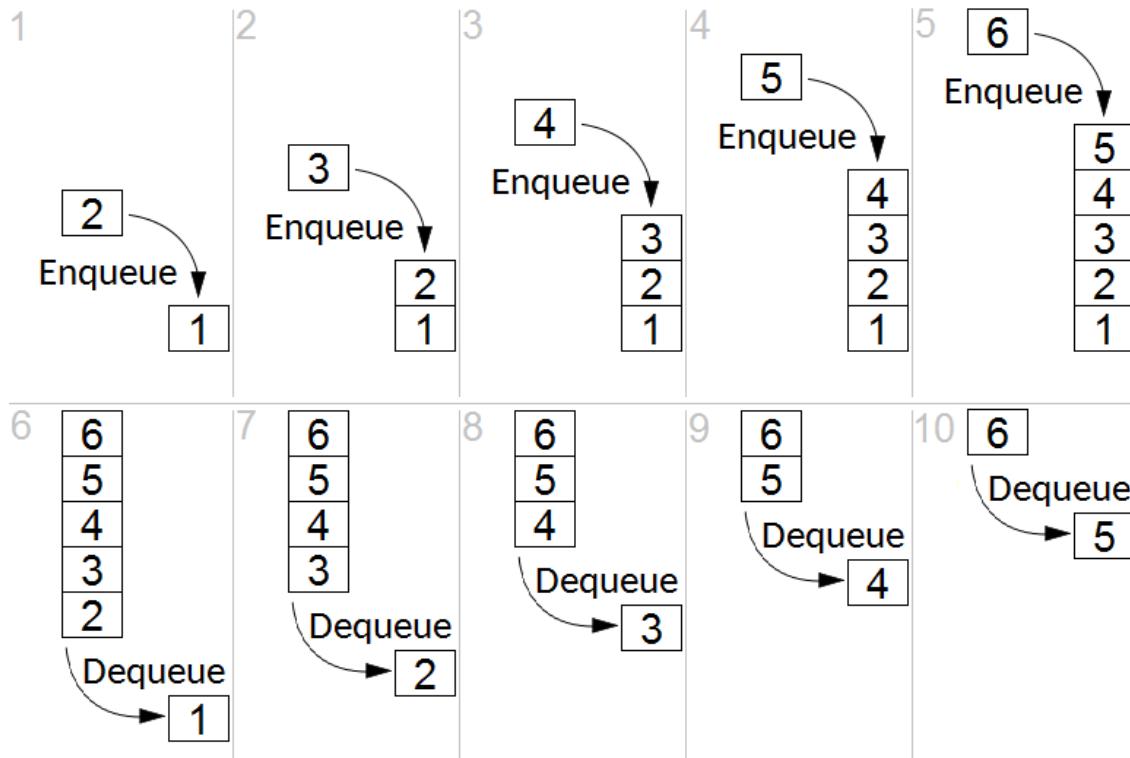
```
[eiger] [dpotter@nid001765 ~]$ module list
Currently Loaded Modules:
 1) craype-x86-rome      4) perftools-base/21.12.0      7) craype/2.7.13      10) cray-libsci/21.08.1.2
 2) libfabric/1.15.0.0    5) xpmem/2.4.4-2.3_9.1_gff0e1d9.shasta  8) cray-dsmml/0.2.2     11) PrgEnv-gnu/8.3.0
 3) craype-network-ofi   6) gcc/11.2.0                      9) cray-mpich/8.1.12
```

```
[eiger] [dpotter@nid001080 ~]$ cc --version
gcc (GCC) 11.2.0 20210728 (Cray Inc.)
Copyright (C) 2021 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

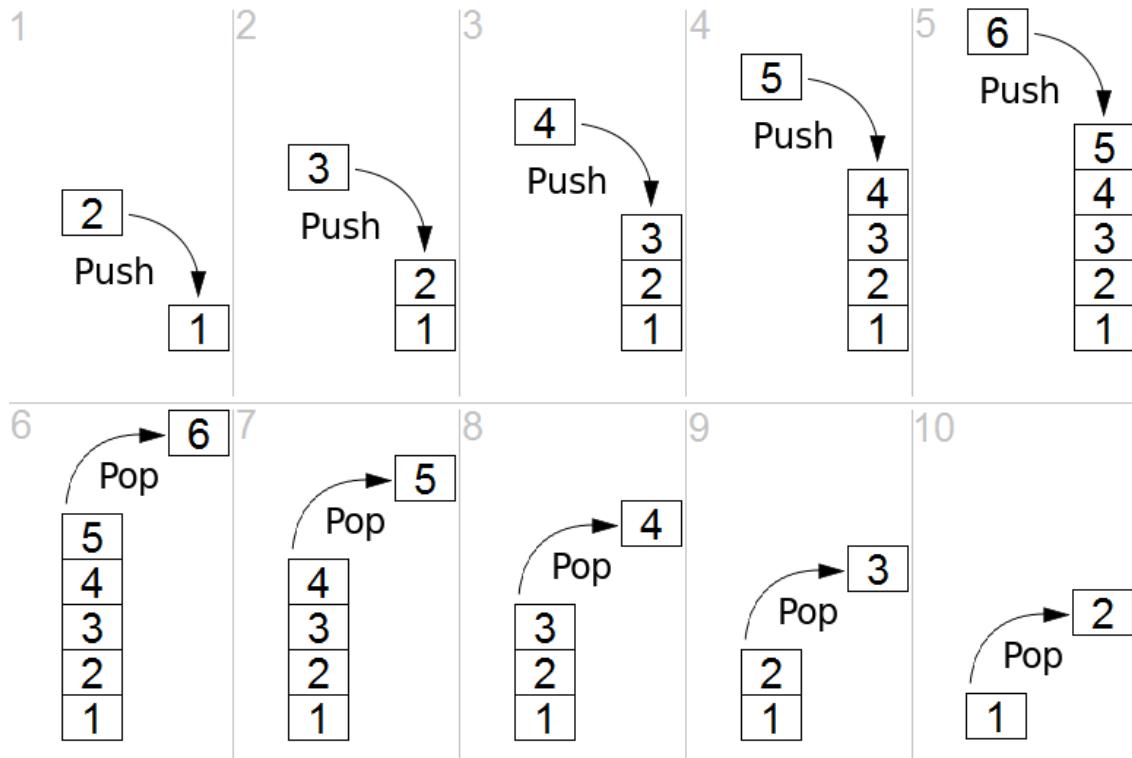
Batch Queue Systems



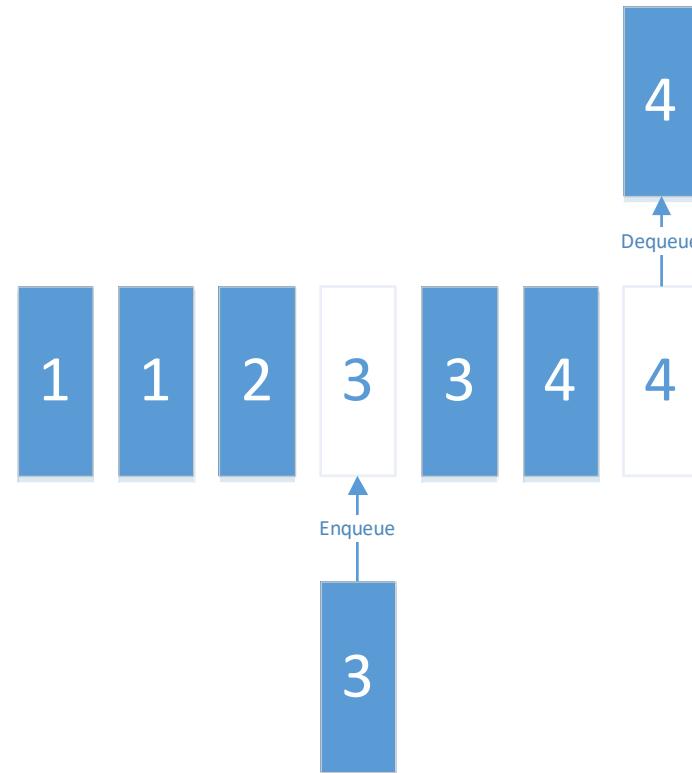
“Queue”
First In,
First Out
(FIFO)



“Stack”
Last In,
First Out
(LIFO)



Priority
Queue



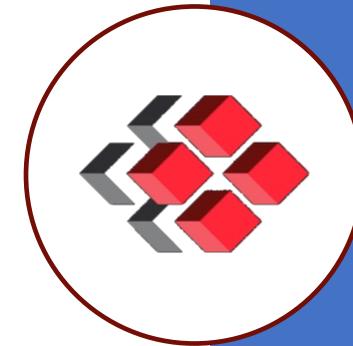


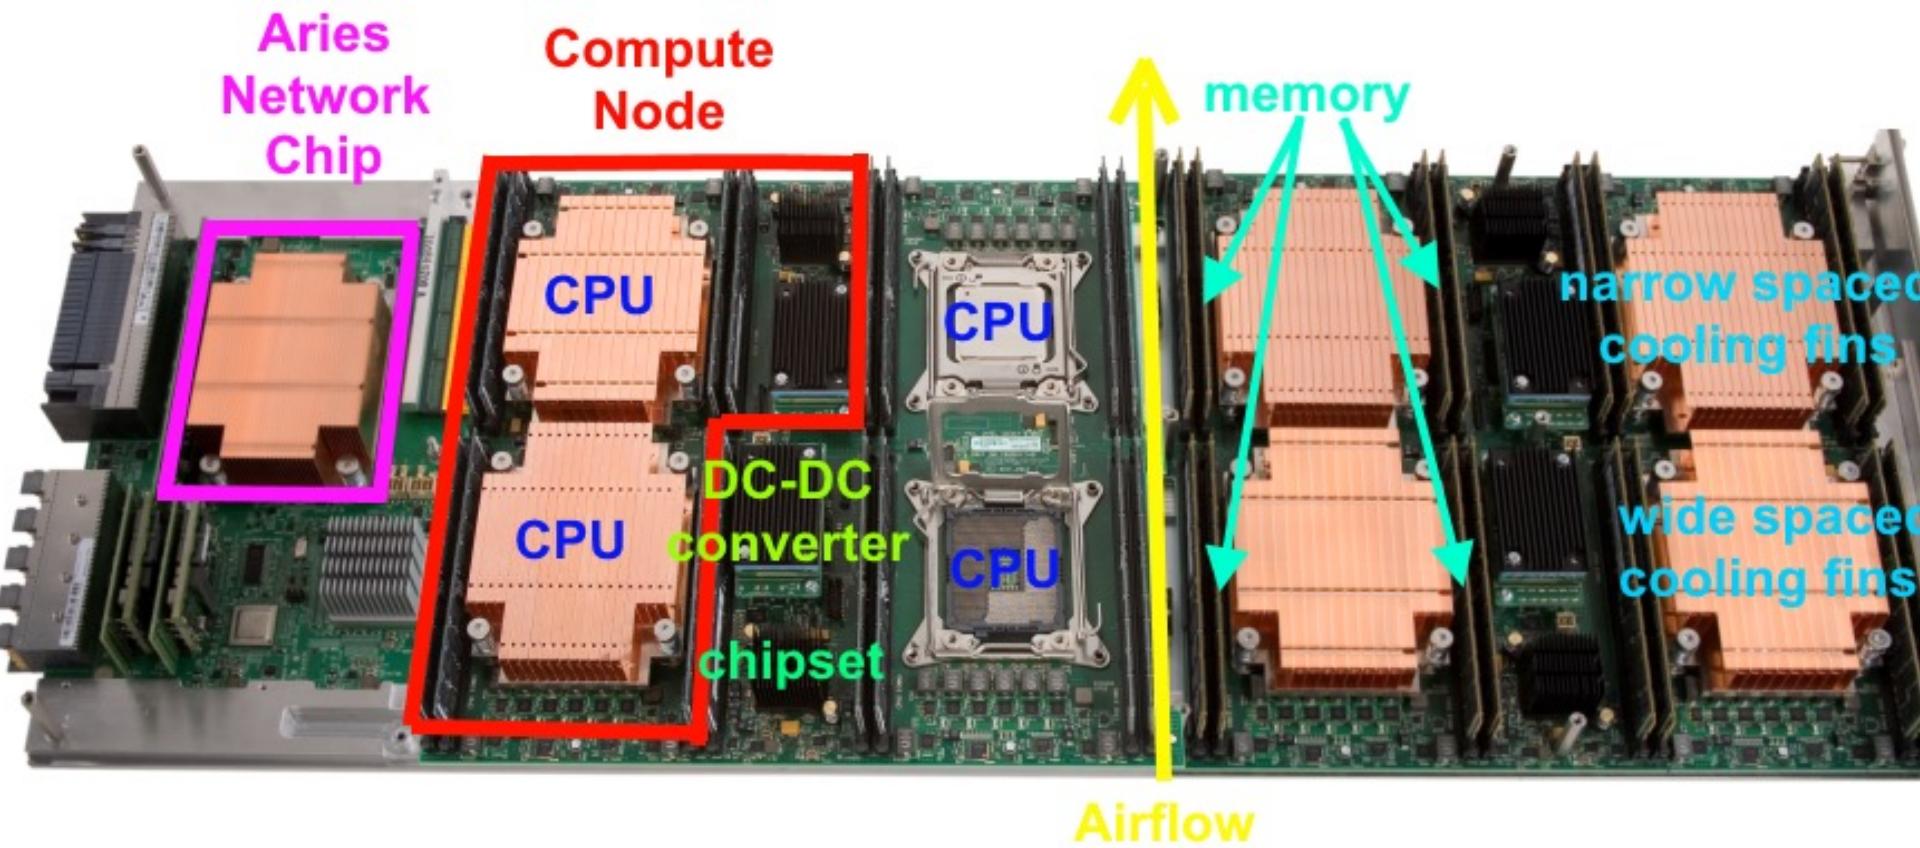
Generate Job Script

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=hpc_test
#SBATCH --time=01:00:00
#SBATCH --nodes=1
#SBATCH --ntasks-per-core=1
#SBATCH --ntasks-per-node=36
#SBATCH --cpus-per-task=1
#SBATCH --partition=normal
#SBATCH --constraint=mc

export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

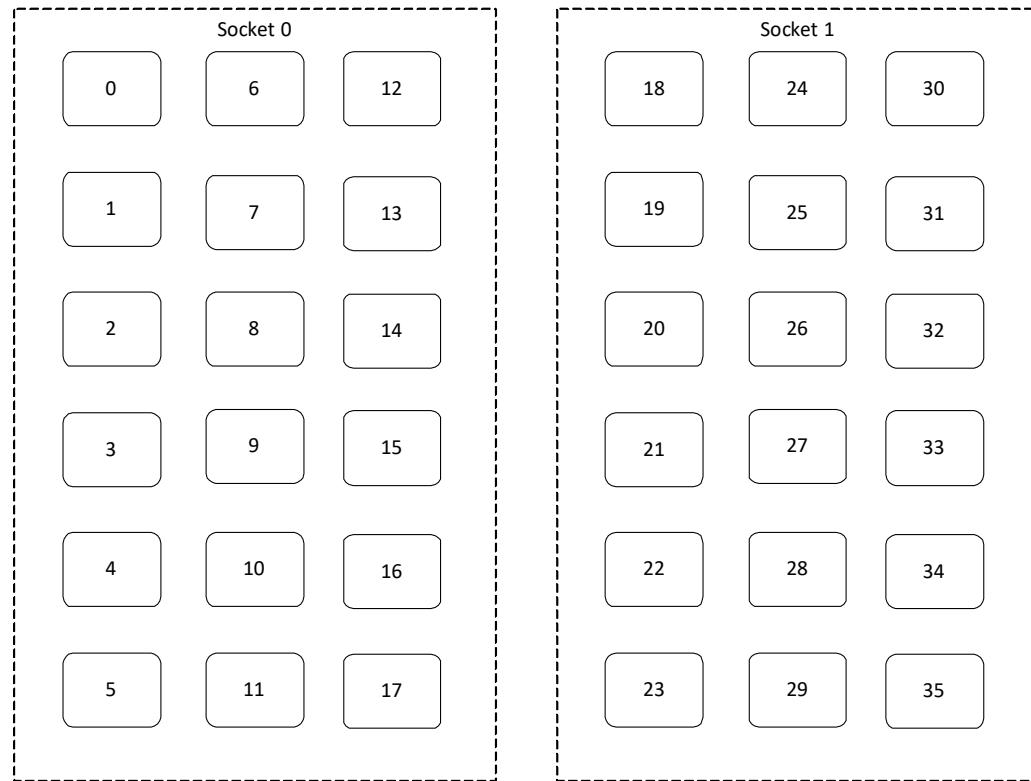
srun hostname
```



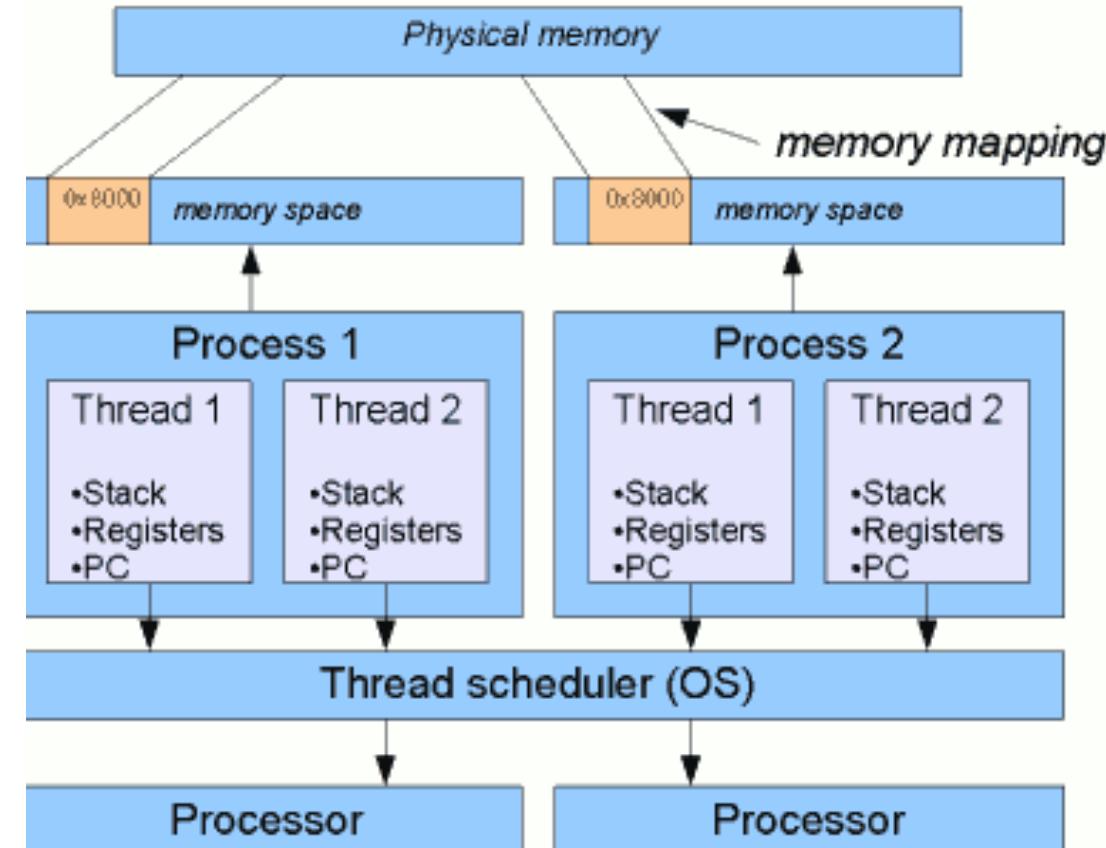


Cray XC40 Blade (daint-mc)

Processor Cores Layout



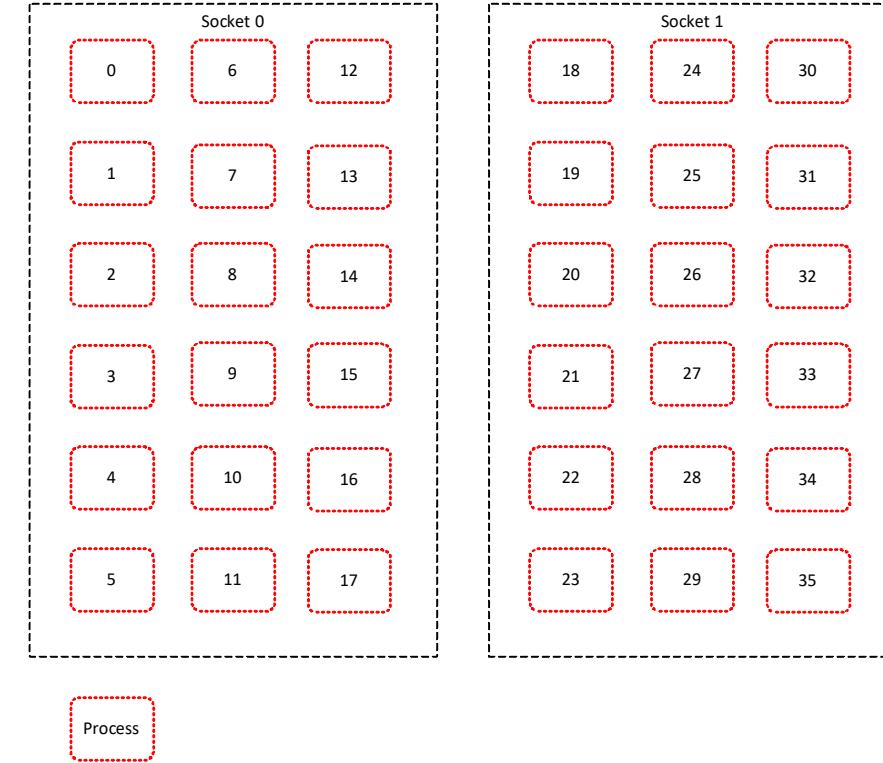
Processes & Threads





MPI Layout

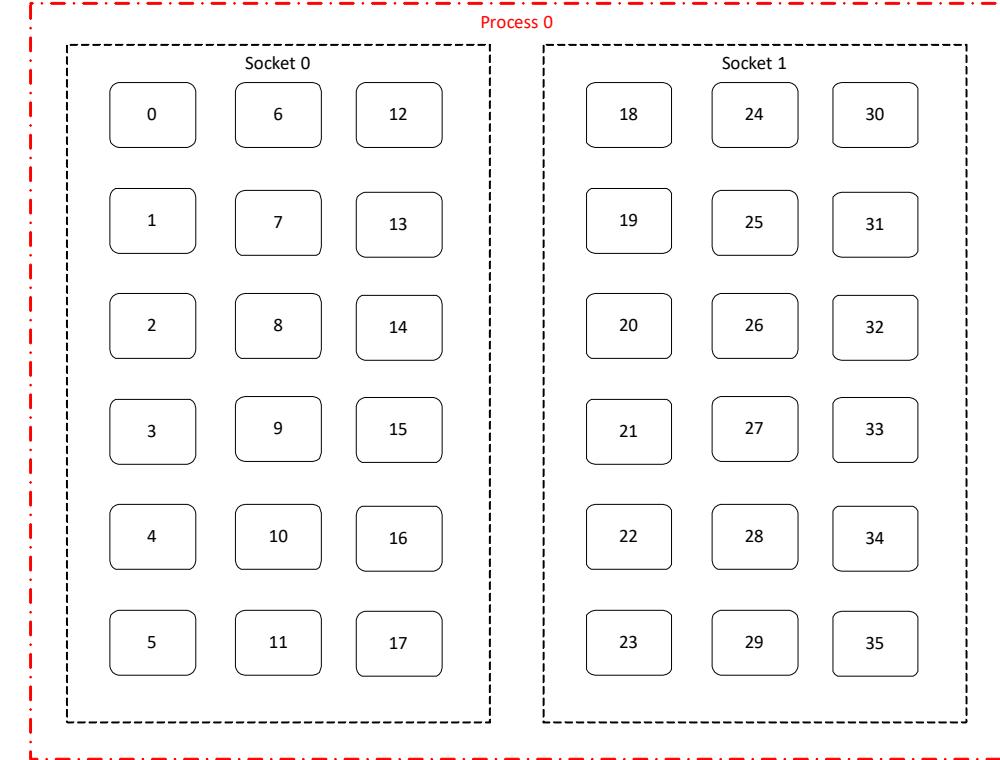
- --ntasks-per-core=1
- --ntasks-per-node=36
- --cpus-per-task=1





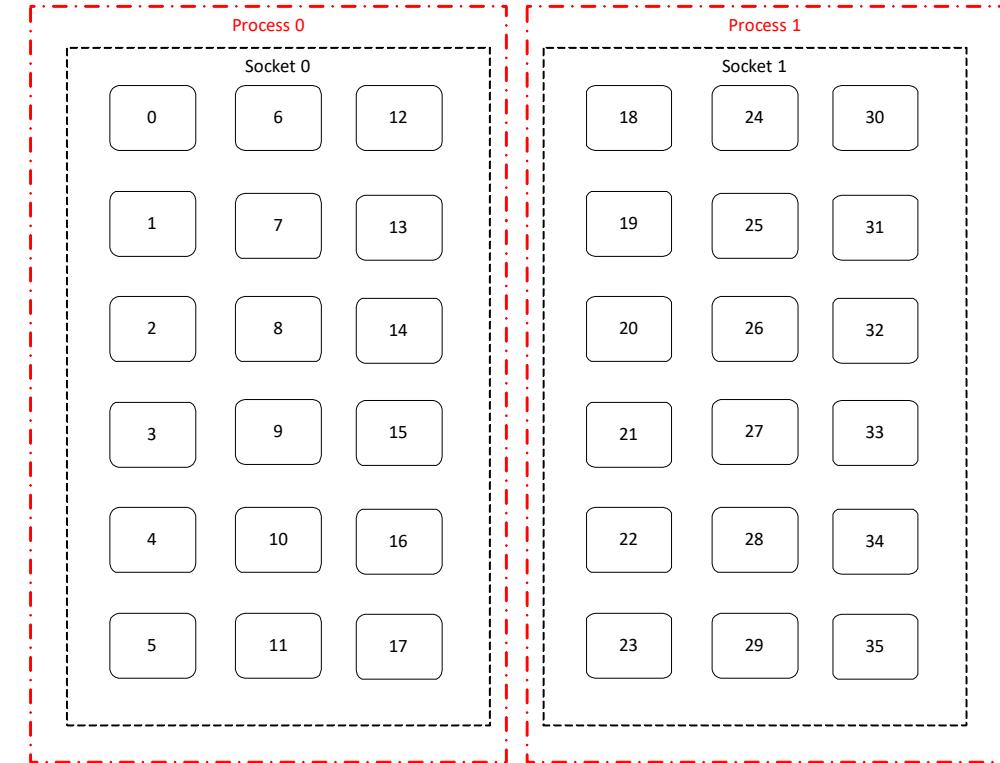
OpenMP Threads Layout

- `--ntasks-per-core=1`
- `--ntasks-per-node=1`
- `--cpus-per-task=36`



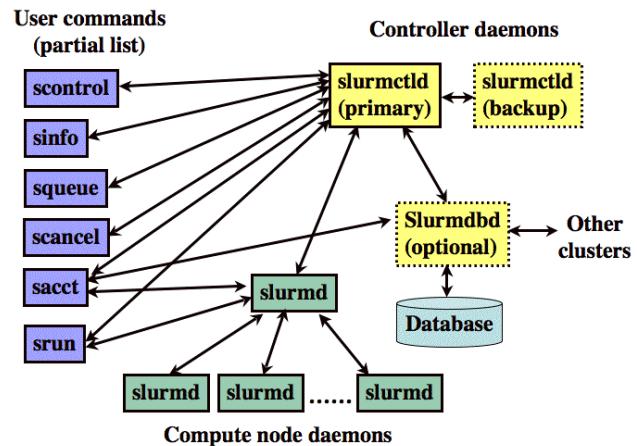
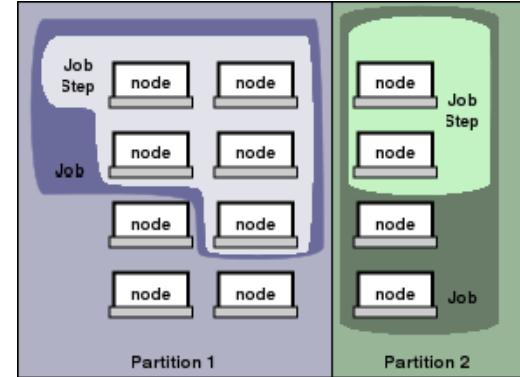
Hybrid Layout

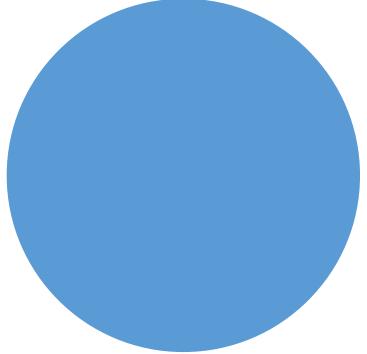
- --ntasks-per-core=1
- --ntasks-per-node=2
- --cpus-per-task=18



SLURM Job Submission

- “sbatch” to submit a job
- “squeue” to see jobs
 - Pending jobs: -t pending
 - Running jobs: -t running
 - Your jobs: -u username
- “sinfo” to see partition information
 - Normal partition: -p normal
 - Debug partition: -p debug
- “srun” to run a job step
 - Always use “srun” in scripts
 - You can use it for interactive sessions
 - Not recommended if avoidable
- “scancel” to cancel a job
- “accounting” for usage
 - CHF 0.376 per “node hour”

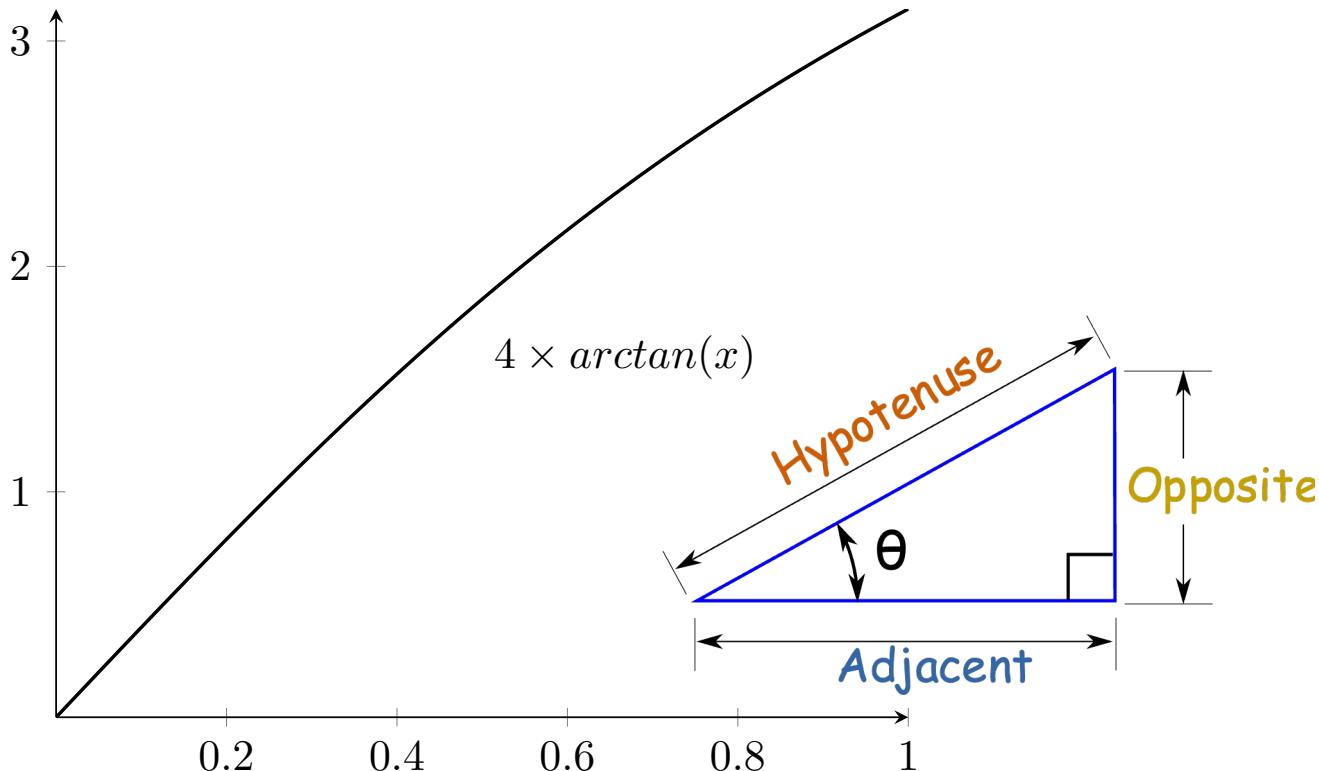




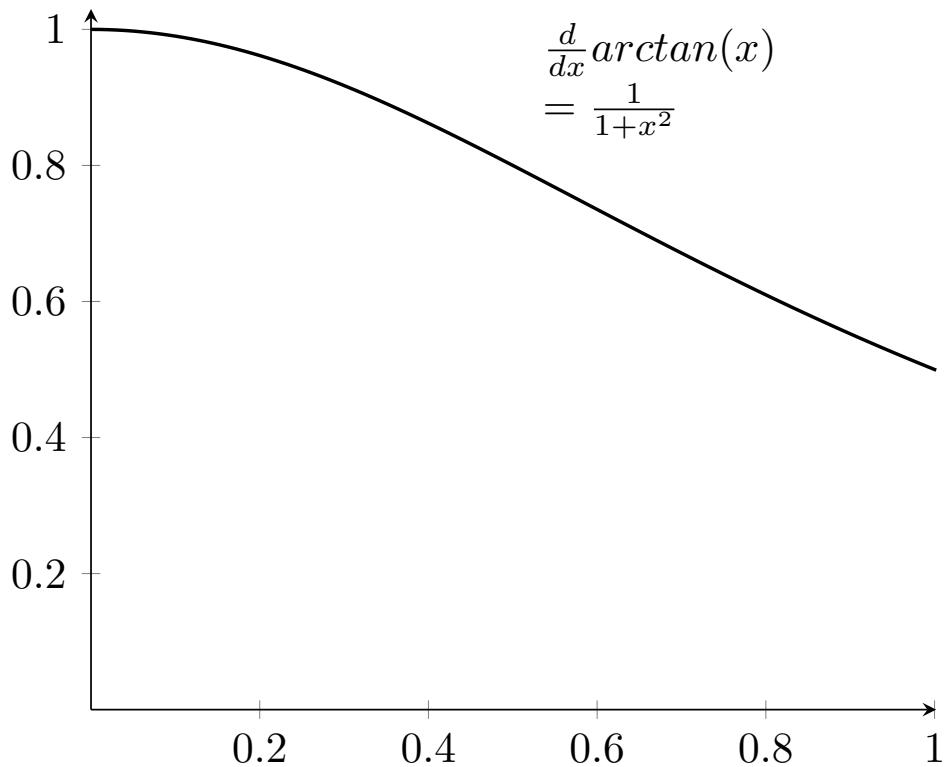
Sample Program

Calculate pi

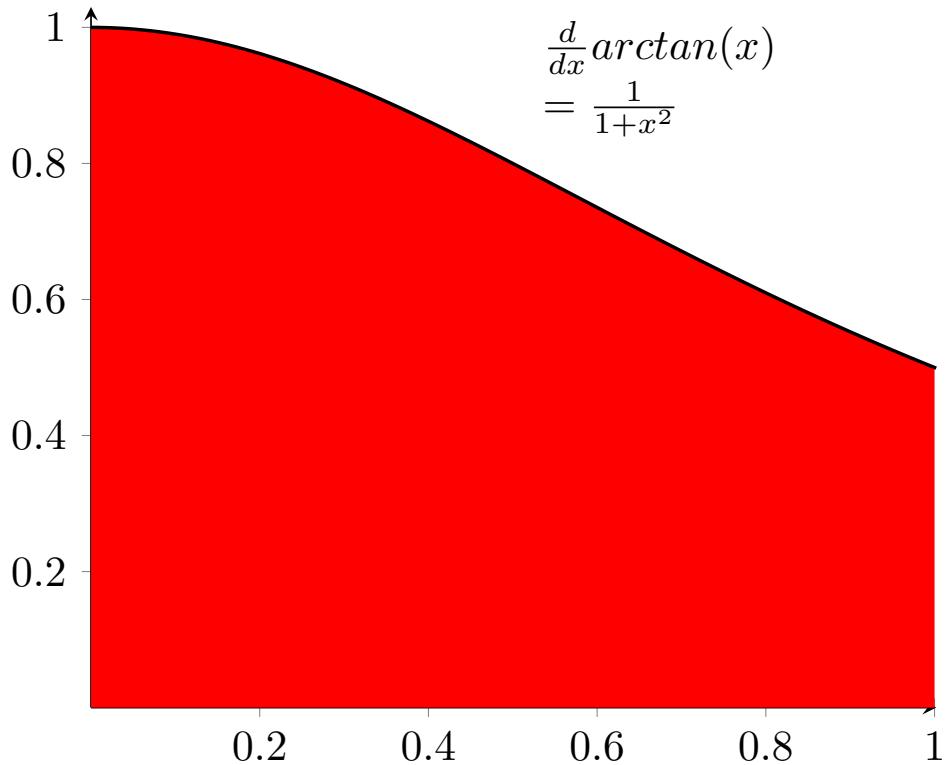
Calculate pi
using arctan



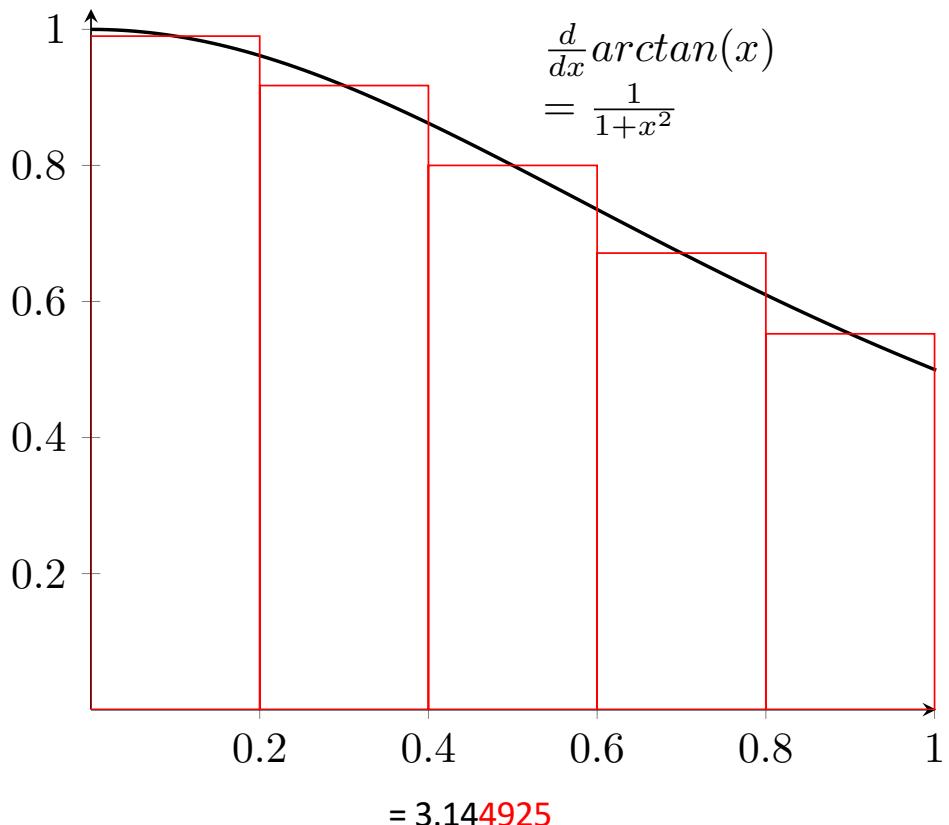
Derivative is
well known



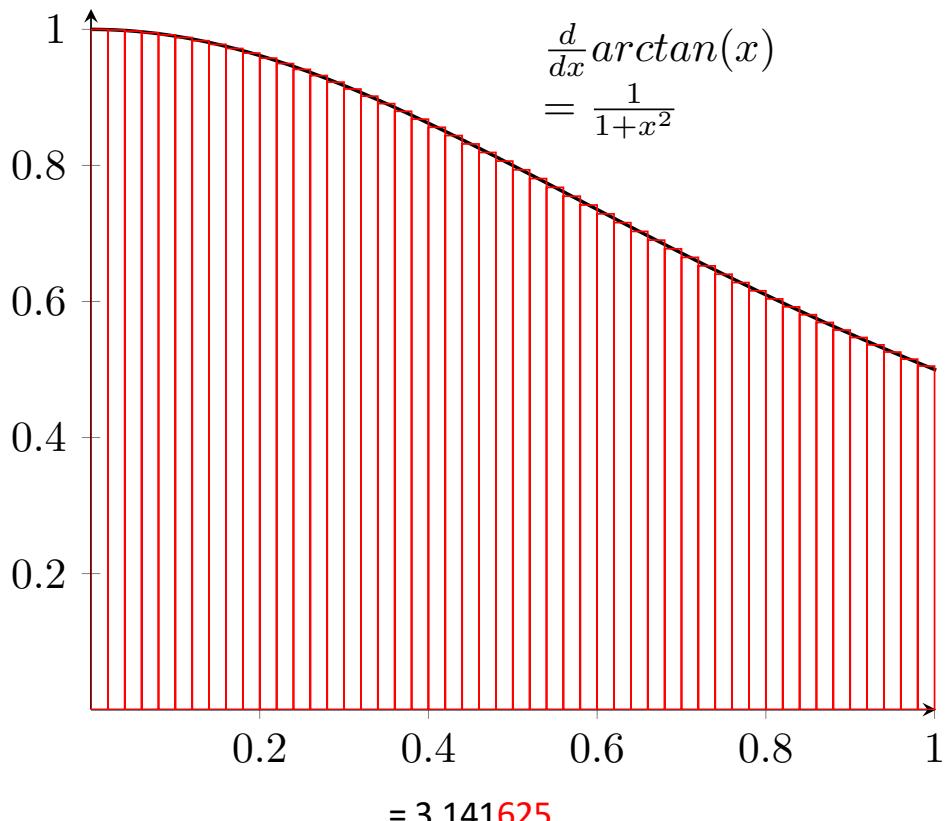
Integrate
to get $\frac{\pi}{4}$



Integrate
numerically

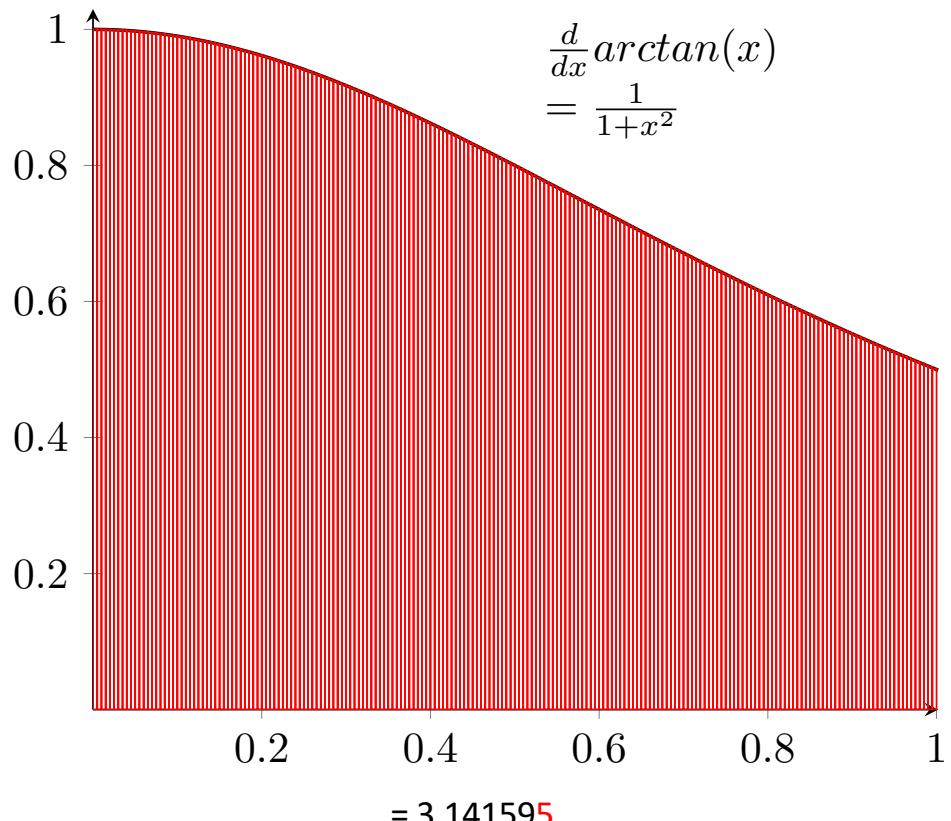


Improve
Accuracy



High Performance Computing

We haven't
even
started!





Integration in Python

```
N=5
dx=1.0/N
s = 0
for x in range(0, N):
    s += dx * 1 / ( 1 + ((x+0.5)*dx)**2 )
print(s*4)
```

```
dhcp-94-191:cpi$ python integrate.py
3.1449258640033277
```



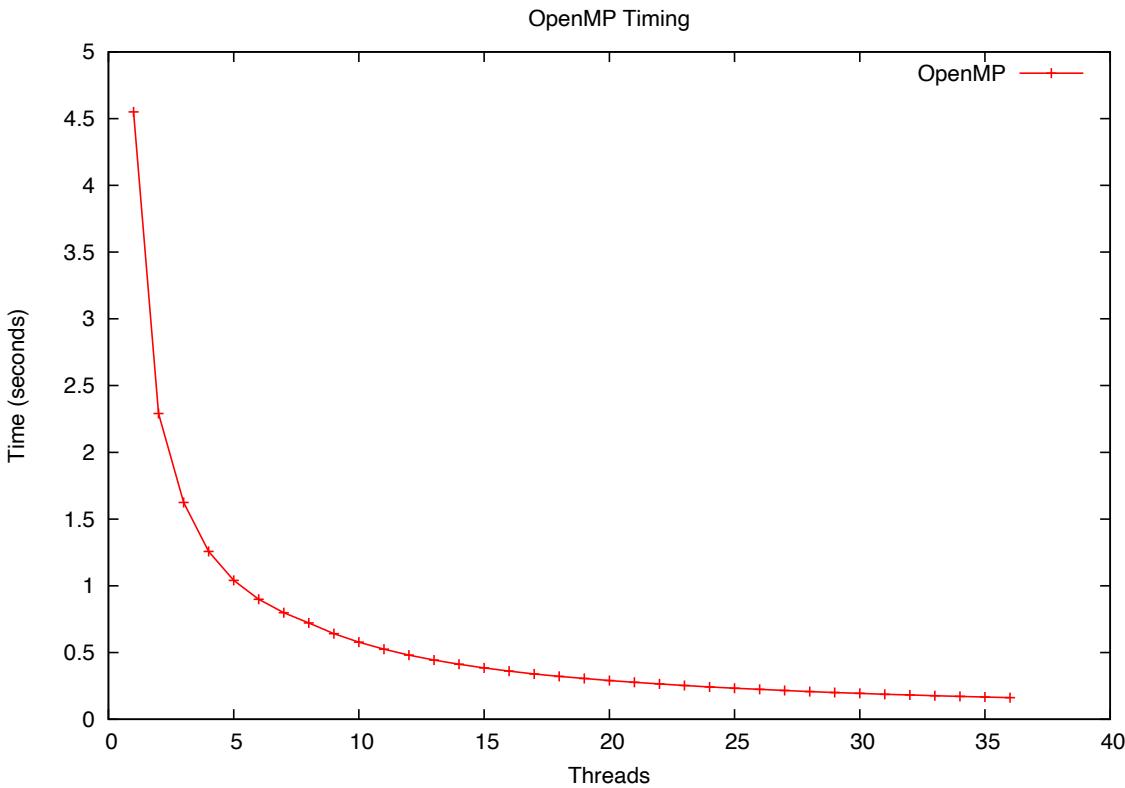
Timing.

Is this
good?

Threads	Seconds	Threads	Seconds
1	4.550	19	0.3053
2	2.291	20	0.2897
3	1.624	21	0.2768
4	1.258	22	0.2644
5	1.041	23	0.2528
6	0.8988	24	0.2423
7	0.7989	25	0.2326
8	0.7217	26	0.2237
9	0.6415	27	0.2154
10	0.5774	28	0.2077
11	0.5249	29	0.2006
12	0.4811	30	0.1939
13	0.4441	31	0.1876
14	0.4124	32	0.1818
15	0.3849	33	0.1763
16	0.3609	34	0.1711
17	0.3397	35	0.1662
18	0.3208	36	0.1616

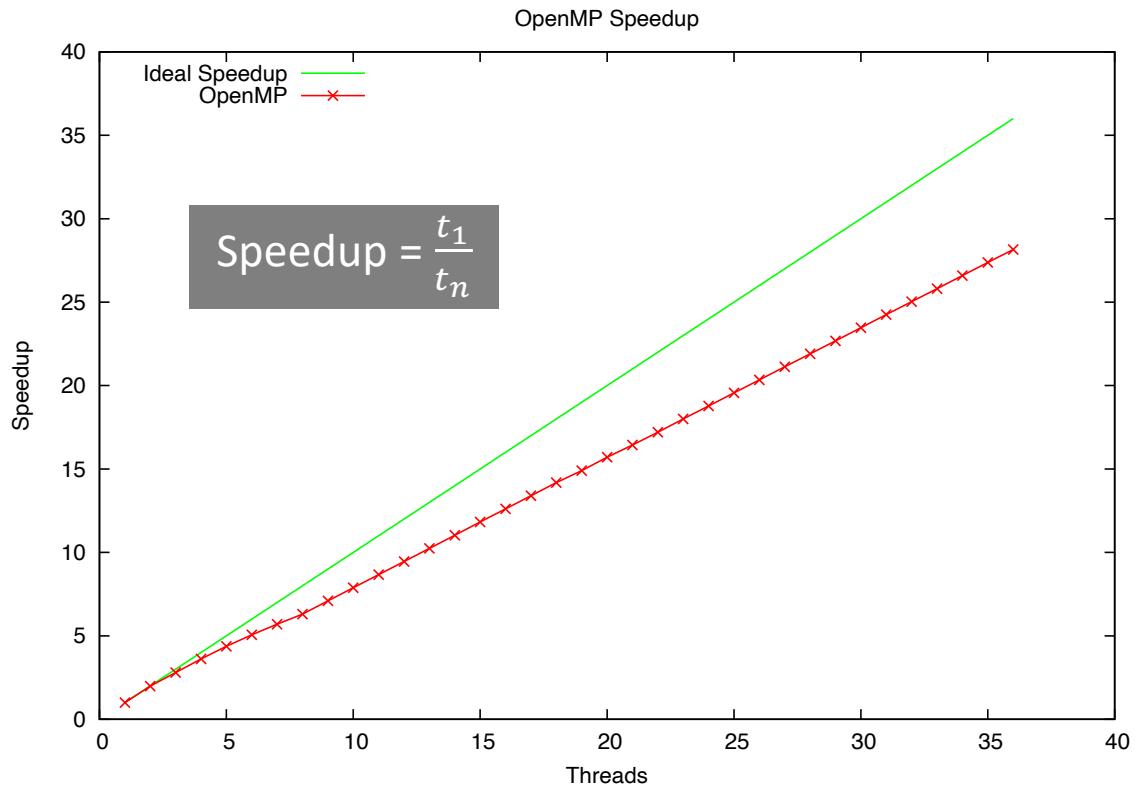
Timing Plot

```
set title 'OpenMP Timing'  
set xlabel 'Threads'  
set ylabel 'Time (seconds)'  
set key top right  
plot "cpi_openmp.dat" u 1:2 w lp  
lw 2 t 'OpenMP'
```

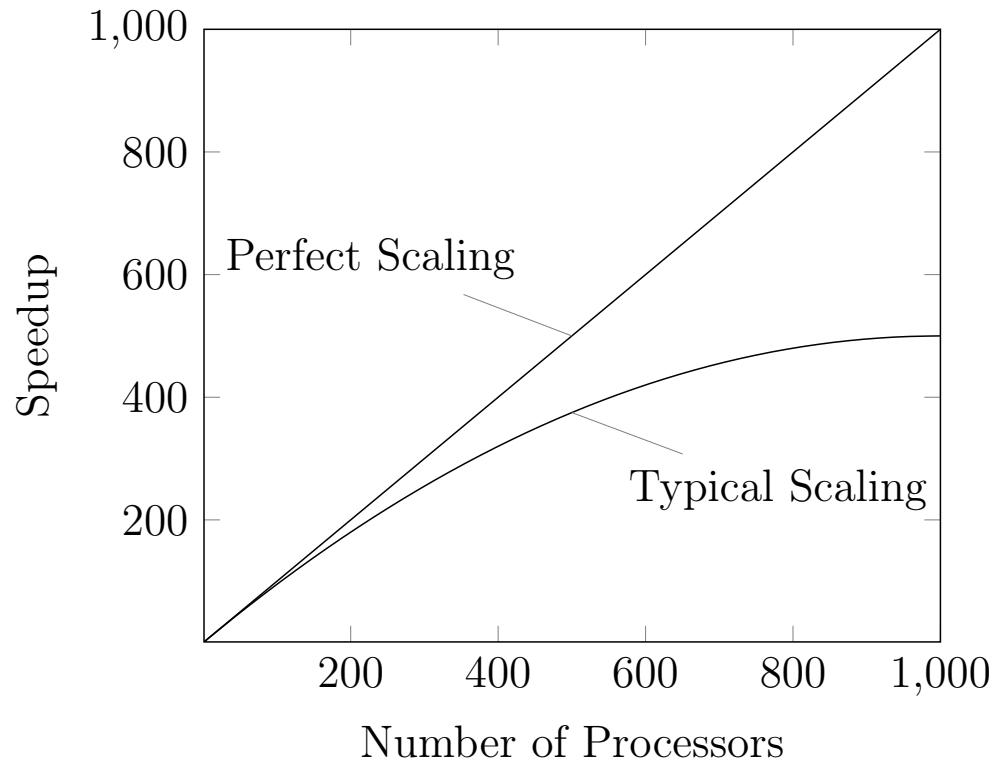


“Speedup” Plot

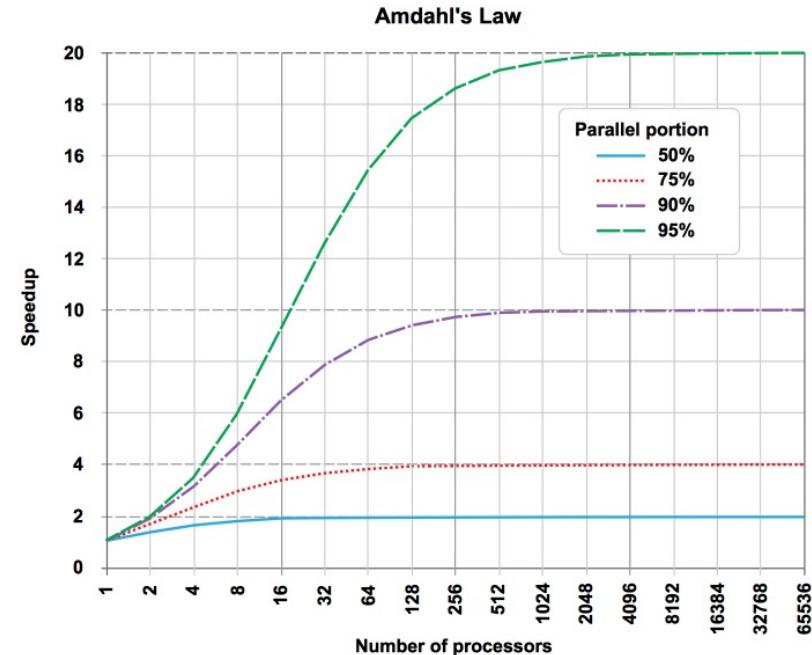
```
set title 'OpenMP Speedup'  
set xlabel 'Threads'  
set ylabel 'Speedup'  
set key top left  
plot x lc 2 lw 2 t 'Ideal Speedup',  
    "cpi_openmp.dat" u 1:(4.55/$2)\  
    w lp lc 1 lw 2 t 'OpenMP'
```



How does it scale?



Why won't it scale?

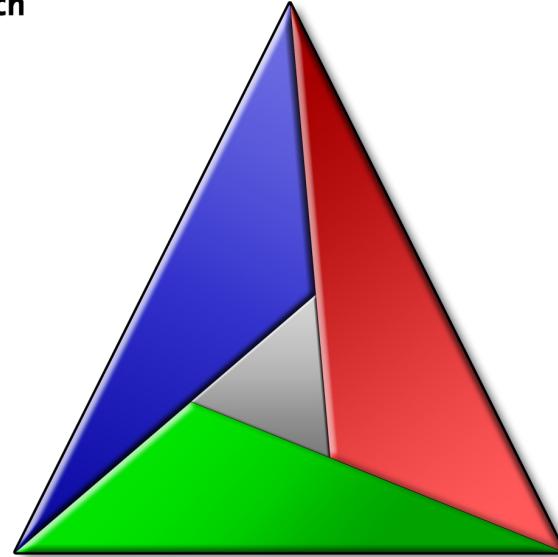




Makefile (for the C versions)

```
Target          Dependencies          Triggered when newer
all:    cpi_openmp cpi_mpi
cpi_openmp:   cpi_openmp.c
            cc -g -O3 -o cpi_openmp -fopenmp cpi_openmp.c
cpi_mpi:     cpi_mpi.c
            cc -g -O3 -o cpi_mpi cpi_mpi.c
Build Rule(s)

clean:
        rm -f cpi_openmp cpi_mpi
Common “clean” target
```



Advanced Make Systems

General Recipe

```
SRC=$HOME/source  
DST=$HOME/install  
  
cd $SRC  
.configure\  
--prefix=$DST  
  
make  
make install
```

```
SRC=$HOME/source  
DST=$HOME/install  
mkdir $HOME/build  
cd $HOME/build  
cmake\  
-DCMAKE_INSTALL_PREFIX:PATH=$DST\  
$SRC  
make  
make install
```