



University of Zurich^{UZH}

High Performance Computing Lecture 3

Douglas Potter

Jozef Bucko

Noah Kubli



Where are
we?

We can run codes

- Submit on a SLURM queue

We can compile codes

- More on this today


Understanding Parallel

- Main topic today

GOAL: change a code

- Maybe write your own

Exercise:
Compute PI
in Parallel





Integration in Python

```
N=5
dx=1.0 / N
s = 0
for x in range(0, N):
    s += dx / ( 1 + ((x+0.5)*dx)**2 )
print(s*4)
```

```
dhcp-94-191:cpu$ python integrate.py
3.144925864003328
```



OpenMP Results (multi-threaded)

N=1,000,000,000

```
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes={nodes}
#SBATCH --ntasks-per-node={processes}
#SBATCH --cpus-per-task={threads}
#SBATCH --partition=normal
#SBATCH --constraint=mc

export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

srun cpi_openmp
```



OpenMP Results (multi-threaded)

N=1,000,000,000

```
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=36
#SBATCH --partition=normal
#SBATCH --constraint=mc

export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

srun cpi_openmp
```



OpenMP Results (multi-threaded)

N=1,000,000,000

running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.1415926535899⁰¹ computed in 2.291 seconds
running on 3 threads: PI = 3.1415926535899⁶² computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589⁸²¹ computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589⁵⁹⁶ computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589⁶⁸² computed in 0.8988 seconds
running on 7 threads: PI = 3.141592653589⁶³ computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589⁷⁶⁹ computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589⁶⁵⁶ computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589⁷⁹⁴ computed in 0.5774 seconds
running on 11 threads: PI = 3.141592653589⁶⁶ computed in 0.5249 seconds
running on 12 threads: PI = 3.141592653589⁸⁶ computed in 0.4811 seconds
running on 13 threads: PI = 3.1415926535898⁶⁵ computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589⁷⁸⁸ computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589⁸⁰⁵ computed in 0.3849 seconds
running on 16 threads: PI = 3.1415926535898³² computed in 0.3609 seconds
running on 17 threads: PI = 3.1415926535898³⁹ computed in 0.3397 seconds
running on 18 threads: PI = 3.1415926535898¹⁴ computed in 0.3208 seconds
running on 19 threads: PI = 3.1415926535898²⁶ computed in 0.3053 seconds
running on 20 threads: PI = 3.1415926535898⁵⁵ computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589⁷⁷⁵ computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589⁸²³ computed in 0.2644 seconds
running on 23 threads: PI = 3.1415926535898⁶⁶ computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589⁷⁹² computed in 0.2423 seconds
running on 25 threads: PI = 3.141592653589⁷⁸ computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589⁸³² computed in 0.2237 seconds
running on 27 threads: PI = 3.1415926535898³⁵ computed in 0.2154 seconds
running on 28 threads: PI = 3.1415926535898¹⁶ computed in 0.2077 seconds
running on 29 threads: PI = 3.1415926535898¹⁹ computed in 0.2006 seconds
running on 30 threads: PI = 3.1415926535898⁹³ computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589⁷⁴⁴ computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589⁷⁵⁸ computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589⁸⁰⁶ computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589⁹²⁶ computed in 0.1711 seconds
running on 35 threads: PI = 3.141592653589⁷⁹ computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589⁸²² computed in 0.1616 seconds



MPI Results

```
This is Process-1/36 running on nid00564
This is Process-3/36 running on nid00564
This is Process-5/36 running on nid00564
This is Process-6/36 running on nid00564
This is Process-8/36 running on nid00564
This is Process-11/36 running on nid00564
This is Process-13/36 running on nid00564
This is Process-14/36 running on nid00564
This is Process-17/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-22/36 running on nid00564
This is Process-23/36 running on nid00564
This is Process-25/36 running on nid00564
This is Process-26/36 running on nid00564
This is Process-28/36 running on nid00564
This is Process-29/36 running on nid00564
This is Process-31/36 running on nid00564
This is Process-33/36 running on nid00564
This is Process-35/36 running on nid00564
This is Process-0/36 running on nid00564
This is Process-2/36 running on nid00564
This is Process-4/36 running on nid00564
This is Process-7/36 running on nid00564
This is Process-9/36 running on nid00564
This is Process-10/36 running on nid00564
This is Process-12/36 running on nid00564
This is Process-15/36 running on nid00564
This is Process-16/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-20/36 running on nid00564
This is Process-21/36 running on nid00564
This is Process-24/36 running on nid00564
This is Process-27/36 running on nid00564
This is Process-30/36 running on nid00564
This is Process-32/36 running on nid00564
This is Process-34/36 running on nid00564
This program uses 36 processes
The number of intervals = 1000000000
```

pi is approximately 3.1415926535898406
Error is 0.00000000000000475
wall clock time = 0.096959

OpenMP Result
running on 36 threads: PI = 3.141592653589822 computed in
0.1616 seconds

Two Nodes

(some lines removed)

```
This is Process-62/72 running on nid01175
This is Process-22/72 running on nid01174
This is Process-63/72 running on nid01175
This is Process-23/72 running on nid01174
This is Process-64/72 running on nid01175
This is Process-24/72 running on nid01174
This is Process-65/72 running on nid01175
This is Process-25/72 running on nid01174
This is Process-66/72 running on nid01175
This is Process-26/72 running on nid01174
This is Process-67/72 running on nid01175
This is Process-29/72 running on nid01174
This is Process-68/72 running on nid01175
This is Process-30/72 running on nid01174
This is Process-69/72 running on nid01175
This is Process-31/72 running on nid01174
This is Process-70/72 running on nid01175
This is Process-32/72 running on nid01174
This is Process-71/72 running on nid01175
This is Process-33/72 running on nid01174
This is Process-39/72 running on nid01175
This is Process-34/72 running on nid01174
This is Process-40/72 running on nid01175
This is Process-35/72 running on nid01174
This is Process-41/72 running on nid01175
This is Process-0/72 running on nid01174
This is Process-43/72 running on nid01175
This is Process-5/72 running on nid01174
This is Process-44/72 running on nid01175
This is Process-7/72 running on nid01174
This is Process-45/72 running on nid01175
This is Process-11/72 running on nid01174
This is Process-46/72 running on nid01175
This is Process-13/72 running on nid01174
This is Process-47/72 running on nid01175
This is Process-19/72 running on nid01174
This is Process-51/72 running on nid01175
This is Process-20/72 running on nid01174
This is Process-59/72 running on nid01175
This is Process-27/72 running on nid01174
This is Process-28/72 running on nid01174
This program uses 72 processes
The number of intervals = 1000000000
```

pi is approximately 3.1415926535898109,
Error is 0.00000000000000178
wall clock time = 0.024506

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=(nodes)
#SBATCH --ntasks-per-node=(processes)
#SBATCH --cpus-per-task=(threads)
#SBATCH --partition=normal
#SBATCH --constraint=mc

srun cpi_mpi
```




MPI Results

```
This is Process-1/36 running on nid00564
This is Process-3/36 running on nid00564
This is Process-5/36 running on nid00564
This is Process-6/36 running on nid00564
This is Process-8/36 running on nid00564
This is Process-11/36 running on nid00564
This is Process-13/36 running on nid00564
This is Process-14/36 running on nid00564
This is Process-17/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-22/36 running on nid00564
This is Process-23/36 running on nid00564
This is Process-25/36 running on nid00564
This is Process-26/36 running on nid00564
This is Process-28/36 running on nid00564
This is Process-29/36 running on nid00564
This is Process-31/36 running on nid00564
This is Process-33/36 running on nid00564
This is Process-35/36 running on nid00564
This is Process-0/36 running on nid00564
This is Process-2/36 running on nid00564
This is Process-4/36 running on nid00564
This is Process-7/36 running on nid00564
This is Process-9/36 running on nid00564
This is Process-10/36 running on nid00564
This is Process-12/36 running on nid00564
This is Process-15/36 running on nid00564
This is Process-16/36 running on nid00564
This is Process-18/36 running on nid00564
This is Process-20/36 running on nid00564
This is Process-21/36 running on nid00564
This is Process-24/36 running on nid00564
This is Process-27/36 running on nid00564
This is Process-30/36 running on nid00564
This is Process-32/36 running on nid00564
This is Process-34/36 running on nid00564
This program uses 36 processes
The number of intervals = 1000000000
```

pi is approximately 3.1415926535898406
Error is 0.00000000000000475
wall clock time = 0.096959

OpenMP Result
running on 36 threads: PI = 3.141592653589822 computed in
0.1616 seconds

Two Nodes

(some lines removed)

```
This is Process-62/72 running on nid01175
This is Process-22/72 running on nid01174
This is Process-63/72 running on nid01175
This is Process-23/72 running on nid01174
This is Process-64/72 running on nid01175
This is Process-24/72 running on nid01174
This is Process-65/72 running on nid01175
This is Process-25/72 running on nid01174
This is Process-66/72 running on nid01175
This is Process-26/72 running on nid01174
This is Process-67/72 running on nid01175
This is Process-29/72 running on nid01174
This is Process-68/72 running on nid01175
This is Process-30/72 running on nid01174
This is Process-69/72 running on nid01175
This is Process-31/72 running on nid01174
This is Process-70/72 running on nid01175
This is Process-32/72 running on nid01174
This is Process-71/72 running on nid01175
This is Process-33/72 running on nid01174
This is Process-39/72 running on nid01175
This is Process-34/72 running on nid01174
This is Process-40/72 running on nid01175
This is Process-35/72 running on nid01174
This is Process-41/72 running on nid01175
This is Process-0/72 running on nid01174
This is Process-43/72 running on nid01175
This is Process-5/72 running on nid01174
This is Process-44/72 running on nid01175
This is Process-7/72 running on nid01174
This is Process-45/72 running on nid01175
This is Process-11/72 running on nid01174
This is Process-46/72 running on nid01175
This is Process-13/72 running on nid01174
This is Process-47/72 running on nid01175
This is Process-19/72 running on nid01174
This is Process-51/72 running on nid01175
This is Process-20/72 running on nid01174
This is Process-59/72 running on nid01175
This is Process-27/72 running on nid01174
This is Process-28/72 running on nid01174
This program uses 72 processes
The number of intervals = 1000000000
```

pi is approximately 3.1415926535898109,
Error is 0.00000000000000178
wall clock time = 0.024506

```
#!/bin/bash -l
#SBATCH --account=uzh8
#SBATCH --job-name=openmp_test
#SBATCH --time=01:00:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=36
#SBATCH --cpus-per-task=1
#SBATCH --partition=normal
#SBATCH --constraint=mc

srun cpi_mpi
```

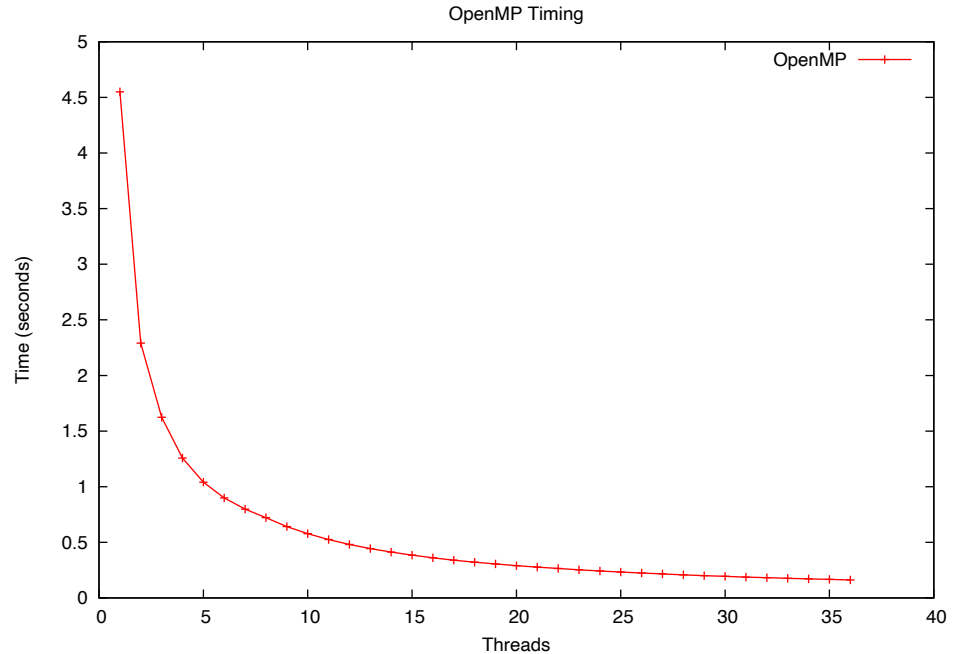
Last
week

Same
results
in table
format

| THREADS | | SECONDS | | THREADS | | SECONDS | |
|---------|--|---------|--|---------|--|---------|--|
| 1 | | 4.550 | | 19 | | 0.3053 | |
| 2 | | 2.291 | | 20 | | 0.2897 | |
| 3 | | 1.624 | | 21 | | 0.2768 | |
| 4 | | 1.258 | | 22 | | 0.2644 | |
| 5 | | 1.041 | | 23 | | 0.2528 | |
| 6 | | 0.8988 | | 24 | | 0.2423 | |
| 7 | | 0.7989 | | 25 | | 0.2326 | |
| 8 | | 0.7217 | | 26 | | 0.2237 | |
| 9 | | 0.6415 | | 27 | | 0.2154 | |
| 10 | | 0.5774 | | 28 | | 0.2077 | |
| 11 | | 0.5249 | | 29 | | 0.2006 | |
| 12 | | 0.4811 | | 30 | | 0.1939 | |
| 13 | | 0.4441 | | 31 | | 0.1876 | |
| 14 | | 0.4124 | | 32 | | 0.1818 | |
| 15 | | 0.3849 | | 33 | | 0.1763 | |
| 16 | | 0.3609 | | 34 | | 0.1711 | |
| 17 | | 0.3397 | | 35 | | 0.1662 | |
| 18 | | 0.3208 | | 36 | | 0.1616 | |

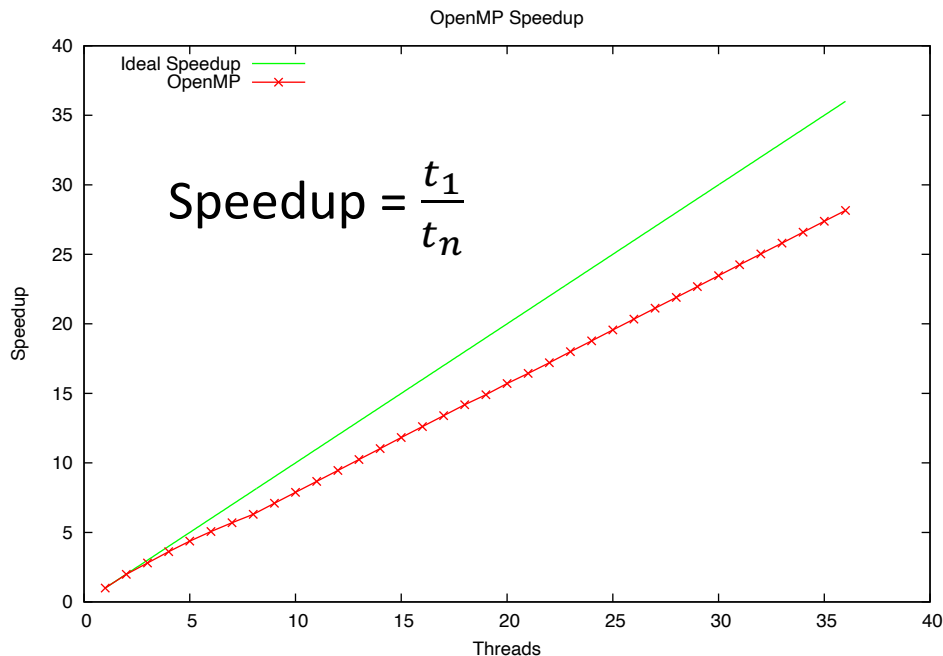
Timing Plot

```
set title 'OpenMP Timing'
set xlabel 'Threads'
set ylabel 'Time (seconds)'
set key top right
plot "cpi_openmp.dat" \
  u 1:2 w lp lw 2 t 'OpenMP'
```



“Speedup” Plot

```
set title 'OpenMP Speedup'  
set xlabel 'Threads'  
set ylabel 'Speedup'  
set key top left  
plot x lc 2 lw 2 t 'Ideal Speedup',\  
"cpi_openmp.dat" u 1:(4.55/$2)\  
w lp lc 1 lw 2 t 'OpenMP'
```



Parsing the Text (Results)



running on 1 threads: $\pi = 3.141592653589971$ computed in 4.55 seconds

| | | | | | |
|----|-----------------|----|----------|----|----------|
| 1 | 4.550000 | 13 | 0.444100 | 25 | 0.232600 |
| 2 | 2.291000 | 14 | 0.412400 | 26 | 0.223700 |
| 3 | 1.624000 | 15 | 0.384900 | 27 | 0.215400 |
| 4 | 1.258000 | 16 | 0.360900 | 28 | 0.207700 |
| 5 | 1.041000 | 17 | 0.339700 | 29 | 0.200600 |
| 6 | 0.898800 | 18 | 0.320800 | 30 | 0.193900 |
| 7 | 0.798900 | 19 | 0.305300 | 31 | 0.187600 |
| 8 | 0.721700 | 20 | 0.289700 | 32 | 0.181800 |
| 9 | 0.641500 | 21 | 0.276800 | 33 | 0.176300 |
| 10 | 0.577400 | 22 | 0.264400 | 34 | 0.171100 |
| 11 | 0.524900 | 23 | 0.252800 | 35 | 0.166200 |
| 12 | 0.481100 | 24 | 0.242300 | 36 | 0.161600 |


```
./parse.pl <slurm-6188431.out
```

How to Plot the Results

```
dpotter@daint104:~/hpc1a> cat parse.pl
#!/usr/bin/perl
use strict;

while(<STDIN>) {
    if (/running on (\d+) threads: PI.*computed in ([0-9]*\.[0-9]+) seconds/) {
        printf("%2d %f\n", $1, $2);
    }
}
```

```
dpotter@daint104:~/hpc1a> cat speedup.gp
set title 'OpenMP Speedup'
set xlabel 'Threads'
set ylabel 'Speedup'
set key top left
plot x t 'Ideal Speedup', "cpi_openmp.dat" u 1:(4.55/$2) w l t 'OpenMP'
```



```
./parse.pl <slurm-6188431.out >cpi_openmp.dat
```

```
gnuplot> set output "speedup.eps"
gnuplot> set terminal postscript enhanced solid color
gnuplot> load "speedup.gp"
gnuplot> set terminal x11
gnuplot> set output
```

Parsing with Python

```
dpotter@daint104:~/hpc1a> cat parse.py
```

```
#!/usr/bin/env python
```

```
import re, sys
```

```
for line in re.finditer(
    r"running on ([0-9]+) threads: PI.*computed in ([0-9]*\.[0-9]+) seconds",
    sys.stdin.read(), re.MULTILINE):
    print("{} {}".format(line.group(1), line.group(2)))
```

```
dpotter@daint104:~/hpc1a> cat speedup.gp
```


```
set title 'OpenMP Speedup'
```

```
set xlabel 'Threads'
```

```
set ylabel 'Speedup'
```

```
set key top left
```

```
plot x t 'Ideal Speedup', "cpi_openmp.dat" u 1:(4.55/$2) w l t 'OpenMP'
```



```
./parse.py <slurm-6188431.out >cpi_openmp.dat
```

```
gnuplot> set output "speedup.eps"
gnuplot> set terminal postscript enhanced solid color
gnuplot> load "speedup.gp"
gnuplot> set terminal x11
gnuplot> set output
```



Parsing Process

```
[dpotter@ela4 hpc1a]$ head -n 36 slurm-6188431.out
running on 1 threads: PI = 3.141592653589971 computed in 4.55 seconds
running on 2 threads: PI = 3.141592653589901 computed in 2.291 seconds
running on 3 threads: PI = 3.141592653589962 computed in 1.624 seconds
running on 4 threads: PI = 3.141592653589821 computed in 1.258 seconds
running on 5 threads: PI = 3.141592653589596 computed in 1.041 seconds
running on 6 threads: PI = 3.141592653589682 computed in 0.8988 seconds
running on 7 threads: PI = 3.14159265358963 computed in 0.7989 seconds
running on 8 threads: PI = 3.141592653589769 computed in 0.7217 seconds
running on 9 threads: PI = 3.141592653589656 computed in 0.6415 seconds
running on 10 threads: PI = 3.141592653589794 computed in 0.5774 seconds
running on 11 threads: PI = 3.14159265358966 computed in 0.5249 seconds
running on 12 threads: PI = 3.14159265358986 computed in 0.4811 seconds
running on 13 threads: PI = 3.141592653589865 computed in 0.4441 seconds
running on 14 threads: PI = 3.141592653589788 computed in 0.4124 seconds
running on 15 threads: PI = 3.141592653589805 computed in 0.3849 seconds
running on 16 threads: PI = 3.141592653589832 computed in 0.3609 seconds
running on 17 threads: PI = 3.141592653589839 computed in 0.3397 seconds
running on 18 threads: PI = 3.141592653589814 computed in 0.3208 seconds
running on 19 threads: PI = 3.141592653589826 computed in 0.3053 seconds
running on 20 threads: PI = 3.141592653589855 computed in 0.2897 seconds
running on 21 threads: PI = 3.141592653589775 computed in 0.2768 seconds
running on 22 threads: PI = 3.141592653589823 computed in 0.2644 seconds
running on 23 threads: PI = 3.141592653589866 computed in 0.2528 seconds
running on 24 threads: PI = 3.141592653589792 computed in 0.2423 seconds
running on 25 threads: PI = 3.14159265358978 computed in 0.2326 seconds
running on 26 threads: PI = 3.141592653589832 computed in 0.2237 seconds
running on 27 threads: PI = 3.141592653589835 computed in 0.2154 seconds
running on 28 threads: PI = 3.141592653589816 computed in 0.2077 seconds
running on 29 threads: PI = 3.141592653589819 computed in 0.2006 seconds
running on 30 threads: PI = 3.141592653589893 computed in 0.1939 seconds
running on 31 threads: PI = 3.141592653589744 computed in 0.1876 seconds
running on 32 threads: PI = 3.141592653589758 computed in 0.1818 seconds
running on 33 threads: PI = 3.141592653589806 computed in 0.1763 seconds
running on 34 threads: PI = 3.141592653589926 computed in 0.1711 seconds
running on 35 threads: PI = 3.14159265358979 computed in 0.1662 seconds
running on 36 threads: PI = 3.141592653589822 computed in 0.1616 seconds
```

```
[dpotter@ela4 hpc1a]$ python parse.py <slurm-6188431.out
1 4.55
2 2.291
3 1.624
4 1.258
5 1.041
6 0.8988
7 0.7989
8 0.7217
9 0.6415
10 0.5774
11 0.5249
12 0.4811
13 0.4441
14 0.4124
15 0.3849
16 0.3609
17 0.3397
18 0.3208
19 0.3053
20 0.2897
21 0.2768
22 0.2644
23 0.2528
24 0.2423
25 0.2326
26 0.2237
27 0.2154
28 0.2077
29 0.2006
30 0.1939
31 0.1876
32 0.1818
33 0.1763
34 0.1711
35 0.1662
36 0.1616
```


More on Compilers

High Performance Computing





Compiler Suites

[Cray Compiler](#) ([PrgEnv-cray](#))

[GNU Compiler](#) ([PrgEnv-gnu](#))

- Free for everyone

[Intel Compiler](#) ([PrgEnv-intel](#))

- [Free for Students](#)

[Portland Compiler](#) ([PrgEnv-pgi](#))

- Important for OpenACC (GPU)

[clang](#) (Apple mostly)

- Open Source & GNU Compatible

[Visual Studio](#) (Windows)

- [Free for us through Microsoft Imagine](#)

```
#include <stdio.h>
#include <sys/time.h>

static long steps = 1000000000;

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;
    char *p;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

cpi.c
(serial
version)

Output Name

```
dhcp-94-191:cpu$ ls
cpu.c
dhcp-94-191:cpu$ gcc cpu.c
dhcp-94-191:cpu$ ls
a.out cpu.c
dhcp-94-191:cpu$ ./a.out
PI = 3.141592653589971 computed in 11.9 seconds
dhcp-94-191:cpu$ rm a.out
dhcp-94-191:cpu$ ls
cpu.c
dhcp-94-191:cpu$ gcc -o cpu cpu.c
dhcp-94-191:cpu$ ls
cpu cpu.c
dhcp-94-191:cpu$ ./cpu
PI = 3.141592653589971 computed in 11.82 seconds
```

Compiler Optimization

```
dhcp-94-191:cpi$ gcc -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 11.84 seconds
dhcp-94-191:cpi$ gcc -O -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 2.505 seconds
dhcp-94-191:cpi$ gcc -O0 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 11.82 seconds
dhcp-94-191:cpi$ gcc -O1 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 2.497 seconds
dhcp-94-191:cpi$ gcc -O2 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 1.19 seconds
dhcp-94-191:cpi$ gcc -O3 -o cpi cpi.c ; ./cpi
PI = 3.141592653589971 computed in 1.174 seconds
dhcp-94-191:cpi$ gcc -O3 -ffast-math -o cpi cpi.c ; ./cpi
PI = 3.141592653589652 computed in 0.6 seconds
dhcp-94-191:cpi$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c ; ./cpi
PI = 3.141592653589729 computed in 0.5518 seconds
3.14159265358979323846 (real value of pi)
```



> 20 Times
Improvement

What do the levels mean?

-O1 Optimize.

- Optimizing compilation takes somewhat more time, and a lot more memory for a large function.

-O2 Optimize even more.

- GCC performs nearly all supported optimizations that do not involve a space-speed tradeoff.
- The compiler does not perform loop unrolling or function inlining when you specify -O2.
- As compared to -O, this option **increases both compilation time** and the performance of the generated code.

-O3 Optimize yet more.

- -O3 turns on all optimizations specified by -O2 and also turns on the -finline-functions, -funswitch-loops, -fpredictive-commoning, -fgcse-after-reload and -ftree-vectorize options.

-O0 No Optimization.

- Reduce compilation time and make debugging produce the expected results.
- **This is the default.**

Fast Math

Associative Property of Addition

$$(a + b) + c = a + (b + c)$$

Distributive Property of Multiplication

$$a(b + c) = ab + ac$$

Sometimes reordering can
improve performance

Fast Math (the Truth)

Associative Property of Addition

$$(a + b) + c \neq a + (b + c)$$

Distributive Property of Multiplication

$$a(b + c) \neq ab + ac$$

Sometimes reordering can
change the result

Beware of Associative Math

```
dhcp-94-191:cpi$ python
```

```
>>> a=12
```

```
>>> b=12
```

```
>>> a==b
```

```
True
```

```
>>> a=0.1 + 0.2 + 0.3
```

```
>>> b=0.1 +(0.2 + 0.3)
```

```
>>> a==b
```

```
False
```

```
>>> print(a)
```

```
0.6
```

```
>>> print(b)
```

```
0.6
```

```
>>> print("%.17f" %a)
```

```
0.600000000000000009
```

```
>>> print("%.17f" %b)
```

```
0.59999999999999998
```

```
>>> c=0.6
```

```
>>> print("%.17f" %c)
```

```
0.59999999999999998
```

“Modules” in C

```
#include <stdio.h>
#include <sys/time.h>

static long steps = 1000000000;

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}

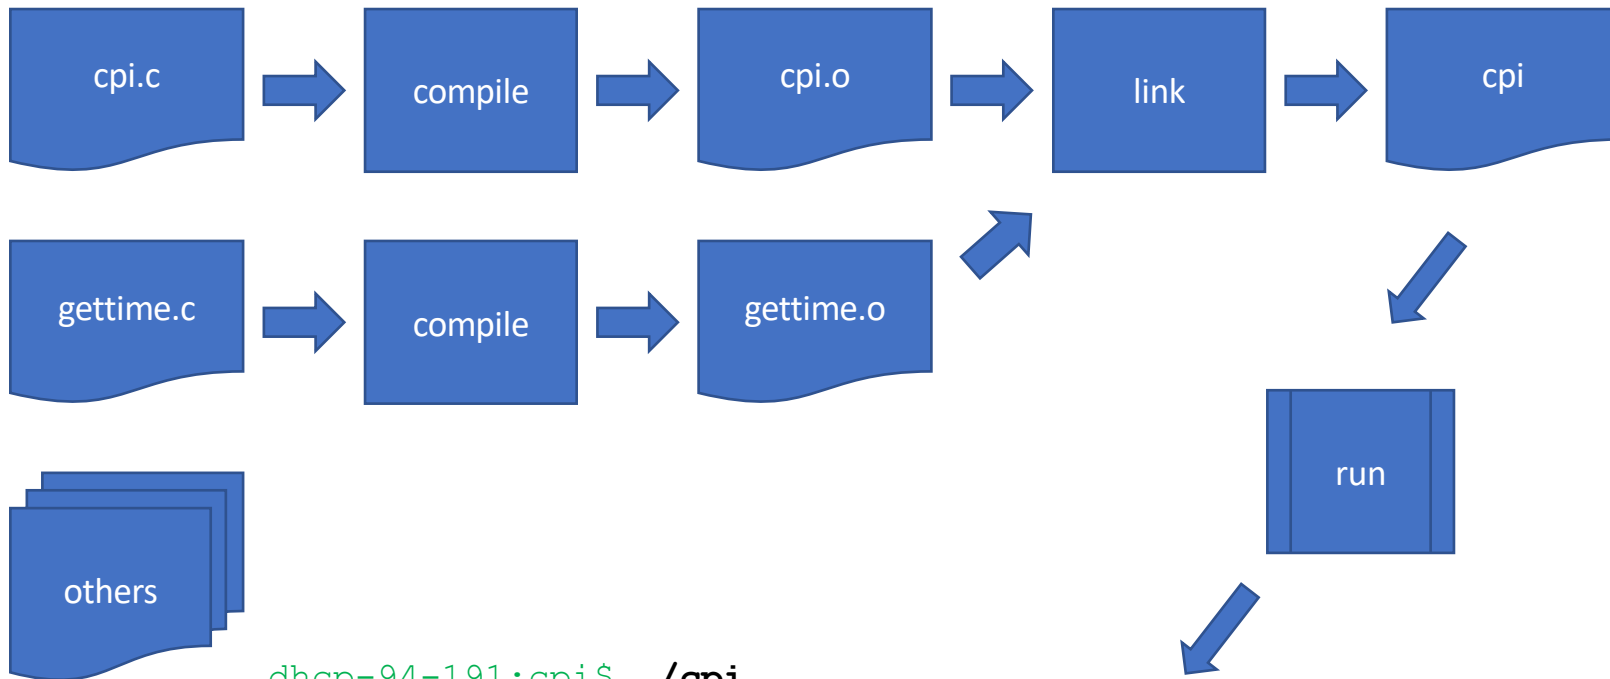
int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Separate Compilation



```
dhcp-94-191:cpi$ ./cpi
```

```
PI = 3.141592653589971 computed in 1.174 seconds
```



gettime.c

```
#include <sys/time.h>

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}
```



cpi.c

```
#include <stdio.h>

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Compiling with two files

```
dhcp-94-191:cpu2$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c gettime.c
```

```
cpi.c: In function 'main':
```

```
cpi.c:13:20: warning: implicit declaration of function 'getTime'; did you mean 'getline'? [-Wimplicit-function-declaration]
```

```
    double start = getTime();
```

```
                ^~~~~~
```

```
                getline
```



University of
Zurich ^{UZH}

gettime.h

```
double getTime(void);
```



cpi.c

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();
    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

gettime.h

```
double getTime(void);
```




gettime.c

```
#include <sys/time.h>
#include "gettime.h"

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}
```

gettime.h

```
double getTime(void);
```



Compile and Link Separately

```
dhcp-94-191:cp12$ gcc -O3 -ffast-math -mavx2 -o cpi cpi.c gettime.c
dhcp-94-191:cp12$ ls
cpi      cpi.c    gettime.c gettime.h
dhcp-94-191:cp12$ rm cpi
```

```
dhcp-94-191:cp12$ ls
cpi.c    gettime.c gettime.h
dhcp-94-191:cp12$ gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
dhcp-94-191:cp12$ ls
cpi.c    cpi.o    gettime.c gettime.h
dhcp-94-191:cp12$ gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c
dhcp-94-191:cp12$ ls
cpi.c    cpi.o    gettime.c gettime.h gettime.o
dhcp-94-191:cp12$ gcc -o cpi cpi.o gettime.o
dhcp-94-191:cp12$ ./cpi
PI = 3.141592653589729 computed in 0.567 seconds
```



Makefile

```

cpu
    : cpu.o gettime.o
    gcc -o cpu cpu.o gettime.o

cpu.o
    : cpu.c gettime.h
    gcc -O3 -ffast-math -mavx2 -c -o cpu.o cpu.c

gettime.o
    : gettime.c gettime.h
    gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c

clean:
    rm -f cpu cpu.o gettime.o
```



Compile and Link Separately

```
dhcp-94-191:cpu2$ make clean
rm -f cpi cpi.o gettime.o
dhcp-94-191:cpu2$ make
gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c
gcc -o cpi cpi.o gettime.o
dhcp-94-191:cpu2$ make
make: `cpi' is up to date.
dhcp-94-191:cpu2$ touch cpi.c
dhcp-94-191:cpu2$ make
gcc -O3 -ffast-math -mavx2 -c -o cpi.o cpi.c
gcc -o cpi cpi.o gettime.o
```

“touch” – change modification time.
Same as if you edited the file.



Makefile

```

cpu      : cpu.o gettime.o
            gcc -o cpu cpu.o gettime.o

cpu.o    : cpu.c gettime.h
            gcc -O3 -ffast-math -mavx2 -c -o cpu.o cpu.c

gettime.o: gettime.c gettime.h
            gcc -O3 -ffast-math -mavx2 -c -o gettime.o gettime.c

clean:
            rm -f cpu cpu.o gettime.o
```



Makefile with default rules

```
cp1      : cp1.o gettime.o
```

```
cp1.o    : cp1.c gettime.h
```

```
gettime.o: gettime.c gettime.h
```

```
clean:
```

```
    rm -f cp1 cp1.o gettime.o
```

```
dhcp-94-191:cp12$ make clean
```

```
rm -f cp1 cp1.o gettime.o
```

```
dhcp-94-191:cp12$ make
```

```
cc      -c -o cp1.o cp1.c
```

```
cc      -c -o gettime.o gettime.c
```

```
cc      cp1.o gettime.o      -o cp1
```



Makefile with default rules

```
CFLAGS=-O3 -ffast-math -mavx2
```

```
CC=gcc
```

```
cpu      : cpu.o gettime.o
```

```
cpu.o    : cpu.c gettime.h
```

```
gettime.o: gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpu cpu.o gettime.o
```

```
dhcp-94-191:cpu2$ make clean
```

```
rm -f cpu cpu.o gettime.o
```

```
dhcp-94-191:cpu2$ make
```

```
gcc -O3 -ffast-math -mavx2      -c -o cpu.o cpu.c
```

```
gcc -O3 -ffast-math -mavx2      -c -o gettime.o gettime.c
```

```
gcc      cpu.o gettime.o      -o cpu
```



Message Passing

Why use MPI?

One Node
isn't enough

- Maybe you have run out of memory
- Maybe it takes too long to calculate


Alternatives
do exist

- Partitioned Global Address Space (PGAS)
- HPX (High Performance ParallelX)
- Charm++

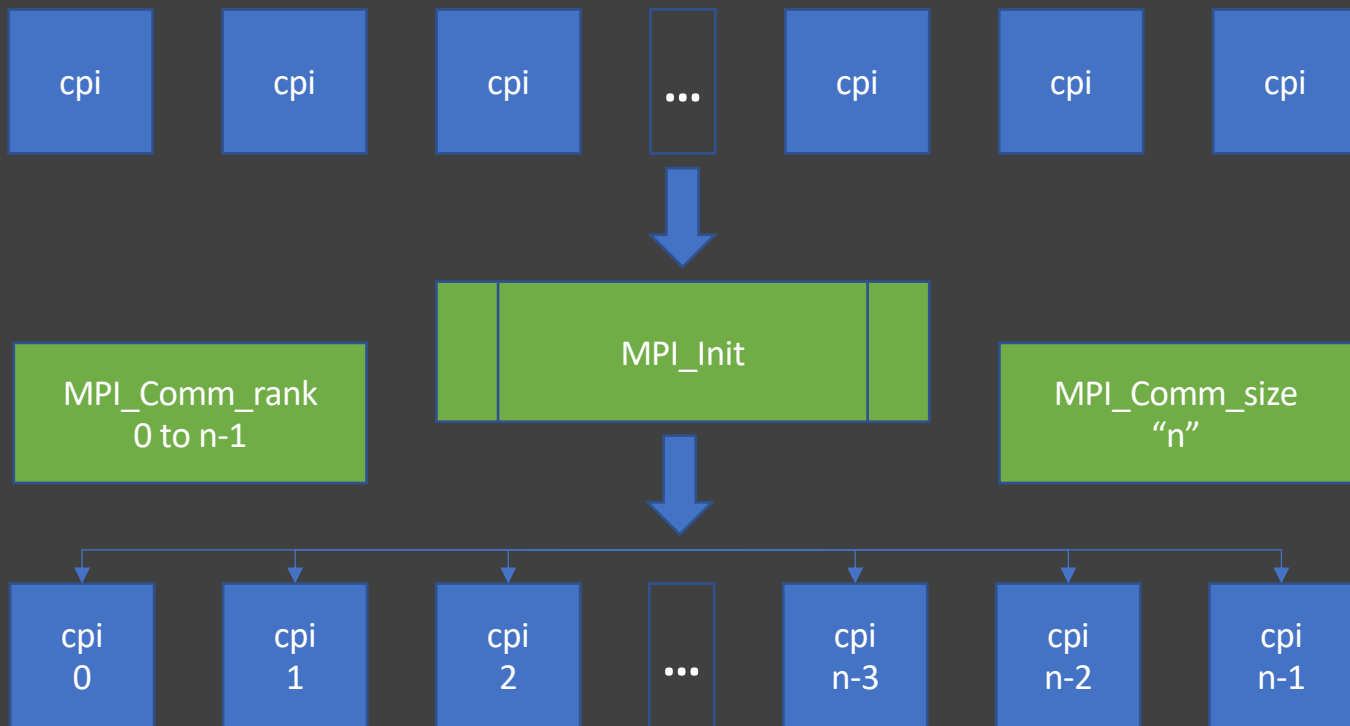
MPI is still
omnipresent

- It has been around since 1991
- Version 3.1 was approved in 2015
- Version 4.0 release candidate now (2021)

MPI Process Layout

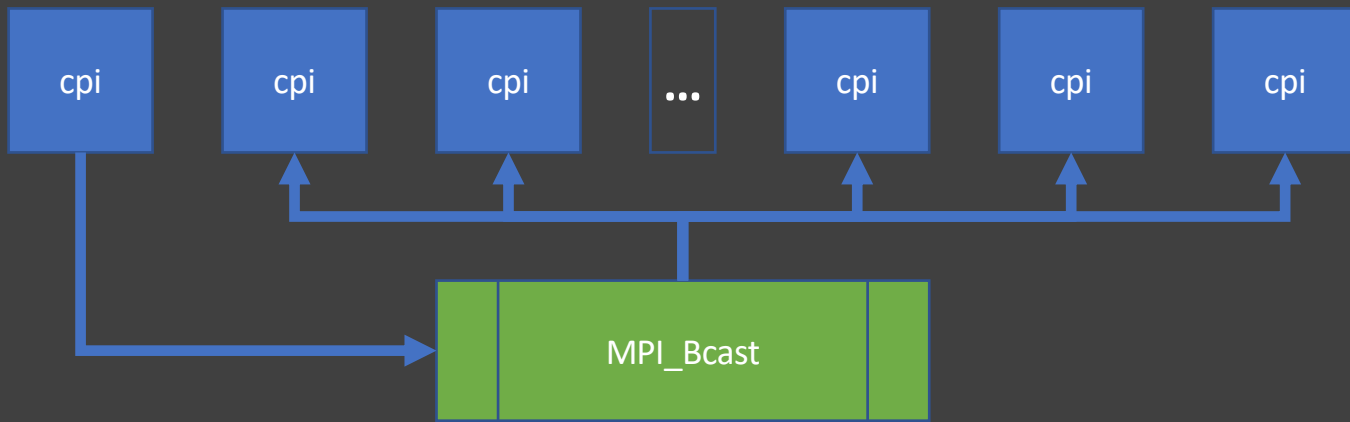


MPI Initialization





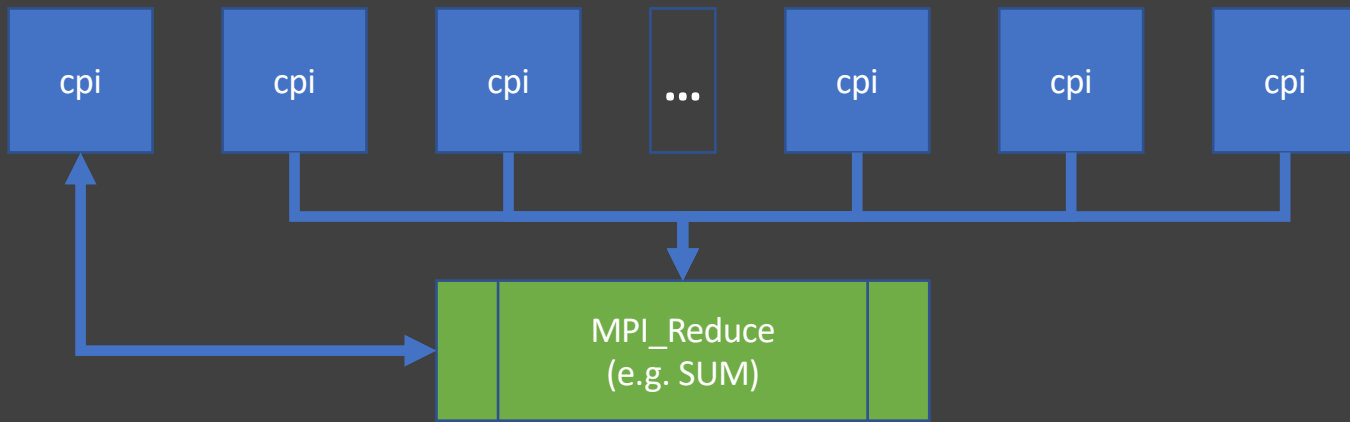
MPI Broadcast



All MPI Ranks Call MPI_Bcast together.
Everyone gets the value from (usually) rank 0.



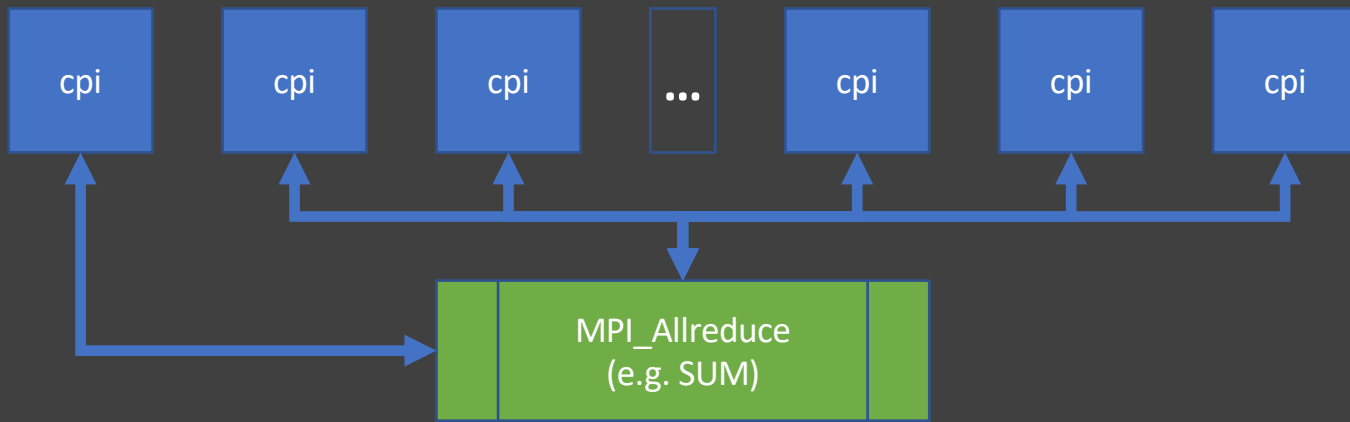
MPI Reduce



All MPI Ranks Call MPI_Reduce together.
All values from all ranks are summed together.
A single rank (usually rank 0) gets the result.



MPI All Reduce





All MPI Ranks Call MPI_Reduce together.
All values from all ranks are summed together.
All ranks get the result.



Partial Reduction (e.g., sum)

| Rank | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|--------|---|---|----|----|----|----|----|----|
| Value | 4 | 3 | 8 | 7 | 2 | 6 | 3 | 1 |
| Scan | 4 | 7 | 15 | 22 | 24 | 30 | 33 | 34 |
| Exscan | 0 | 4 | 7 | 15 | 22 | 24 | 30 | 33 |





The MPI version of cpi

```
MPI_Init(&argc,&argv);          /* Connect processes to each other */
MPI_Comm_size(MPI_COMM_WORLD,&numprocs); /* Get total number of processes */
MPI_Comm_rank(MPI_COMM_WORLD,&myid);      /* Rank of this process */
MPI_Get_processor_name(name, &resultlen);

printf("This is Process-%d/%d running on %s \n",myid,numprocs,name);
MPI_Barrier(MPI_COMM_WORLD);

if(myid == 0) {
    printf("This program uses %d processes\n", numprocs);
    n = 1000000000;
}

MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); /* Send n from 0 to all others */

sum = 0.0;
h = 1.0/n;
for (i=myid+0.5; i<n; i+=numprocs) {
    sum += dx_arctan(i*h);
}
mypi = 4.0*h*sum;

/* Compute the Derivative of ArcTan */
double dx_arctan(double x) {
    return 1.0 / (1.0 + x*x);
}

/* Add all partial sums to each other and send to rank 0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

if (myid == 0) {
    printf("pi is approximately %.16f\n",pi);
}
MPI_Finalize();                /* Disconnect all processes */
```





[MPI Publish name](#)

[MPI Comm c2f](#)

[MPI Comm call errhandler](#)

[MPI Comm compare](#)

[MPI Comm connect](#)

[MPI Comm create](#)

[MPI Comm create errhandler](#)

[MPI Comm create group](#)

[MPI Comm create keyval](#)

[MPI Comm delete attr](#)

[MPI Comm disconnect](#)

[MPI Comm dup](#)

[MPI Comm dup_with_info](#)

[MPI Comm f2c](#)

[MPI Comm free](#)

[MPI Comm free keyval](#)

[MPI Comm get attr](#)

[MPI Comm get errhandler](#)

[MPI Comm get info](#)

[MPI Comm get name](#)

[MPI Comm get parent](#)

[MPI Comm group](#)

[MPI Comm idup](#)

[MPI Comm join](#)

[MPI Comm rank](#)

[MPI Comm remote group](#)

[MPI Comm remote size](#)

[MPI Comm set attr](#)

[MPI Comm set errhandler](#)

[MPI Comm set info](#)

[MPI Comm set name](#)

[MPI Comm size](#)

[MPI Comm spawn](#)

[MPI Comm spawn multiple](#)

[MPI Comm split](#)

[MPI Comm split type](#)

[MPI Comm test inter](#)

[MPI Compare and swap](#)

[MPI Dims create](#)

[MPI Dist_graph create](#)

[MPI Dist_graph create adjacent](#)

[MPI Dist_graph neighbors](#)

[MPI Dist_graph neighbors count](#)

[MPI File write ordered end](#)

[MPI File write shared](#)

[MPI Finalize](#)

[MPI Finalized](#)

[MPI Free mem](#)

[MPI Gather](#)

[MPI Gatherv](#)

[MPI Get](#)

[MPI Get accumulate](#)

[MPI Get address](#)

[MPI Get count](#)

[MPI Get elements](#)

[MPI Get elements x](#)

[MPI Get library version](#)

[MPI Get processor name](#)

[MPI Get version](#)

[MPI Graph create](#)

[MPI Graph_get](#)

[MPI Graph map](#)

[MPI Graph neighbors](#)

[MPI Graph_neighbors_count](#)

[MPI Graphdims_get](#)

[MPI Grequest complete](#)

[MPI Grequest start](#)

[MPI Group c2f](#)

[MPI Group compare](#)

[MPI Group difference](#)

[MPI Group excl](#)

[MPI Group f2c](#)

[MPI Group free](#)

[MPI Group incl](#)

[MPI Group intersection](#)

[MPI Group range excl](#)

[MPI Group range incl](#)

[MPI Group rank](#)

[MPI Group size](#)

[MPI Group translate ranks](#)

[MPI Group union](#)

[MPI Iallgather](#)

[MPI Iallgatherv](#)

[MPI Iallreduce](#)

[MPI Ialltoall](#)

[MPI Ialltoallv](#)

[MPI Put](#)

[MPI Query thread](#)

[MPI Raccumulate](#)

[MPI Recv](#)

[MPI Recv_init](#)

[MPI Reduce](#)

[MPI Reduce local](#)

[MPI Reduce scatter](#)

[MPI Reduce scatter block](#)

[MPI Register datrep](#)

[MPI Request c2f](#)

[MPI Request f2c](#)

[MPI Request free](#)

[MPI Request_get_status](#)

[MPI Rget](#)

[MPI Rget accumulate](#)

[MPI Rput](#)

[MPI Rsend](#)

[MPI Rsend_init](#)

[MPI Scan](#)

[MPI Scatter](#)

[MPI Scatterv](#)

[MPI Send](#)

[MPI Send_init](#)

[MPI Sendrecv](#)

[MPI Sendrecv replace](#)

[MPI Sizeof](#)

[MPI Ssend](#)

[MPI Ssend_init](#)

[MPI Start](#)

[MPI Startall](#)

[MPI Status c2f](#)

[MPI Status f2c](#)

[MPI Status set cancelled](#)

[MPI Status set elements](#)

[MPI Status set elements x](#)

[MPI T category changed](#)

[MPI T category_get categories](#)

[MPI T category_get cvars](#)

[MPI T category_get info](#)

[MPI T category_get num](#)

[MPI T category_get pvars](#)

API
tion 3 man pages)

[MPI File call errhandler](#)

[MPI File close](#)

[MPI File create errhandler](#)

[MPI File delete](#)

[MPI File f2c](#)

[MPI File get amode](#)

[MPI File get atomicity](#)

[MPI File get byte offset](#)

[MPI File get errhandler](#)

[MPI File get_group](#)

[MPI File get info](#)

[MPI File get position](#)

[MPI File get_position shared](#)

[MPI File get size](#)

[MPI File get_type_extent](#)

[MPI File get view](#)

[MPI File iread](#)

[MPI File iread_all](#)

[MPI File iread_at](#)

[MPI File iread_at_all](#)

[MPI File iread_shared](#)

[MPI File iread_shared](#)

[MPI File iread_at](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all_begin](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File iread_at_all_end](#)

[MPI File create errhandler](#)

[MPI File close](#)

[MPI File create errhandler](#)

[MPI File delete](#)

[MPI File f2c](#)

[MPI File get amode](#)

[MPI File get atomicity](#)

[MPI File get byte offset](#)

[MPI File get errhandler](#)

[MPI File get_group](#)

[MPI File get info](#)

[MPI File get position](#)

[MPI File get_position shared](#)

[MPI File get size](#)

[MPI File get_type_extent](#)

[MPI File get view](#)

[MPI File iread](#)

[MPI File iread_all](#)

[MPI File iread_at](#)

[MPI File iread_at_all](#)

[MPI File iread_shared](#)

[MPI File iread_shared](#)

[MPI File iread_at](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI File iread_at_all](#)

[MPI Ineighbor allgather](#)

[MPI Ineighbor allgatherv](#)

[MPI Ineighbor alltoall](#)

[MPI Ineighbor alltoallv](#)

[MPI Ineighbor alltoallw](#)

[MPI Info c2f](#)

[MPI Info create](#)

[MPI Info delete](#)

[MPI Info dup](#)

[MPI Info env](#)

[MPI Info f2c](#)

[MPI Info free](#)

[MPI Info get](#)

[MPI Info get_nkeys](#)

[MPI Info get_nthkey](#)

[MPI Info get_valuelen](#)

[MPI Info set](#)

[MPI Init](#)

[MPI Init_thread](#)

[MPI Initialized](#)

[MPI Intercomm create](#)

[MPI Intercomm merge](#)

[MPI Iprobe](#)

[MPI Irecv](#)

[MPI Ireduce](#)

[MPI Ireduce scatter](#)

[MPI Ireduce scatter block](#)

[MPI Isend](#)

[MPI Is_thread_main](#)

[MPI Iscan](#)

[MPI Iscatter](#)

[MPI Iscatterv](#)

[MPI Isend](#)

[MPI Issend](#)

[MPI Keyval create](#)

[MPI Keyval free](#)

[MPI Lookup_name](#)

[MPI Message c2f](#)

[MPI Message f2c](#)

[MPI Mprobe](#)

[MPI Mrecv](#)

[MPI Neighbor allgather](#)

[MPI Neighbor allgatherv](#)

[MPI Neighbor alltoall](#)

[MPI Neighbor alltoallv](#)

[MPI Neighbor alltoallw](#)

[MPI Op c2f](#)

[MPI Op commutative](#)

[MPI Op create](#)

[MPI Op f2c](#)

[MPI Op free](#)

[MPI Open_port](#)

[MPI Pack](#)

[MPI Pack_external](#)

[MPI T init thread](#)

[MPI T pvar_get_info](#)

[MPI T pvar_get_num](#)

[MPI T pvar_handle_alloc](#)

[MPI T pvar_handle_free](#)

[MPI T pvar_read](#)

[MPI T pvar_readreset](#)

[MPI T pvar_reset](#)

[MPI T pvar_session_create](#)

[MPI T pvar_session_free](#)

[MPI T pvar_start](#)

[MPI T pvar_stop](#)

[MPI T pvar_write](#)

[MPI Test](#)

[MPI Test_cancelled](#)

[MPI Testall](#)

[MPI Testany](#)

[MPI Testsome](#)

[MPI Topo_test](#)

[MPI Type c2f](#)

[MPI Type commit](#)

[MPI Type contiguous](#)



[MPI Type create_darray](#)

[MPI Type create_f90_co](#)

[MPI Type create_f90_int](#)

[MPI Type create_f90_re](#)

[MPI Type create_hindex](#)



OpenMP



High Performance Computing

The Serial Version of cpi

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

Start of OpenMP Version

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double x;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    #pragma omp parallel for
    for (i=0; i < steps; i++) {
        x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```



Need to Update our Makefile

```
CFLAGS=-Wall -O3 -ffast-math -mavx2 -fopenmp
```

```
CC=gcc
```

```
cpi          : cpi.o gettime.o
```

```
cpi.o        : cpi.c gettime.h
```

```
gettime.o    : gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpi cpi.o gettime.o
```



Not quite good enough

```
dhcp-94-191:cpi3$ make
gcc -Wall -O3 -ffast-math -mavx2 -fopenmp -c -o cpi.o cpi.c
gcc -Wall -O3 -ffast-math -mavx2 -fopenmp -c -o gettime.o gettime.c
gcc cpi.o gettime.o -o cpi
Undefined symbols for architecture x86_64:
  "_GOMP_parallel", referenced from:
      _main in cpi.o
  "_omp_get_num_threads", referenced from:
      _main._omp_fn.0 in cpi.o
  "_omp_get_thread_num", referenced from:
      _main._omp_fn.0 in cpi.o
ld: symbol(s) not found for architecture x86_64
collect2: error: ld returned 1 exit status
make: *** [cpi] Error 1
```



Makefile need linker flags too

```
CFLAGS=-Wall -O3 -ffast-math -mavx2 -fopenmp
```

```
LDFLAGS=-fopenmp
```

```
CC=gcc
```

```
cpi                : cpi.o gettime.o
```

```
cpi.o              : cpi.c gettime.h
```

```
gettime.o          : gettime.c gettime.h
```

```
clean:
```

```
    rm -f cpi cpi.o gettime.o
```

Better. Not good, but better.

```
dhcp-94-191:cpu3$ make  
gcc -fopenmp cpu.o gettime.o -o cpu  
dhcp-94-191:cpu3$ ./cpu  
PI = 0.5675882184166633 computed in 0.2785 seconds
```

(Laptop)

This is pi

This is not pi

```
dhcp-94-191:cpu3$ OMP_NUM_THREADS=1 ./cpu  
PI = 3.141592653589729 computed in 0.5467 seconds  
dhcp-94-191:cpu3$ OMP_NUM_THREADS=2 ./cpu  
PI = 1.287002217586625 computed in 0.2806 seconds  
dhcp-94-191:cpu3$ OMP_NUM_THREADS=4 ./cpu  
PI = 0.9799146525073925 computed in 0.2782 seconds  
dhcp-94-191:cpu3$ OMP_NUM_THREADS=4 ./cpu  
PI = 0.5675882184166633 computed in 0.279 seconds  
dhcp-94-191:cpu3$ OMP_NUM_THREADS=4 ./cpu  
PI = 0.9799146525073925 computed in 0.2772 seconds
```


Where is the problem?

```
#include <stdio.h>
#include "gettime.h"

static long steps = 1000000000;

int main (int argc, const char *argv[]) {
    int i;
    double pi;

    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();

    #pragma omp parallel for
    for (i=0; i < steps; i++) {
        double x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
    pi = step * sum;
    double delta = getTime() - start;
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
}
```

Where is the problem?

```
#include <stdio.h>
#include "gettime.h"
```

```
static long steps = 1000000000;
```

```
int main (int argc, const char *argv[]) {
    int i;
    double pi;
```

```
    double step = 1.0/(double) steps;
    double sum = 0.0;
    double start = getTime();
```

Like
MPI_Reduce

```
#pragma omp parallel for reduction(+ : sum)
```

```
    for (i=0; i < steps; i++) {
        double x = (i+0.5)*step;
        sum += 4.0 / (1.0+x*x);
    }
```

All others must be "handled"

```
    pi = step * sum;
```

```
    double delta = getTime() - start;
```

```
    printf("PI = %.16g computed in %.4g seconds\n", pi, delta);
```

```
}
```

Loop variable is fine



Now it all works

```
dhcp-94-191:cp13$ make  
gcc -Wall -O3 -ffast-math -mavx2 -fopenmp -c -o cpi.o cpi.c  
gcc -fopenmp cpi.o gettimeofday.o -o cpi
```



```
dhcp-94-191:cp13$ ./cpi  
PI = 3.141592653589774 computed in 0.2805 seconds
```

```
dhcp-94-191:cp13$ OMP_NUM_THREADS=4 ./cpi  
PI = 3.141592653589774 computed in 0.2776 seconds
```

```
dhcp-94-191:cp13$ OMP_NUM_THREADS=3 ./cpi  
PI = 3.141592653589908 computed in 0.2786 seconds
```

```
dhcp-94-191:cp13$ OMP_NUM_THREADS=2 ./cpi  
PI = 3.141592653589974 computed in 0.2841 seconds
```

```
dhcp-94-191:cp13$ OMP_NUM_THREADS=20 ./cpi  
PI = 3.141592653589802 computed in 0.2766 seconds
```




Parallel Pitfalls

What you're glad that
you don't need to know
(too much)




Critical Section: $\text{sum} += \text{value}$





Solution: Locking




Deadlock

- Think for some time
- Pick up left fork
- Pick up right fork
- Eat for some time
- Put down right fork
- Put down left fork
- Continue Thinking



Edsger W. Dijkstra



Requirements for a Deadlock

Mutual Exclusion

- Resource cannot be shared

Hold and Wait

- Must hold a resource while waiting for another

No Pre-emption

- Held resources cannot be given away


Circular Wait

- The wait dependency must be circular

Starvation

(No Hold and Wait)



- Think for some time
- Pick up left fork
- Right fork not available
- ... or left fork not available
- Put down left fork
- Think for a while
- Eventually try again



Resource Hierarchy (No Circular Wait)

- Number each resource
- Choose resources in order
 - Confucius is different
- Cannot deadlock
- Cannot starve?







Load Balancing

How to divide
up the work


Divide 100 rectangles
into four domains




Option 1







Option 2



Option 3



Which is the best choice?





The MPI version of cpi

```
MPI_Init(&argc,&argv); /* Connect processes to each other */
MPI_Comm_size(MPI_COMM_WORLD,&numprocs); /* Get total number of processes */
MPI_Comm_rank(MPI_COMM_WORLD,&myid); /* Rank of this process */
MPI_Get_processor_name(name, &resultlen);

printf("This is Process-%d/%d running on %s \n",myid,numprocs,name);
MPI_Barrier(MPI_COMM_WORLD);

if(myid == 0) {
    printf("This program uses %d processes\n", numprocs);
    n = 1000000000;
}


MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); /* Send n from 0 to all others */


sum = 0.0;
h = 1.0/n;
for (i=myid+0.5; i<n; i+=numprocs) {
    sum += dx_arctan(i*h);
}
mypi = 4.0*h*sum;

/* Add all partial sums to each other and send to rank 0 */
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);


if (myid == 0) {
    printf("pi is approximately %.16f\n",pi);
}

MPI_Finalize(); /* Disconnect all processes */
```






By Phonikal



By DaveSGage

Fast Fourier Transform



<http://2decomp.org>

Particle Data (N-Body)

