



University of Zurich^{UZH}

High Performance Computing | Lecture 8

Douglas Potter
Jozef Bucko
Noah Kubli

Compiler Suites

[Cray Compiler](#)
([PrgEnv-cray](#))

Available on Cray Computers

[GNU Compiler](#)
([PrgEnv-gnu](#))

Free for everyone

[Intel Compiler](#)
([PrgEnv-intel](#))

[Free for Students](#)

[Portland](#)
([PrgEnv-pgi](#))

Important for OpenACC (GPU)

[Clang](#)
(Apple & Cray)

Open Source & GNU Compatible

[Visual Studio](#)
(Windows)

Not usually High Performance

Compile with the Cray “wrappers”

```
dpotter@daint106:~> cc --version
Cray clang version 11.0.0  (1563bb12fb3eabb03515d4cc4aeaedd757aea56a)
Target: x86_64-unknown-linux-gnu
Thread model: posix
InstalledDir: /opt/cray/pe/cce/11.0.0/cce-clang/x86_64/share/../bin

dpotter@daint106:~> CC --version
Cray clang version 11.0.0  (1563bb12fb3eabb03515d4cc4aeaedd757aea56a)
Target: x86_64-unknown-linux-gnu
Thread model: posix
InstalledDir: /opt/cray/pe/cce/11.0.0/cce-clang/x86_64/share/../bin

dpotter@daint106:~> ftn --version
Cray Fortran : Version 11.0.0
dpotter@daint106:~>
```

Compile with the Cray “wrappers”

```
dpotter@daint106:~> module swap PrgEnv-cray PrgEnv-gnu
dpotter@daint106:~> cc --version
gcc (GCC) 10.1.0 20200507 (Cray Inc.)
Copyright (C) 2020 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
dpotter@daint106:~> CC --version
g++ (GCC) 10.1.0 20200507 (Cray Inc.)
...
```

```
dpotter@daint106:~> ftn --version
GNU Fortran (GCC) 10.1.0 20200507 (Cray Inc.)
...
```

Compile with the Cray “wrappers”

```
dpotter@daint106:~> module swap PrgEnv-gnu PrgEnv-pgi
dpotter@daint106:~> cc --version

pgcc (aka pgcc18) 20.1-0 LLVM 64-bit target on x86-64 Linux -tp haswell-64
PGI Compilers and Tools
Copyright (c) 2020, NVIDIA CORPORATION. All rights reserved.

dpotter@daint106:~> CC --version

pgc++ 20.1-0 LLVM 64-bit target on x86-64 Linux -tp haswell-64
...
dpotter@daint106:~> ftn --version

pgf90 20.1-0 LLVM 64-bit target on x86-64 Linux -tp haswell-64
...
```

On Some Systems you use the MPI Wrappers

```
ubuntu@dpotte-gpu:~$ mpicc --version
gcc (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0
Copyright (C) 2017 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
ubuntu@dpotte-gpu:~$ mpicxx --version
g++ (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0
...
```

```
ubuntu@dpotte-gpu:~$ mpif90 --version
GNU Fortran (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0
...
```

MPI compilation

Compilation directives:

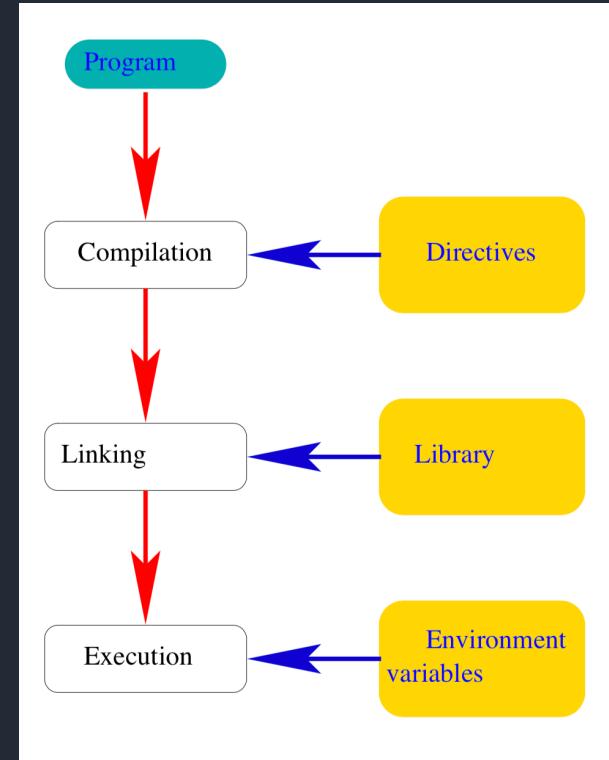
- You need to include “mpi.h” to call MPI functions.
 - This tells the compiler the parameters and their type.
- You need to tell the compiler where it can find “mpi.h” (`-I`).

Linking:

- You need to link against the MPI library (`-l`).
- You need to tell the linker where to find this library (`-L`).

Wrappers:

- The compiler wrapper (`cc` or `mpicc`) usually adds the necessary flags for you (`-L`, `-l`, and `-I`).



man gcc

-llibrary

Search the library named library when linking.

-Ldir

Add directory dir to the list of directories to be searched for -l.

-I dir

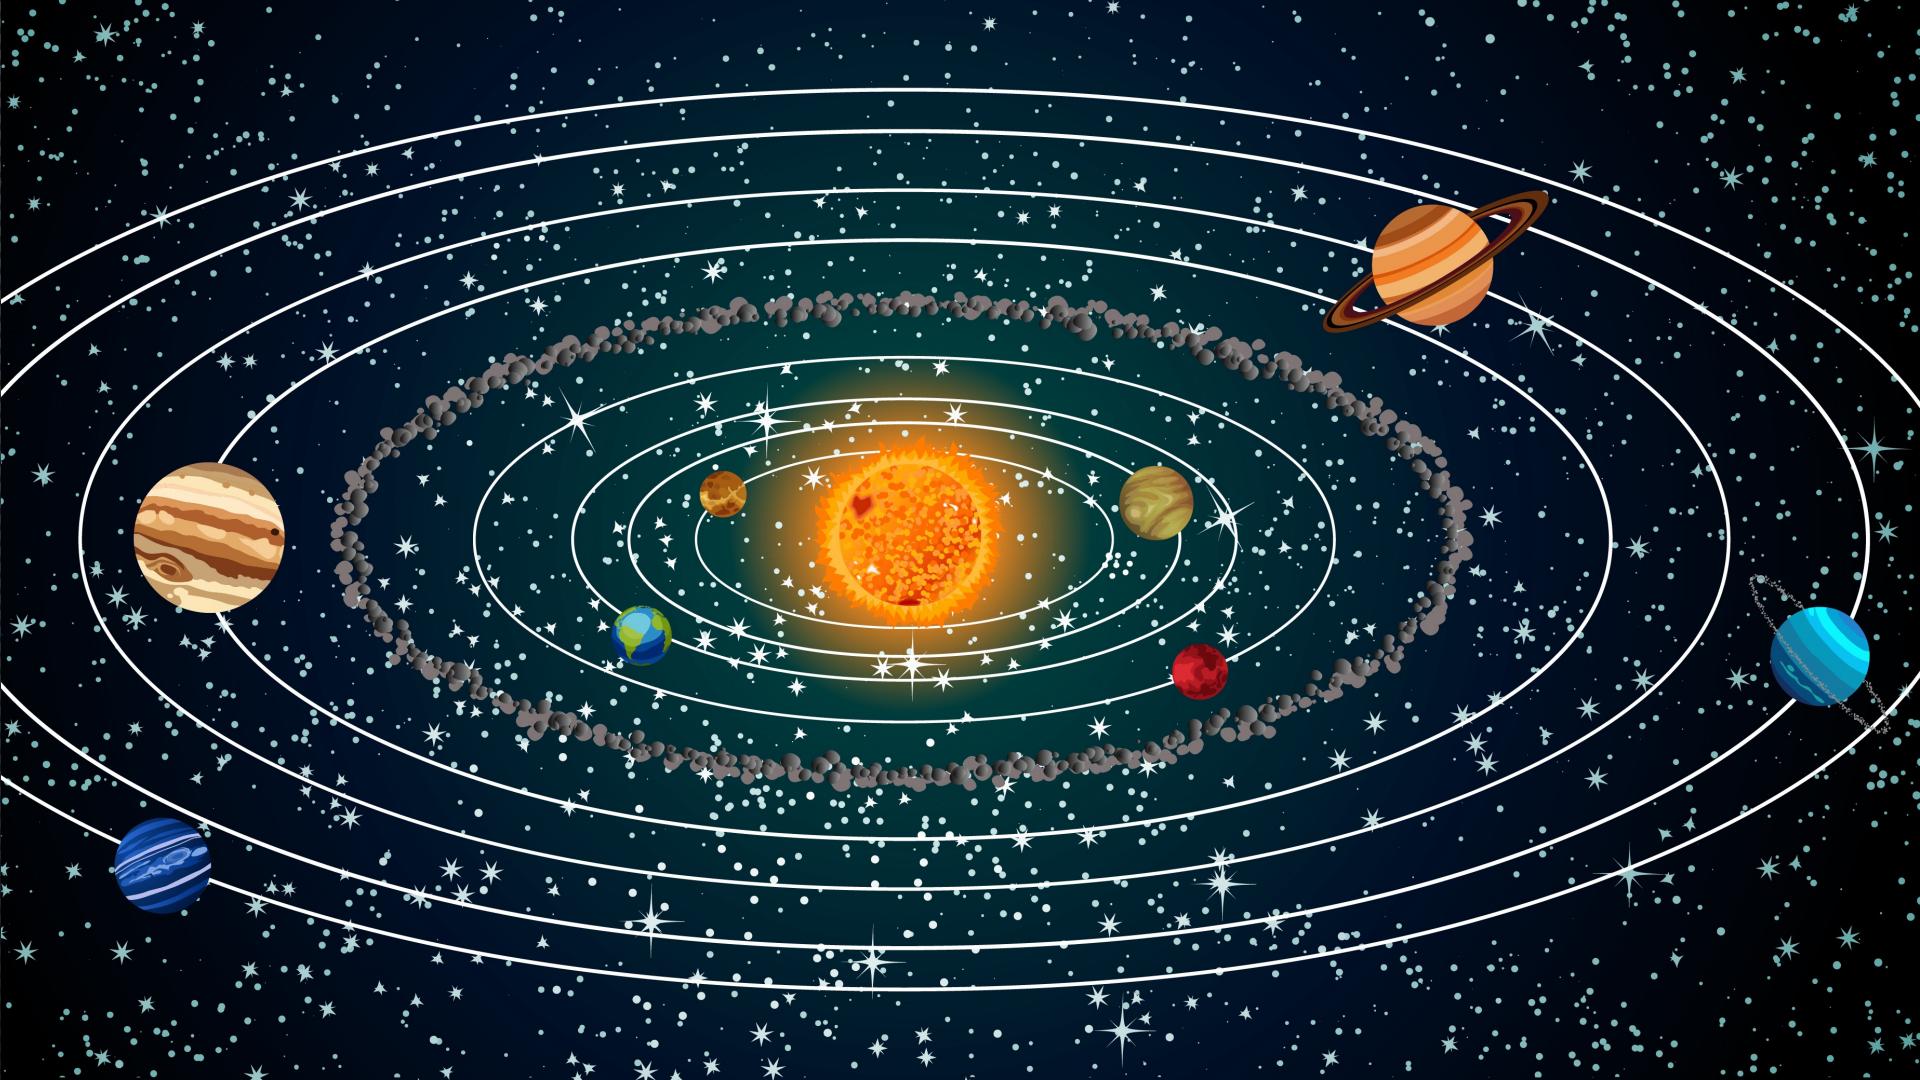
Add the directory dir to the list of directories to be searched for header files during preprocessing.

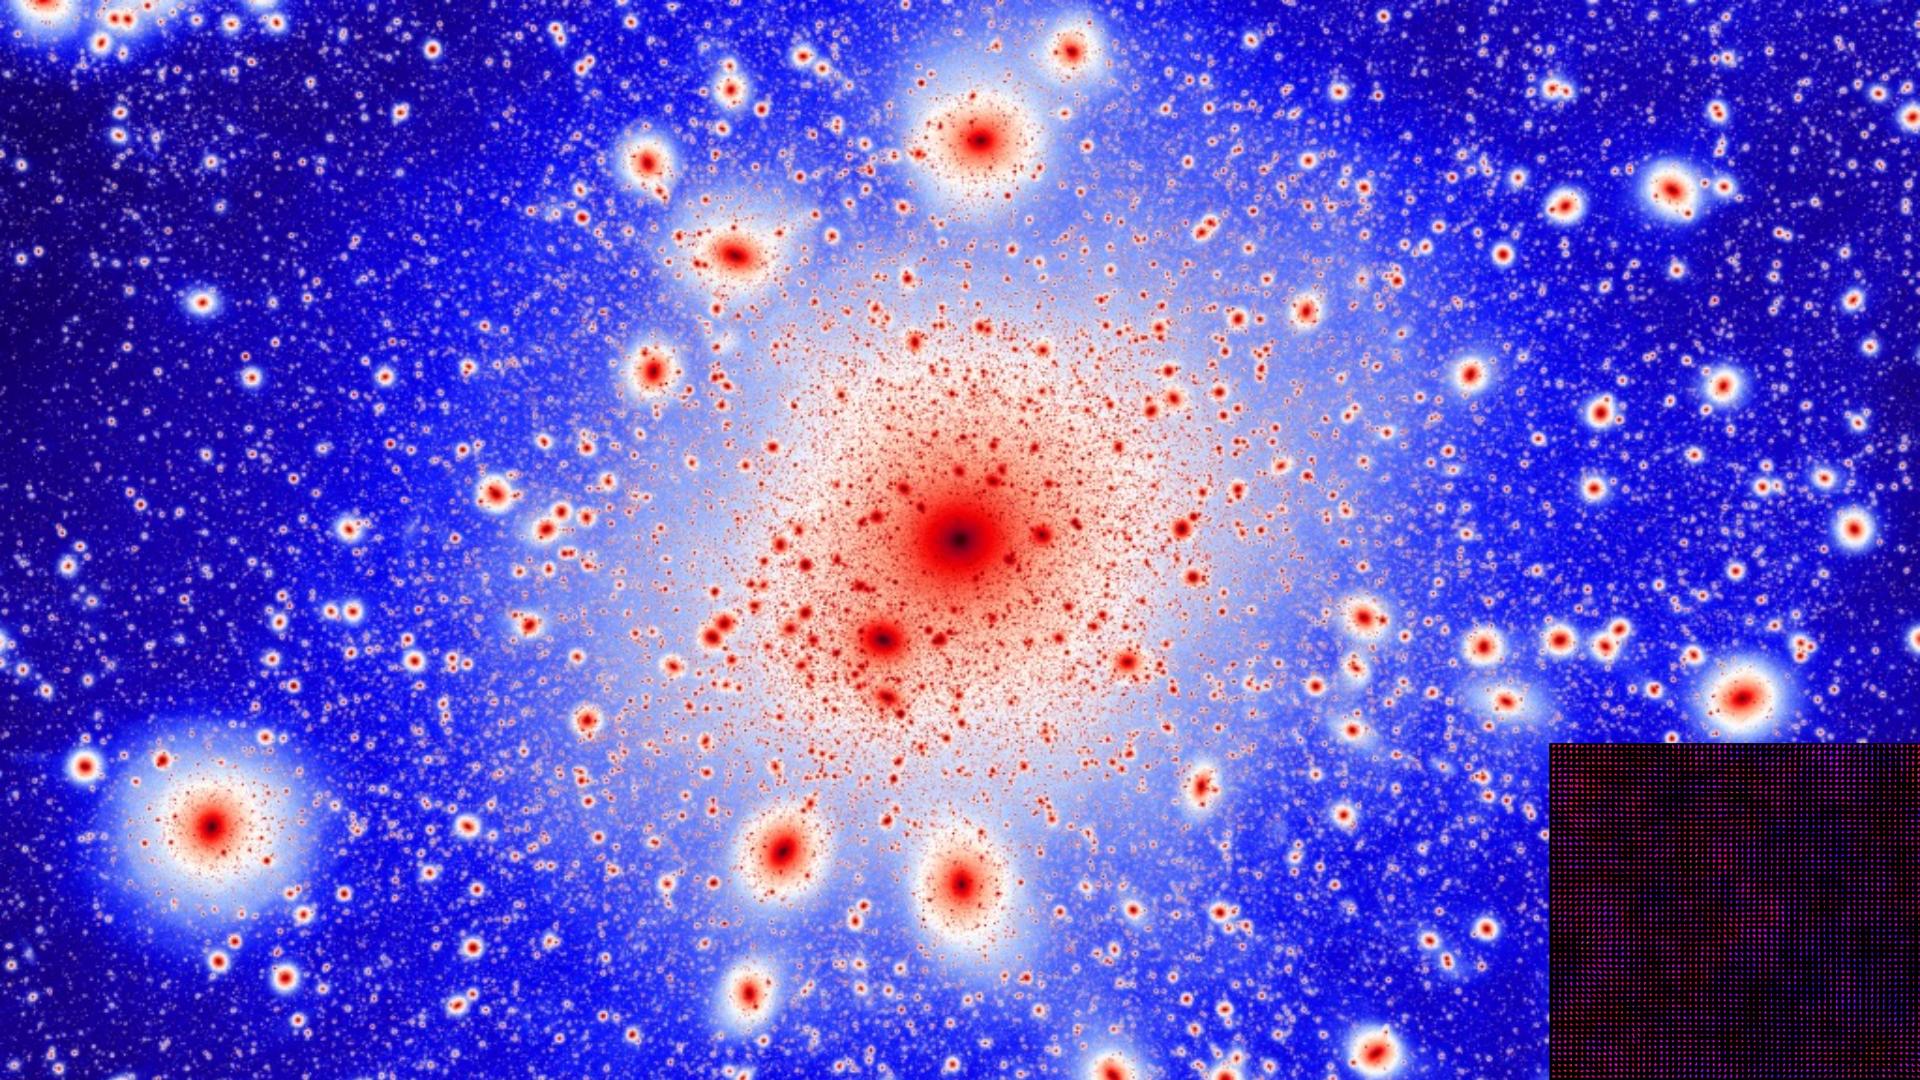
Example:

```
gcc -L/path/to/mpi/lib -lmpi -I/path/to/mpi/include -O3 -o mpitest mpitest.c
```

GRAVITY

THE N-BODY METHOD





EQUATIONS OF MOTION

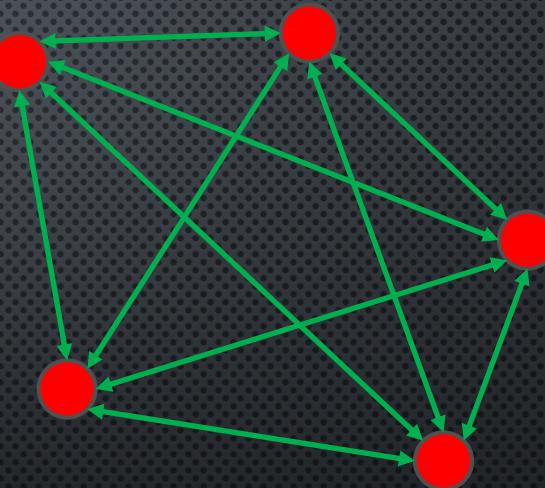
$$\mathbf{a} = \frac{d\mathbf{v}_i}{dt} = \sum_{j \neq i}^N Gm_j \frac{\mathbf{r}_j - \mathbf{r}_i}{|\mathbf{r}_j - \mathbf{r}_i|^3}$$

EQUATIONS OF MOTION

$$\mathbf{a} = \frac{d\mathbf{v}_i}{dt} = \sum_{j \neq i}^N \cancel{\mathbf{x}_i \mathbf{x}_j} \frac{\mathbf{r}_j - \mathbf{r}_i}{|\mathbf{r}_j - \mathbf{r}_i|^3}$$

METHOD

PAIRWISE INTERACTIONS



The n-body code

```
#include <vector>

struct particle {
    float x, y, z;      // position
    float vx, vy, vz;   // velocity
    float ax, ay, az;   // acceleration
};

typedef std::vector<particle> particles;
int main(int argc, char *argv[]) {
    int N=20'000;        // number of particles
    particles plist;    // vector of particles
    ic(plist,N);        // initialize starting position/velocity
    forces(plist);      // calculate the forces
    return 0;
}
```

Random Initial Conditions

```
#include <random>

// Initial conditions
void ic(particles &plist, int n) {
    std::random_device rd; //Will be used to obtain a seed for the random number engine
    std::mt19937 gen(rd()); //Standard mersenne_twister_engine seeded with rd()
    std::uniform_real_distribution<float> dis(0.0, 1.0);

    plist.clear();          // Remove all particles
    plist.reserve(n);       // Make room for "n" particle
    for( auto i=0; i<n; ++i) {
        particle p { dis(gen),dis(gen),dis(gen),0,0,0,0,0,0 } ;
        plist.push_back(p);
    }
}
```

Force Calculations

```

void forces(particles &plist) {
    int n = plist.size();
    for(int i=0; i<n; ++i) {                                // We want to calculate the force on all particles
        plist[i].ax = plist[i].ay = plist[i].az = 0;      // start with zero acceleration
        for(int j=0; j<n; ++j) {                          // Depends on all other particles
            if (i==j) continue; // Skip self interaction
            auto dx = plist[j].x - plist[i].x;
            auto dy = plist[j].y - plist[i].y;
            auto dz = plist[j].z - plist[i].z;
            auto r = sqrt(dx*dx + dy*dy + dz*dz);
            auto ir3 = 1 / (r*r*r);
            plist[i].ax += dx * ir3;
            plist[i].ay += dy * ir3;
            plist[i].az += dz * ir3;
        }
    }
}

```

$$\mathbf{a} = \frac{d\mathbf{v}_i}{dt} = \sum_{j \neq i}^N \cancel{\mathbf{r}_j} \frac{\mathbf{r}_j - \mathbf{r}_i}{|\mathbf{r}_j - \mathbf{r}_i|^3}$$

How does this perform? No idea!

```
dpotter@daint106:~/hpc1a/Y7> CC --version
g++ (GCC) 10.1.0 20200507 (Cray Inc.)
Copyright (C) 2020 Free Software Foundation, Inc.

This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
dpotter@daint106:~/hpc1a/Y7> CC -o nbody nbody.cxx
dpotter@daint106:~/hpc1a/Y7> ./nbody
dpotter@daint106:~/hpc1a/Y7>
```

The “time” command

```
dpotter@daint106:~/hpc1a/Y7> time ./nbody
```

```
real    0m20.061s
user    0m20.006s
sys     0m0.030s
```

```
dpotter@daint106:~/hpccla/Y7> /usr/bin/time ps -fu dpotter f
```

```
UID        PID  PPID  C STIME TTY          STAT   TIME  CMD
dpotter  24319 23916  0 10:08 ?          S      0:00 sshd: dpotter@pts/46
dpotter  24320 24319  0 10:08 pts/46    Ss     0:00  \_ -bash
dpotter  16118 24320  0 10:59 pts/46    S+    0:00      \_ /usr/bin/time ps -fu dpotter f
dpotter  16119 16118  0 10:59 pts/46    R+    0:00      \_ ps -fu dpotter f
0.01user 0.03system 0:00.06elapsed 85%CPU (0avgtext+0avgdata 5484maxresident)k
0inputs+0outputs (0major+635minor)pagefaults 0swaps
```

Oops, forgot to optimize

```
dpotter@daint106:~/hpc1a/Y7> CC -O3 -o nbody nbody.cxx
```

```
dpotter@daint106:~/hpc1a/Y7> time ./nbody
```

```
real    0m4.335s
```

```
user    0m4.303s
```

```
sys     0m0.021s
```

```
dpotter@daint106:~/hpc1a/Y7> CC -O3 -ffast-math -o nbody nbody.cxx
```

```
dpotter@daint106:~/hpc1a/Y7> time ./nbody
```

```
real    0m3.976s
```

```
user    0m3.928s
```

```
sys     0m0.037s
```

New compiler flag “-g” (man gcc)

GCC(1)

GNU

GCC(1)

NAME

```
gcc - GNU project C and C++ compiler
```

```
...
```

```
Options for Debugging Your Program
```

To tell GCC to emit extra information for use by a debugger, in almost all cases you need only to add `-g` to your other options.

GCC allows you to use `-g` with `-O`. The shortcuts taken by optimized code may occasionally be surprising: some variables you declared may not exist at all; flow of control may briefly move where you did not expect it; some statements may not be executed because they compute constant results or their values are already at hand; some statements may execute in different places because they have been moved out of loops. Nevertheless it is possible to debug optimized output. This makes it reasonable to use the optimizer for programs that might have bugs.

If you are not using some other optimization option, consider using `-Og` with `-g`. With no `-O` option at all, some compiler passes that collect information useful for debugging do not run at all, so that `-Og` may result in a better debugging experience.

`-g` Produce debugging information in the operating system's native format (stabs, COFF, XCOFF, or DWARF). GDB can work with this debugging information.

Effect of “-g” with gcc

```
dpotter@daint102:~/hpc1a/Y7> CC -g -O3 -ffast-math -o nbody nbody.cxx
dpotter@daint102:~/hpc1a/Y7> ls -l nbody
-rwxr-xr-x 1 dpotter uzh0 126464 Mar 22 10:53 nbody
dpotter@daint102:~/hpc1a/Y7> strip nbody
dpotter@daint102:~/hpc1a/Y7> ls -l nbody
-rwxr-xr-x 1 dpotter uzh0 17016 Mar 22 10:54 nbody
```

```
dpotter@daint102:~/hpc1a/Y7> CC -O3 -ffast-math -o nbody nbody.cxx
dpotter@daint102:~/hpc1a/Y7> ls -l nbody
-rwxr-xr-x 1 dpotter uzh0 25168 Mar 22 10:53 nbody
dpotter@daint102:~/hpc1a/Y7> strip nbody
dpotter@daint102:~/hpc1a/Y7> ls -l nbody
-rwxr-xr-x 1 dpotter uzh0 17016 Mar 22 10:55 nbody
```

Effect of “-og” with gcc

```
dpotter@daint106:~/hpc1a/Y7> CC -O3 -ffast-math -o nbody nbody.cxx  
dpotter@daint106:~/hpc1a/Y7> time ./nbody
```

```
real    0m3.976s  
user    0m3.928s  
sys     0m0.037s
```

```
dpotter@daint106:~/hpc1a/Y7> CC -Og -ffast-math -o nbody nbody.cxx  
dpotter@daint106:~/hpc1a/Y7> time ./nbody
```

```
real    0m4.862s  
user    0m4.819s  
sys     0m0.024s
```

Using Craypat – must use queue system

```
dpotter@daint106:~/hpc1a/Y7> module load perf-tools-lite
dpotter@daint106:~/hpc1a/Y7> CC -Og -ffast-math -o nbody nbody.cxx
WARNING: PerfTools is saving object files from a temporary directory into directory
'/users/dpotter/.craypat/nbody/18328'
INFO: creating the PerfTools-instrumented executable 'nbody' (lite-samples) ...OK
dpotter@daint106:~/hpc1a/Y7> time ./nbody
pat[FATAL] [0]: PMI_Init failed

real    0m2.144s
user    0m0.022s
sys     0m0.084s
```

Using Craypat

```
dpotter@daint106:~/hpc1a/Y7> cat nbody.job
#!/bin/bash -l
#SBATCH --account=uzg2
#SBATCH --job-name=nbody
#SBATCH --time=00:5:00
#SBATCH --nodes=1
#SBATCH --ntasks-per-core=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --partition=debug
#SBATCH --constraint=gpu

srun nbody
dpotter@daint106:~/hpc1a/Y7> sbatch nbody.job
Submitted batch job 30052444
```

Using Craypat

```
CrayPat/X: Version 20.10.0 Revision 7ec62de47 09/16/20 16:54:26
```

```
DRAM: 64 GiB DDR4-2400 on 2.6 GHz nodes
```

```
#####
#                               #
#      CrayPat-lite Performance Statistics      #
#                               #
#####
```

Avg Process Time:	4.38 secs
High Memory:	22.0 MiBytes 22.0 MiBytes per PE
Observed CPU clock boost:	126.5 %
Percent cycles stalled:	59.7 %
Instr per Cycle:	1.38

```
CrayPat/X: Version 20.10.0 Revision 7ec62de47 09/16/20 16:54:26
```

```
Experiment:          lite  lite-samples
```

```
Number of PEs (MPI ranks):    1
```

```
Numbers of PEs per Node:      1
```

```
Numbers of Threads per PE:    1
```

```
Number of Cores per Socket:   12
```

```
Execution start time: Mon Mar 22 10:30:17 2021
```

```
System name and speed: nid02356 2.601 GHz (nominal)
```

```
Intel Haswell    CPU Family: 6 Model: 63 Stepping: 2
```

Using Craypat

Notes for table 1:

This table shows functions that have significant exclusive sample hits, averaged across ranks.

For further explanation, see the "General table notes" below,

or use: pat_report -v -O samp_profile ...

Table 1: Profile by Function

Samp%	Samp	Imb.	Imb.	Group
	Samp	Samp%		Function=[MAX10]
100.0%	433.0	--	--	Total
<hr/>				
100.0%	433.0	--	--	USER
<hr/>				
83.1%	360.0	--	--	forces
16.9%	73.0	--	--	std::vector<>::operator[]
<hr/>				

Using Craypat

Notes for table 2:

This table shows functions, and line numbers within functions, that have significant exclusive sample hits, averaged across ranks.
 For further explanation, see the "General table notes" below,
 or use: pat_report -v -O samp_profile+src ...

Table 2: Profile by Group, Function, and Line

Samp%	Samp	Imb.	Imb.	Group
		Samp	Samp%	Function=[MAX10]
				Source
				Line

						Total

	100.0%	433.0	--	--	--	USER

	83.1%	360.0	--	--	--	forces

3	80.1%	347.0	--	--	--	dpotter/hpclab/Y7/nbody.cxx

4	2.8%	12.0	--	--	--	line.32
4	1.8%	8.0	--	--	--	line.36
4	4.4%	19.0	--	--	--	line.37
4	65.4%	283.0	--	--	--	line.39
4	3.2%	14.0	--	--	--	line.40
4	2.5%	11.0	--	--	--	line.41

Using Craypat

```
dpotter@daint106:~/hpc1a/Y7> cat -n nbody.cxx | sed -n 28,43p
```

```
28     void forces(particles &plist) {
29         int n = plist.size();
30         for(int i=0; i<n; ++i) {
31             plist[i].ax = plist[i].ay = plist[i].az = 0;
32             for(int j=0; j<n; ++j) {
33                 if (i==j) continue;
34                 auto dx = plist[j].x - plist[i].x;
35                 auto dy = plist[j].y - plist[i].y;
36                 auto dz = plist[j].z - plist[i].z;
37                 auto r = sqrt(dx*dx + dy*dy + dz*dz);          4||| 4.4% | 19.0 | -- | -- | line.37
38                 auto ir3 = 1 / (r*r*r);
39                 plist[i].ax += dx * ir3;                      4||| 65.4% | 283.0 | -- | -- | line.39
40                 plist[i].ay += dy * ir3;                      4||| 3.2% | 14.0 | -- | -- | line.40
41                 plist[i].az += dz * ir3;
42             }
43         }
```

Sampling

`PAT_RT_SAMPLING_INTERVAL`

Specifies the interval, in microseconds, at which the instrumented executable is sampled.

To specify a random interval, use the following format:

`lower-bound,upper-bound[,seed]`

After a sample is captured, the interval used for the next sampling interval is generated using `rand(3)` and will be between lower-bound and upper-bound. The initial seed (seed) for the sequence of random numbers is optional. See `srand(3)` for more information.

This environment variable affects sampling experiments. It can also be used to control trace-enhanced sampling experiments, provided the program is instrumented for tracing but the `PAT_RT_EXPERIMENT` environment variable is used to specify a sampling-type experiment, and subject to the `PAT_RT_SAMPLING_MODE` environment variable setting.

Default: 10000 (microseconds)

Using Craypat without optimization “-O0”

34	auto dx = plist[j].x - plist[i].x;	4	3.1%		50.0		--		--		line.34
35	auto dy = plist[j].y - plist[i].y;	4	3.5%		56.0		--		--		line.35
36	auto dz = plist[j].z - plist[i].z;	4	3.6%		59.0		--		--		line.36
37	auto r = sqrt(dx*dx + dy*dy + dz*dz);	4	6.0%		97.0		--		--		line.37
38	auto ir3 = 1 / (r*r*r);	4	25.9%		419.0		--		--		line.38
39	plist[i].ax += dx * ir3;	4	11.2%		181.0		--		--		line.39
40	plist[i].ay += dy * ir3;	4	2.8%		46.0		--		--		line.40
41	plist[i].az += dz * ir3;	4	3.8%		62.0		--		--		line.41

Craypat summary

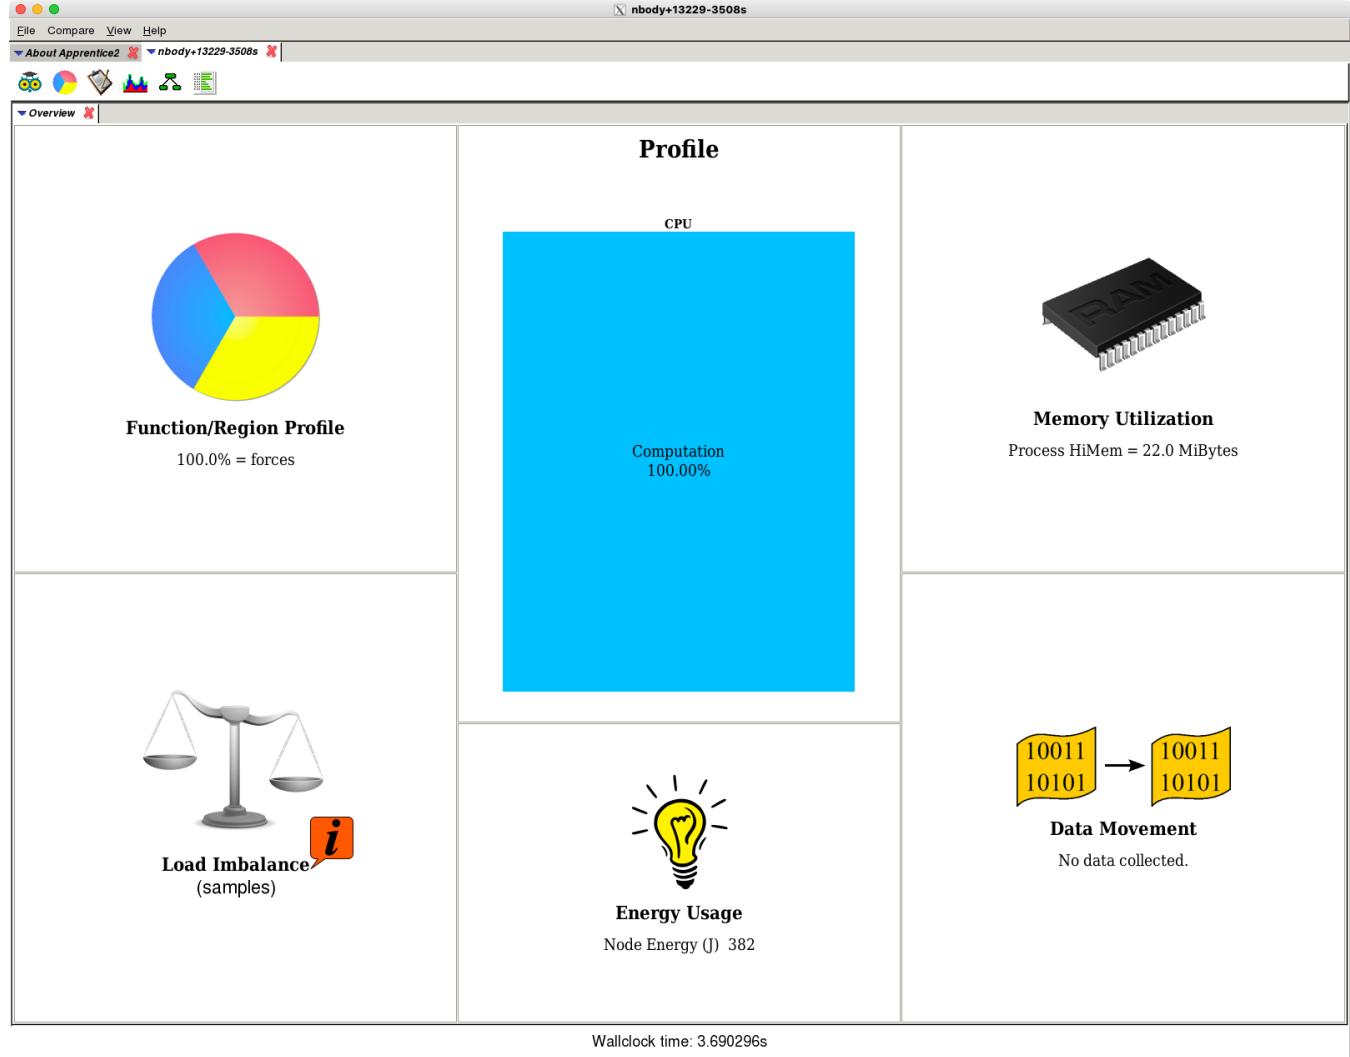
- Load “perftools-lite” module (or “perftools” for advanced)
- Compile with “-g” for debug information
- Consider effect of optimization
- Run using the batch queue system

app2

```
dpotter@daint105:~/hpc1a/Y7> ls -ld nbody+13229-3508s/
drwxr-x--- 5 dpotter uzh0 8192 Mar 22 11:31 nbody+13229-3508s/
dpotter@daint105:~/hpc1a/Y7> ls -l nbody+13229-3508s/
total 171
drwxr-x--- 2 dpotter uzh0 4096 Mar 22 11:31 ap2-files
-rw-r--r-- 1 dpotter uzh0 167936 Mar 22 11:31 index.ap2
drwxr-x--- 2 dpotter uzh0 4096 Mar 22 11:31 rpt-files
drwxr-x--- 2 dpotter uzh0 4096 Mar 22 11:31 xf-files
dpotter@daint105:~/hpc1a/Y7> app2 nbody+13229-3508s/
```



app2





VTUNE™ PROFILER PERFORMANCE ANALYZER

Intel, and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries. *Other names and brands may be claimed as the property of others. © Intel Corporation

High Performance Computing



Intel VTune Profiler

Project Navigator

sample (matrix)
r000hs
r001ue

Configure Analysis

WHERE

Local Host ▾

WHAT

Launch Application ▾

Specify and configure your analysis target: an application or a script to execute.

Application:

/home/ubuntu/code/nbody/nbody

Application parameters:

Use application directory as working directory

Advanced

HOW

Hotspots ▾

Identify the most time consuming functions and drill down to see time spent on each line of source code. Focus optimization efforts on hot code for the greatest performance impact. [Learn more](#)

⚠ Hardware collection of CPU events is not possible on this system. Microarchitecture performance insights will not be available.

User-Mode Sampling ⓘ

Hardware Event-Based Sampling ⓘ

Overhead

Show additional performance insights

Details

▶ ⏪ ⏴ ⏵ ⏵

High Performance Computing



Intel VTune Profiler

Project Navigator

sample (matrix)

r000hs
r001ue
r002hs

Hotspots Hotspots by CPU Utilization

Analysis Configuration Collection Log Summary Bottom-up Caller/Callee Top-down Tree Platform

Elapsed Time: 3.561s

CPU Time: 3.550s
Total Thread Count: 1
Paused Time: 0s

Top Hotspots

This section lists the most active functions in your application. Optimizing these hotspot functions typically results in improving overall application performance.

Function	Module	CPU Time
forces	nbody	3.550s

*N/A is applied to non-summable metrics.

Effective CPU Utilization Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.

Simultaneously Utilized Logical CPUs

High Performance Computing

Hotspots Insights

If you see significant hotspots in the Top Hotspots list, switch to the Bottom-up view for in-depth analysis per function. Otherwise, use the Caller/Callee view to track critical paths for these hotspots.

Explore Additional Insights

Parallelism : 3.1% ↗

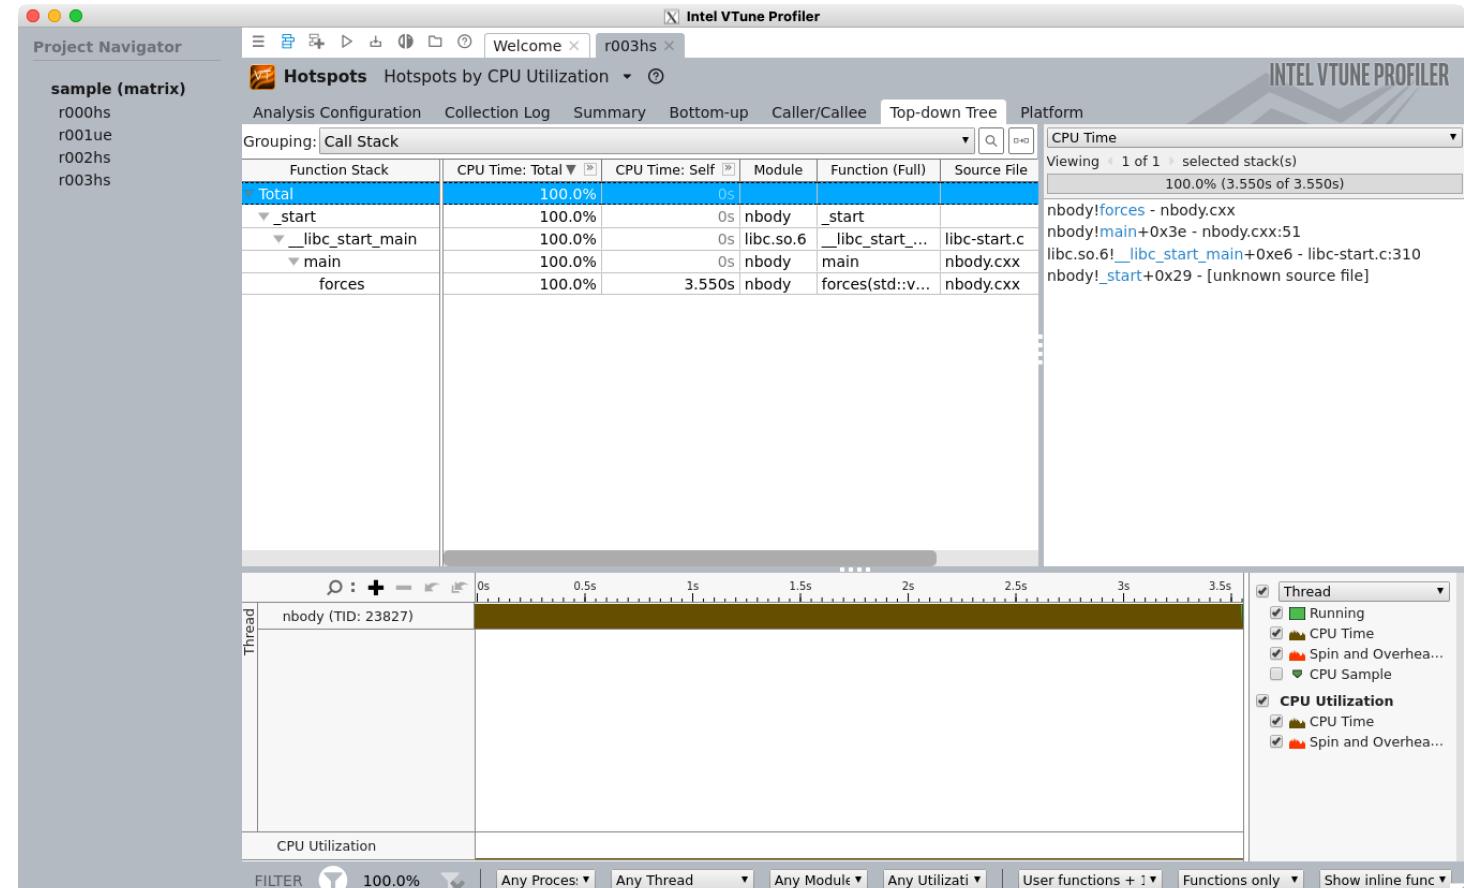
Use Threading to explore more opportunities to increase parallelism in your application.

INTEL VTUNE PROFILER

INSIGHTS



VTUNE





Intel VTune Profiler

Project Navigator Welcome x r003hs x

Hotspots Hotspots by CPU Utilization

Analysis Configuration Collection Log Summary Bottom-up Caller/Callee Top-down Tree Platform nbody.cxx x

Source Assembly

Source	CPU Time: Total	CPU Time: Self
23 }		
24		
25 }		
26		
27		
28 void forces(particles &plist) {		
29 int n = plist.size();		
30 #pragma omp parallel for		
31 for(int i=0; i<n; ++i) { // We want to calculate the force on all particles		
32 plist[i].ax = plist[i].ay = plist[i].az = 0; // start with zero acceleration		
33 for(int j=0; j<n; ++j) { // Depends on all other particles	6.8%	0.240s
34 if (i==j) continue; // Skip self interaction		
35 auto dx = plist[j].x - plist[i].x;		
36 auto dy = plist[j].y - plist[i].y;		
37 auto dz = plist[j].z - plist[i].z;	1.8%	0.064s
38 auto r = sqrt(dx*dx + dy*dy + dz*dz);		
39 auto ir3 = 1 / (r*r*r);	12.7%	0.452s
40 plist[i].ax += dx * ir3;	45.7%	1.622s
41 plist[i].ay += dy * ir3;	14.4%	0.512s
42 plist[i].az += dz * ir3;	18.6%	0.660s
43 }		
44 }		
45 }		
46		
47 int main(int argc, char *argv[]) {		
48 int N=20'000;		
49 particles plist;		
50 ic(plist,N);		
51 forces(plist);		
52 return 0;		
53 }		



Intel VTune Profiler

Project Navigator

sample (matrix)

r000hs
r001ue
r002hs
r003hs
r004hs

Hotspots Hotspots by CPU Utilization

Analysis Configuration Collection Log Summary Bottom-up Caller/Callee Top-down Tree Platform

Elapsed Time : 0.327s

CPU Time : 6.708s

Total Thread Count: 32
Paused Time: 0s

Top Hotspots

This section lists the most active functions in your application. Optimizing these hotspot functions typically results in improving overall application performance.

Function	Module	CPU Time
forces_omp_fn.0	nbody	5.557s
func@0x19070	libgomp.so.1	0.561s
func@0x18ee0	libgomp.so.1	0.552s
GOMP_parallel	libgomp.so.1	0.020s
func@0x17db0	libgomp.so.1	0.018s

*N/A is applied to non-summable metrics.

Effective CPU Utilization Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU utilization value.

Elapsed Time

Simultaneously Utilized Logical CPUs

Average Effective CPU Utilization

Target Utilization

Idle

Poor

Ok

Ideal

Hotspots Insights

If you see significant hotspots in the Top Hotspots list, switch to the Bottom-up view for in-depth analysis per function. Otherwise, use the Caller/Callee view to track critical paths for these hotspots.

Explore Additional Insights

Parallelism : 64.1% ↗

Use **Threading** to explore more opportunities to increase parallelism in your application.

INSIGHTS



Self timing

- Basic concept
 - Get the current time
 - Do some calculations
 - Get the new time
 - The difference is how long it took
- What function(s) can we use?



time()

TIME(2)

Linux Programmer's Manual

TIME(2)

NAME

time - get time in seconds

SYNOPSIS

```
#include <time.h>

time_t time(time_t *tloc);
```

DESCRIPTION

time() returns the time as the number of **seconds** since the Epoch, 1970-01-01 00:00:00 +0000 (UTC).

If tloc is non-NULL, the return value is also stored in the memory pointed to by tloc.

RETURN VALUE

On success, the value of time in seconds since the Epoch is returned. On error, ((time_t) -1) is returned, and errno is set appropriately.



gettimeofday()

GETTIMEOFDAY (2)
GETTIMEOFDAY (2)

Linux Programmer's Manual

NAME

gettimeofday, gettimeofday - get / set time

SYNOPSIS

```
#include <sys/time.h>
int gettimeofday(struct timeval *tv, struct timezone *tz);
int settimeofday(const struct timeval *tv, const struct timezone *tz);
```

DESCRIPTION

The functions gettimeofday() and settimeofday() can get and set the time as well as a timezone.

The tv argument is a struct timeval (as specified in <sys/time.h>), and gives the number of seconds and **microseconds** since the Epoch (see time(2)). The tz argument is a struct timezone:

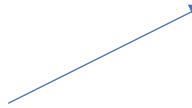
```
struct timezone {
    int tz_minuteswest;      /* minutes west of Greenwich */
    int tz_dsttime;          /* type of DST correction */
};
```



gettime.c

```
#include <sys/time.h>

double getTime(void) {
    struct timeval tv;
    struct timezone tz;
    gettimeofday(&tv, &tz);
    return tv.tv_sec + 1e-6*(double)tv.tv_usec;
}
```



```
struct timeval {
    time_t          tv_sec;      /* seconds since Jan. 1, 1970 */
    suseconds_t     tv_usec;     /* and microseconds */
};
```



clock_gettime() and clock_getres()

SYNOPSIS

```
#include <time.h>
int clock_getres(clockid_t clk_id, struct timespec *res);
int clock_gettime(clockid_t clk_id, struct timespec *tp);
```

DESCRIPTION

The function `clock_getres()` finds the resolution (precision) of the specified clock `clk_id`, and, if `res` is non-NULL, stores it in the `struct timespec` pointed to by `res`. The resolution of clocks depends on the implementation and cannot be configured by a particular process. If the time value pointed to by the argument `tp` of `clock_gettime()` is not a multiple of `res`, then it is truncated to a multiple of `res`.

The functions `clock_gettime()` and `clock_settime()` retrieve and set the time of the specified clock `clk_id`.

All implementations support the system-wide real-time clock, which is identified by `CLOCK_REALTIME`. Its time represents seconds and nanoseconds since the Epoch. When its time is changed, timers for a relative interval are unaffected, but timers for an absolute point in time are affected.

Sufficiently recent versions of glibc and the Linux kernel support the following clocks:

`CLOCK_MONOTONIC`

Clock that cannot be set and represents monotonic time since some unspecified starting point. This clock is not affected by discontinuous jumps in the system time (e.g., if the system administrator manually changes the clock), but is affected by the incremental adjustments performed by `adjtime(3)` and NTP.



High resolution C++ timer

std::chrono::high_resolution_clock

Defined in header <chrono>

class high_resolution_clock; (since C++11)

Class std::chrono::high_resolution_clock represents the clock with the smallest tick period provided by the implementation.

It may be an alias of std::chrono::system_clock or std::chrono::steady_clock, or a third, independent clock.

std::chrono::high_resolution_clock meets the requirements of TrivialClock.



Example C++ program

```
dpotter@daint102:~/hpc1a/Y7> cat c++clock.cxx
#include <iostream>
#include <chrono>
int main(int argc, char*argv[]) {
    using namespace std::chrono;
    using clock = high_resolution_clock;
    double tick = 1.0 * clock::period::num / clock::period::den;
    auto t1 = clock::now();
    std::cout << "Timer resolution is " << tick << " seconds" << std::endl;
    auto t2 = clock::now();
    auto elapsed = duration<duration<double>>(t2 - t1);
    std::cout << "Elapsed time " << elapsed.count() << " seconds" << std::endl;
    return 0;
}

dpotter@daint102:~/hpc1a/Y7> CC -O3 -o c++clock c++clock.cxx
dpotter@daint102:~/hpc1a/Y7> ./c++clock
Timer resolution is 1e-09 seconds
Elapsed time 4.1871e-05 seconds
dpotter@daint102:~/hpc1a/Y7> ./c++clock
Timer resolution is 1e-09 seconds
Elapsed time 3.7361e-05 seconds
```



Example Python

Python 3.9.1 (default, Jan 8 2021, 17:17:43)

[Clang 12.0.0 (clang-1200.0.32.28)] on darwin

Type "help", "copyright", "credits" or "license" for more information.

```
>>> import time
```

```
>>> help(time.perf_counter)
```

perf_counter(...)

perf_counter() -> float

Performance counter for benchmarking.

```
>>> t1=time.perf_counter() ; t2 = time.perf_counter() ; print(t1,t2,t2-t1)
```

31.597328265 31.597329154 8.889999989492026e-07



Advanced -- “cycle.h”

```
/********************************************/  
/* To use the cycle counters in your code, simply #include "cycle.h" (this  
file), and then use the functions/macros:  
  
    ticks getticks(void);  
  
ticks is an opaque typedef defined below, representing the current time.  
You extract the elapsed time between two calls to gettick() via:  
  
    double elapsed(ticks t1, ticks t0);  
  
which returns a double-precision variable in arbitrary units. You  
are not expected to convert this into human units like seconds; it  
is intended only for *comparisons* of time intervals.  
  
(In order to use some of the OS-dependent timer routines like  
Solaris' gethrtime, you need to paste the autoconf snippet below  
into your configure.ac file and #include "config.h" before cycle.h,  
or define the relevant macros manually if you are not using autoconf.)  
*/
```



One possible implementation

```
/*
 * X86-64 cycle counter
 */

#ifndef __GNUC,__ICC__ || __x86_64__
typedef unsigned long long ticks;

static __inline__ ticks getticks(void) {
    unsigned a, d;
asm volatile("rdtsc" : "=a" (a), "=d" (d));
    return ((ticks)a) | (((ticks)d) << 32);
}
```

Time Stamp Counter

From Wikipedia, the free encyclopedia

The **Time Stamp Counter (TSC)** is a 64-bit register present on all [x86](#) processors since the [Pentium](#). It counts the number of [cycles](#) since reset. The instruction [RDTSC](#) returns the TSC in EDX:EAX. In [x86-64](#) mode, [RDTSC](#) also clears the higher 32 bits of [RAX](#) and [RDX](#). Its [opcode](#) is [0F 31](#).^[1] Pentium competitors such as the [Cyrix 6x86](#) did not always have a TSC and may consider [RDTSC](#) an illegal instruction. Cyrix included a Time Stamp Counter in their [MII](#).



Comparison of time

Gravity Calculated, Wallclock: 147.719104 secs

Gflops:2671816.5, Total Gflop:3.95e+08

% computing:	max=	97.957	avg=	87.266	of 48400	std-dev=	2.396
% waiting:	max=	6.545	avg=	3.399	of 48400	std-dev=	0.812
% syncing:	max=	16.694	avg=	9.335	of 48400	std-dev=	2.280



Summary

- Compiler Wrappers (cc,CC,ftn,mpicc,...)
- Timing a program (**time** command)
- Profiling using sampling (craypat)
- Other alternatives
 - Intel VTUNE
 - AMD μProf
- Build in instrumentation (for phases)

Data Layout (C++) – Indexes start at zero!

```
blitz::Array<float,2> r(1000,2)
blitz::Array<float,1> d2(1000);

for(i=0; i<1000; ++i) {
    d2(i) = r(i,0)*r(i,0) + r(i,1)*r(i,1);
}
```

Data Layout (Fortran) – Indexes start at one!

```
real, dimension(1000,2) :: r
```

```
real, dimension(1000) :: d2
```

```
do i = 1, 1000
```

```
    d2(i) = r(i,1)*r(i,1) + r(i,2)*r(i,2)
```

```
end do
```

$$d^2 = x^2 + y^2$$

Data Layout (Fortran)

```
real, dimension(2,1000) :: r
```

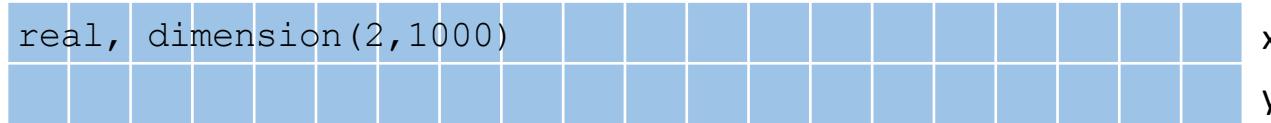
```
real, dimension(1000) :: d2
```

```
do i = 1, 1000
```

```
    d2(i) = r(1,i)*r(1,i) + r(2,i)*r(2,i)
```

```
end do
```

Data Layout (Fortran)



real, dimension(1000,2)

